

# CyCLaDEs: A Decentralized Cache for Triple Pattern Fragments

**Pauline Folz, Hala Skaf-Molli & Pascal Molli**  
ESWC 2016 - May 2016 - Heraklion



UNIVERSITÉ DE NANTES



# Context

LDF clients execute SPARQL queries over LDF servers <sup>[1]</sup>

- **Minimize server side processing** to simple **triple pattern fragment queries (TPF)**
- **Joins** executed on **clients** side



[1] R. Verborgh and al. Querying Datasets on the Web with High Availability. In *13th International Semantic Web Conference (ISWC 2014)*, pages 180-196, 2014.

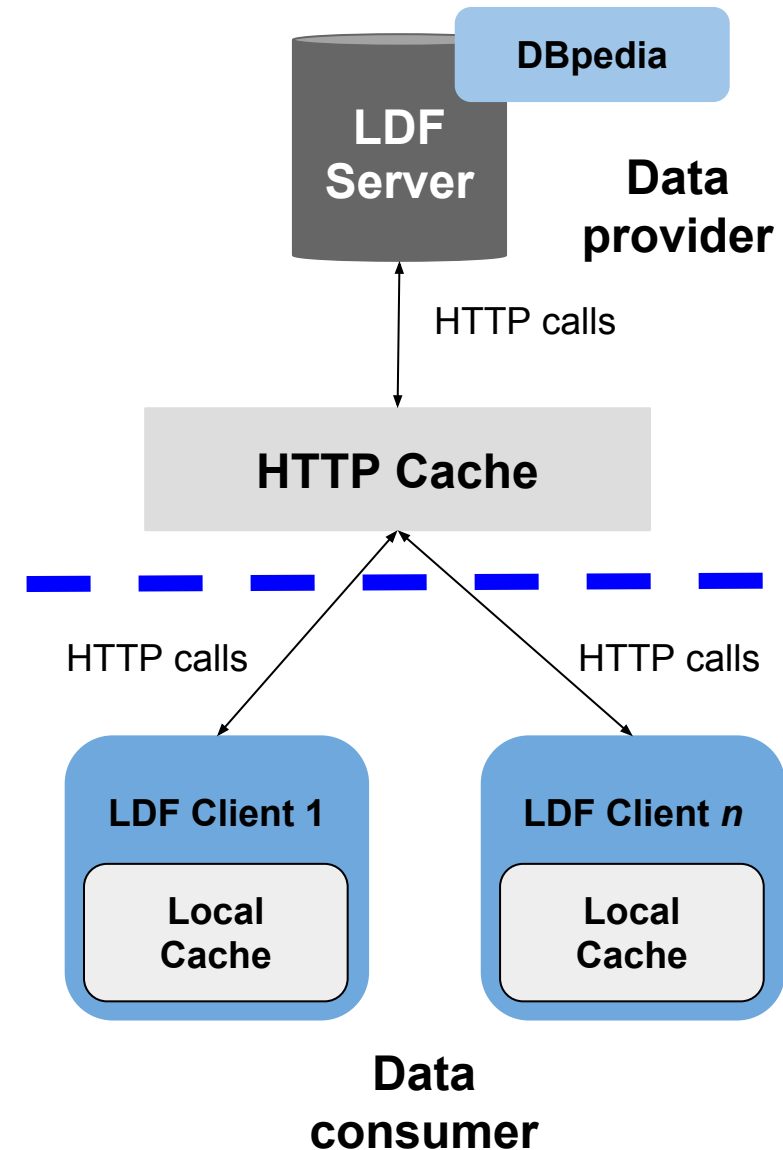
# Context

**Caches** play an important role in the **performances** of LDF server:

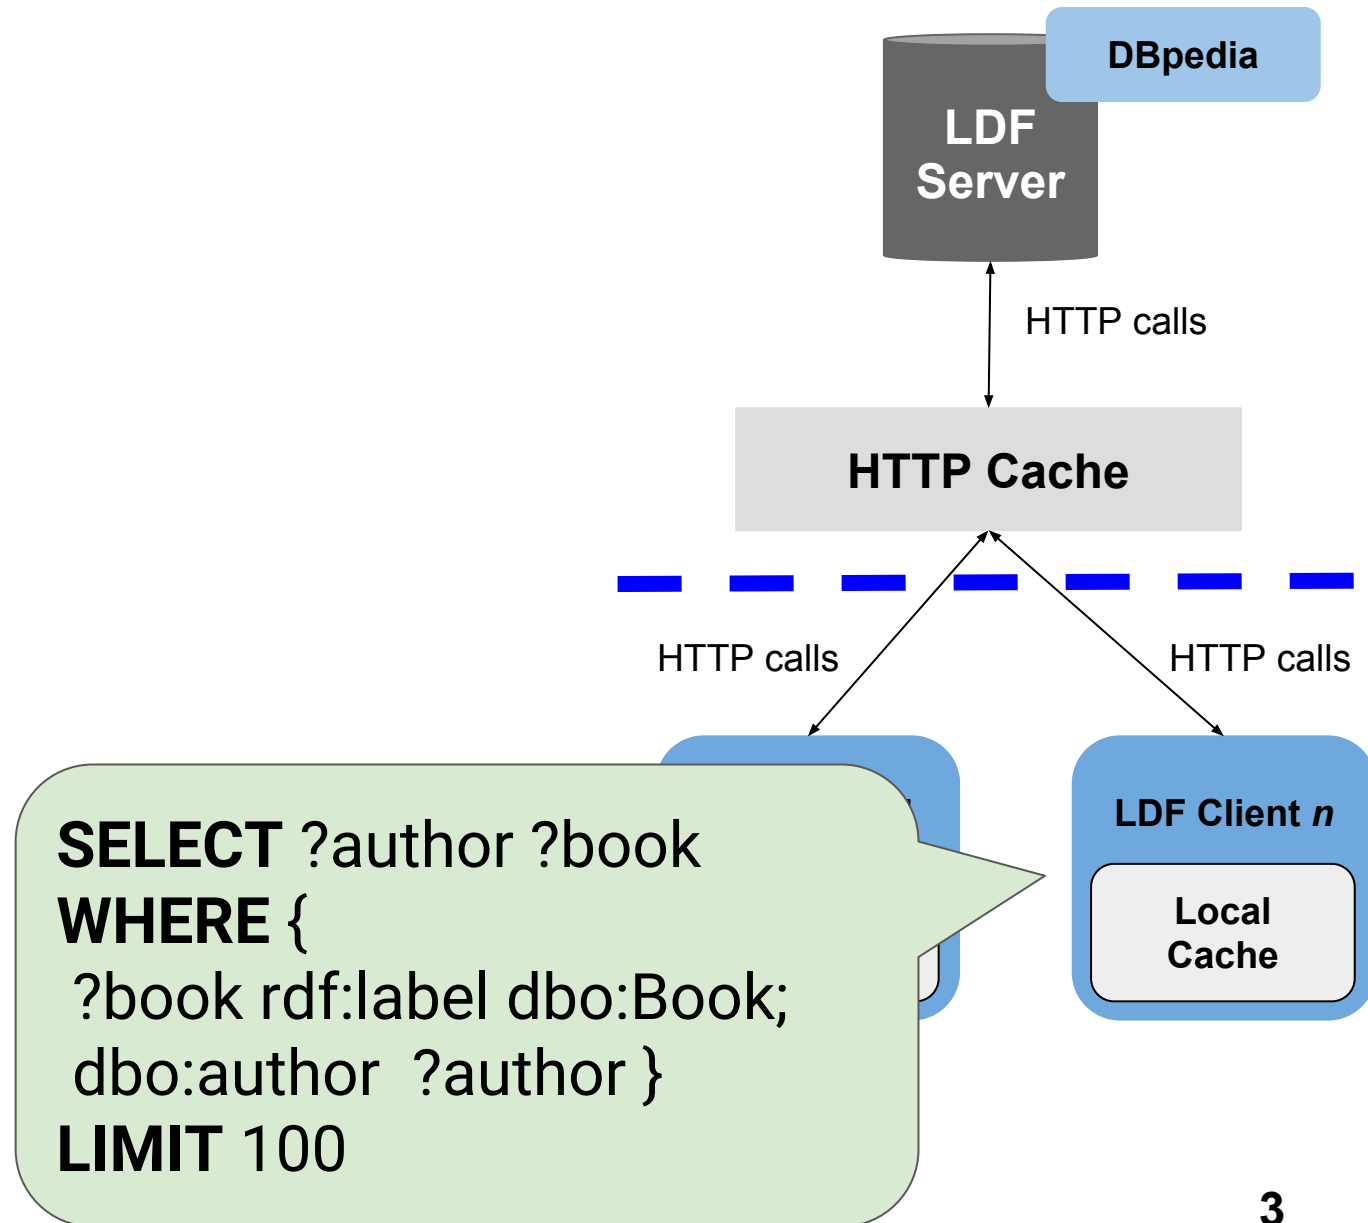
- Caches contain TPF, TPF are more likely to be reused locally and across clients



# Triple Pattern Fragment & Caches



# Triple Pattern Fragment & Caches

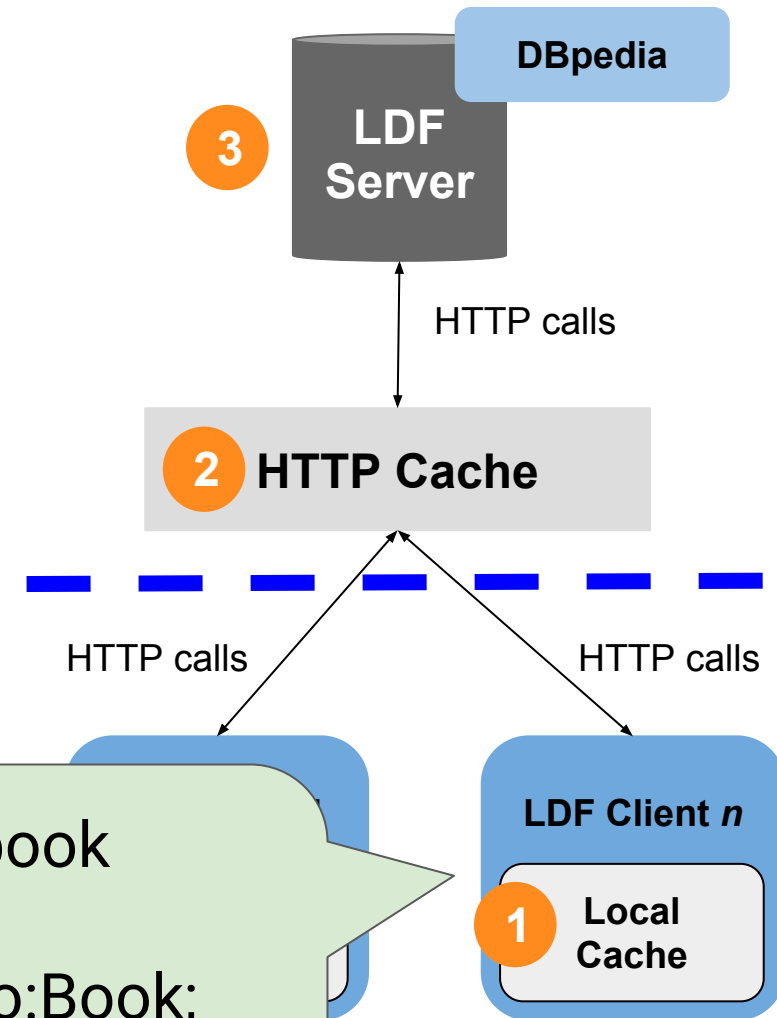


# Triple Pattern Fragment & Caches

152 HTTP calls

0	?book <b>http://www.w3.org/1999/02/22-rdf-syntaxns#type</b> http:.../ontology/Book
1	?book <b>http:.../ontology/author</b> ?author
2	http:.../resource/%22...And Ladies of the Club%22 <b>http:.../ontology/author</b> ?author
3	http:.../resource/%22... <b>http:.../ontology/author</b> ?author
4	http:.../resource/%22...Burglar <b>http:.../ontology/author</b>

```
SELECT ?author ?book
WHERE {
  ?book rdf:label dbo:Book;
  dbo:author ?author }
LIMIT 100
```



# What happens if clients collaborate?



**United Federation of Data Consumers**

# What if clients collaborate?

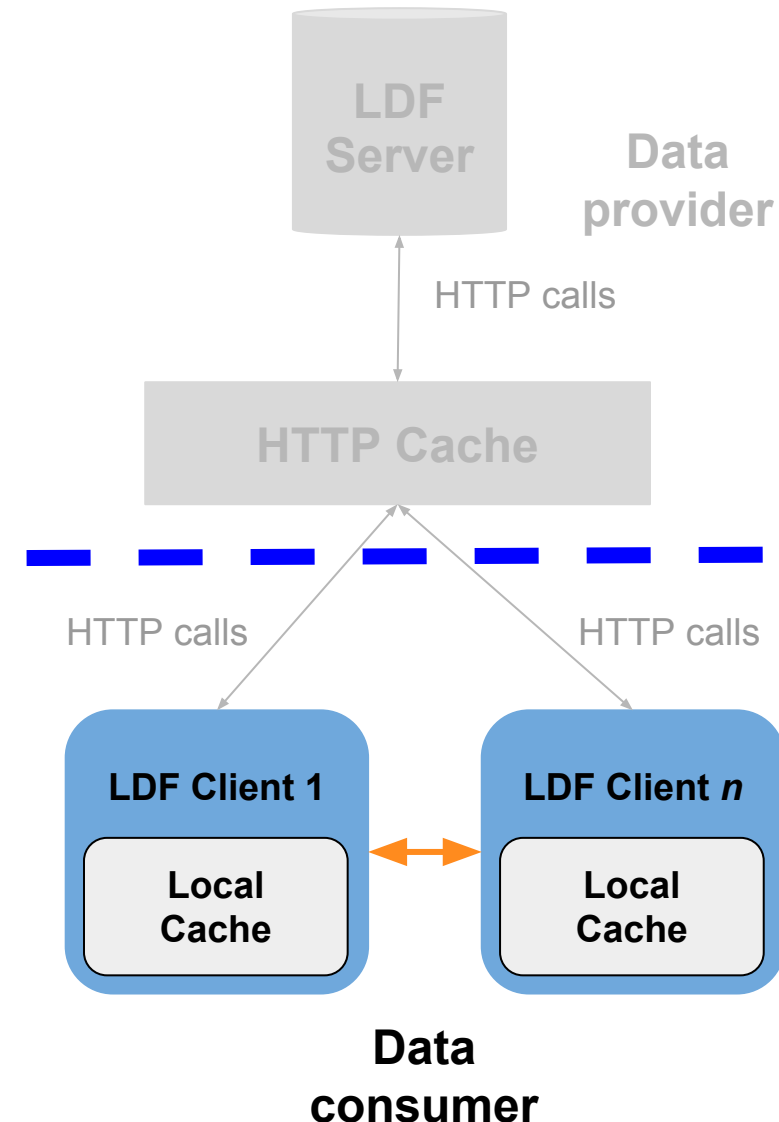
A local cache of a client can be **shared** with other clients

- Reduce the load on the server

## Challenges:

- Network with 1 million of clients

**How a client can find a TPF quickly in another client cache?**






# Related Works

**DHT:** distributes the cache among participants [2]

- + Lookup(TPF): finds TPF if it exists!
- 1 query  $\rightarrow$  20,000 calls  $\rightarrow$  20,000  $\log(n)$  hops,  $n$ : number of peers




[2] M. El Dick and al. Flower-cdn: A hybrid p2p overlay for efficient query processing in cdn. *In Proceedings of the 12th International Conference on Extending Database Technology: Advances in Database Technology*, 2009. ACM.

# Related Works

**Behavioral Cache:** connects a fixed size of similar nodes <sup>[3]</sup>

- + Zero-hop lookup latency
- Not sure to find a TPF
- ~ Experimented with browsing histories

What if we experiment behavioral cache with queries?



# CyCLaDEs Approach

**Assumption:** Clients which performed **similar queries** in the **past** will likely perform **similar queries** in the **future**

# CyCLaDEs Approach

**Assumption:** Clients which performed **similar queries** in the **past** will likely perform **similar queries** in the **future**

## **Approach:**

- Each client builds a **fixed-size** of **best similar neighbors** with **zero-hop latency**

# CyCLaDEs Approach

**Assumption:** Clients which performed **similar queries** in the **past** will likely perform **similar queries** in the **future**

## **Approach:**

- Each client builds a **fixed-size** of **best** similar **neighbors** with **zero-hop latency**
- For each call 1) check **local** cache 2) check **neighbors'** cache in parallel 3) go to the server

# CyCLaDEs Approach

**Assumption:** Clients which performed **similar queries** in the **past** will likely perform **similar queries** in the **future**

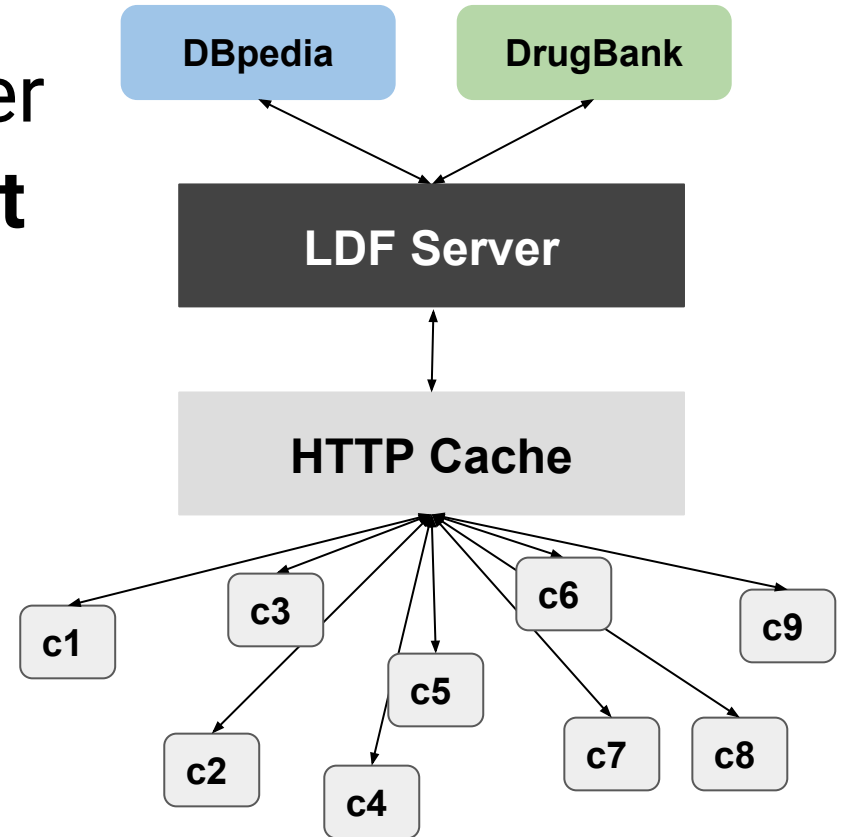
## **Approach:**

- Each client builds a **fixed-size** of **best** similar **neighbors** with **zero-hop latency**
- For each call 1) check **local** cache 2) check **neighbors'** cache in parallel 3) go to the server

**How many neighbors cache hits can we get?**

# LDF: Approach Overview

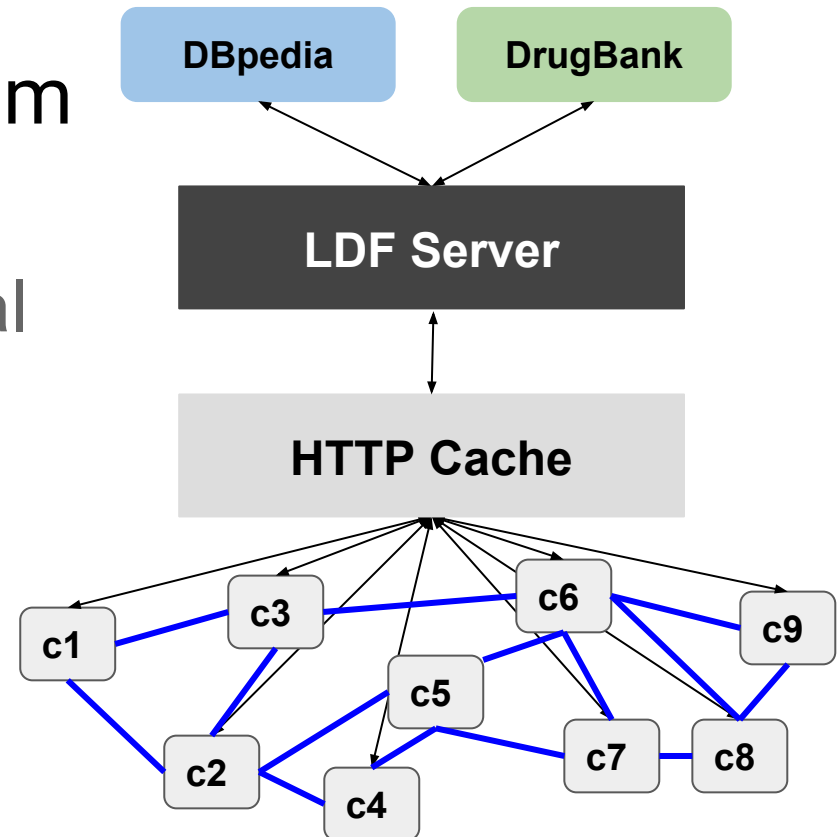
Many clients access LDF server concurrently, but **clients do not collaborate**.



# CyCLaDEs: Approach Overview

Connect nodes through Random Peer Sampling (RPS):

- Each node maintains a partial view on the entire network
- The view contains a random subset of network nodes



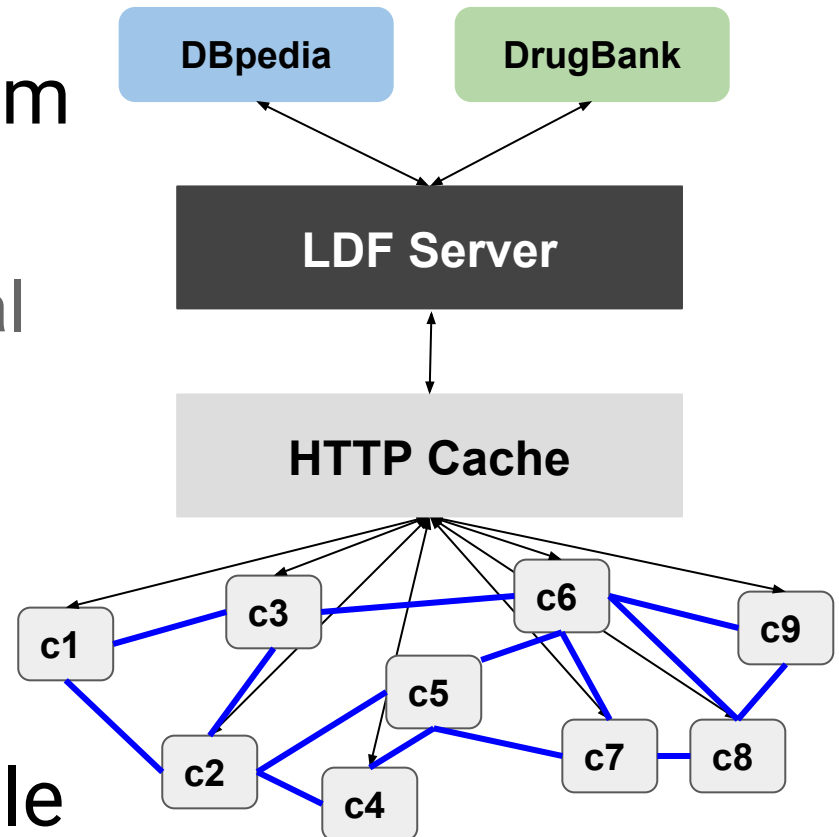


# CyCLaDEs: Approach Overview

Connect nodes through Random Peer Sampling (RPS):

- Each node maintains a partial view on the entire network
- The view contains a random subset of network nodes

Shuffle phases: **renew**  
**periodically** neighbors to handle  
churn and avoid network partition.



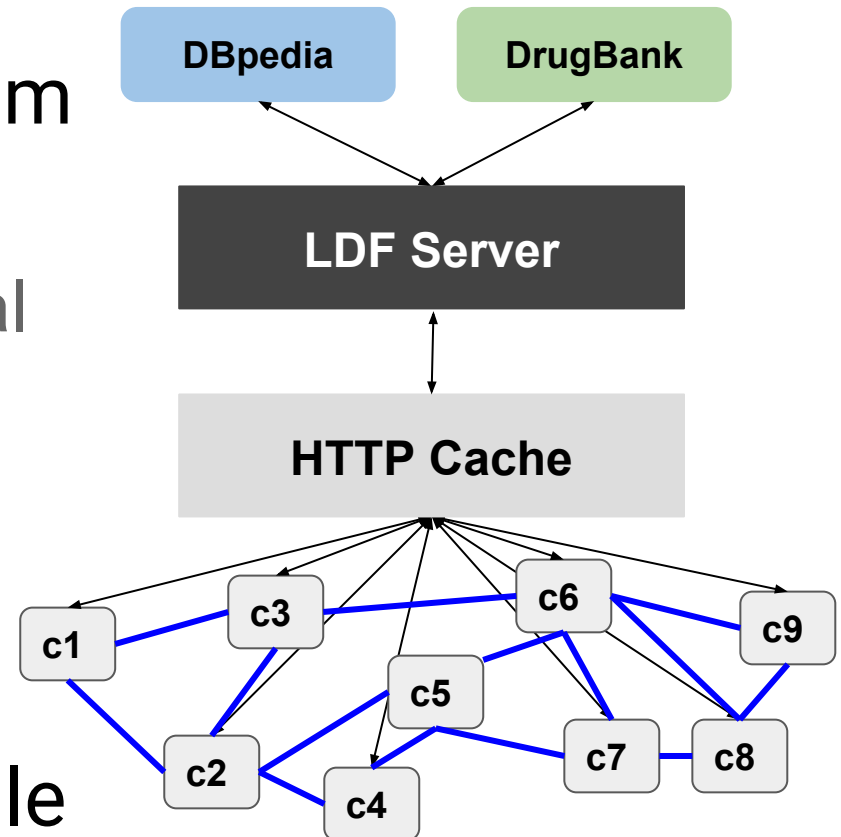
# CyCLaDEs: Approach Overview

Connect nodes through Random Peer Sampling (RPS):

- Each node maintains a partial view on the entire network
- The view contains a random subset of network nodes

Shuffle phases: **renew periodically** neighbors to handle churn and avoid network partition.


→ We use Cyclon<sup>[4]</sup> for RPS.



[4] S. Voulgaris and al. Cyclon: Inexpensive membership management for unstructured p2p overlays. *Journal of Network and Systems Management*, 2005.

# CyCLaDEs: Approach Overview

RPS overlay network ensures connectivity among **all** clients.




# CyCLaDEs: Approach Overview

RPS overlay network ensures connectivity among **all** clients.

C6 is connected to C3:

- C6 → DBpedia
- C3 → DrugBank
- C6 is not **similar** to C3




# CyCLaDEs: Approach Overview

RPS overlay network ensures connectivity among **all** clients.

C6 is connected to C3:

- C6 → DBpedia
- C3 → DrugBank
- C6 is not **similar** to C3


Need a second overlay to handle **similarity** as proposed in Gossple [5].



# CyCLaDEs: Approach Overview

## Cluster Overlay Network (CON):

- Each node has a **profile** based on the **history** of executed queries




# CyCLaDEs: Approach Overview

## Cluster Overlay Network (CON):

- Each node has a **profile** based on the **history** of executed queries

Shuffle phases allow to obtain better neighbors.

Each node ranks and selects **best neighbors** based on similarity of profiles.



# How to profile nodes?

Executing queries produces a **stream of TPF**, cache is a **window** on this **stream**.



**Profile = summary of the recent past =  
frequency of the  $k$  last recently used  
predicates**



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)

---

**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp = (s\ p\ o)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:      $Pr.update(tp.p, f_p + 1, t)$ 
6:   else
7:      $Pr \cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:      $Pr \setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

---

# Node profile - Algorithm

---

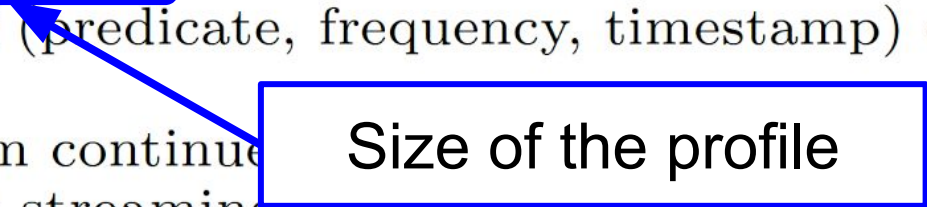
## Algorithm 1 ComputeProfile(s,w,t)

---

**Require:** : w: Window size s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues
3:   Receive the next streaming triple  $tp = (s, p, o)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:     Pr.update( $tp.p, f_p + 1, t$ )
6:   else
7:     Pr  $\cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:     Pr  $\setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)

---

**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp = (s, p, o)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:     Pr.update( $tp.p, f_p + 1, t$ )
6:   else
7:     Pr  $\cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:     Pr  $\setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

Stream of triples

# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)

---

**Require:** : w: Window size, s: Stream of triples, t: timestamp  
**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

1:  $Pr \rightarrow \emptyset$   
2: **while** data stream continues  
3:   Receive the next streaming triple  $tp = (s, p, o)$   
4:   **if**  $(tp.p, f_p, -) \in Pr$  **then**  
5:      $Pr.update(tp.p, f_p + 1, t)$   
6:   **else**  
7:      $Pr \cup (tp.p, 1, t)$   
8:   **if**  $|Pr| > w$  **then**  
9:      $Pr \setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$

Timestamp of node



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)


---

**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp = (s, p, t)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:     Pr.update( $tp.p, f_p + 1, t$ )
6:   else
7:     Pr  $\cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:     Pr  $\setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

For each predicate in the stream



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)

---

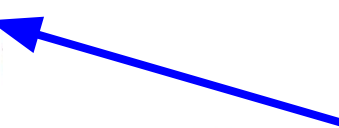
**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp = (s\ p\ o)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:      $Pr.update(tp.p, f_p + 1, t)$ 
6:   else
7:      $Pr \cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:      $Pr \setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

If predicate exist:

- 1) Increment frequency by 1
- 2) Update timestamp



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)

---


**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp = (s\ p\ o)$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:      $Pr.update(tp.p, f_p + 1, t)$ 
6:   else
7:      $Pr \cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:      $Pr \setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

If predicate does not exist:

- 1) Insert predicate in profile
  - a) With frequency = 1
  - b) Current timestamp



# Node profile - Algorithm

---

## Algorithm 1 ComputeProfile(s,w,t)


---

**Require:** : w: Window size, s: Stream of triples, t: timestamp

**Ensure:** : Pr : set of (predicate, frequency, timestamp) of size w

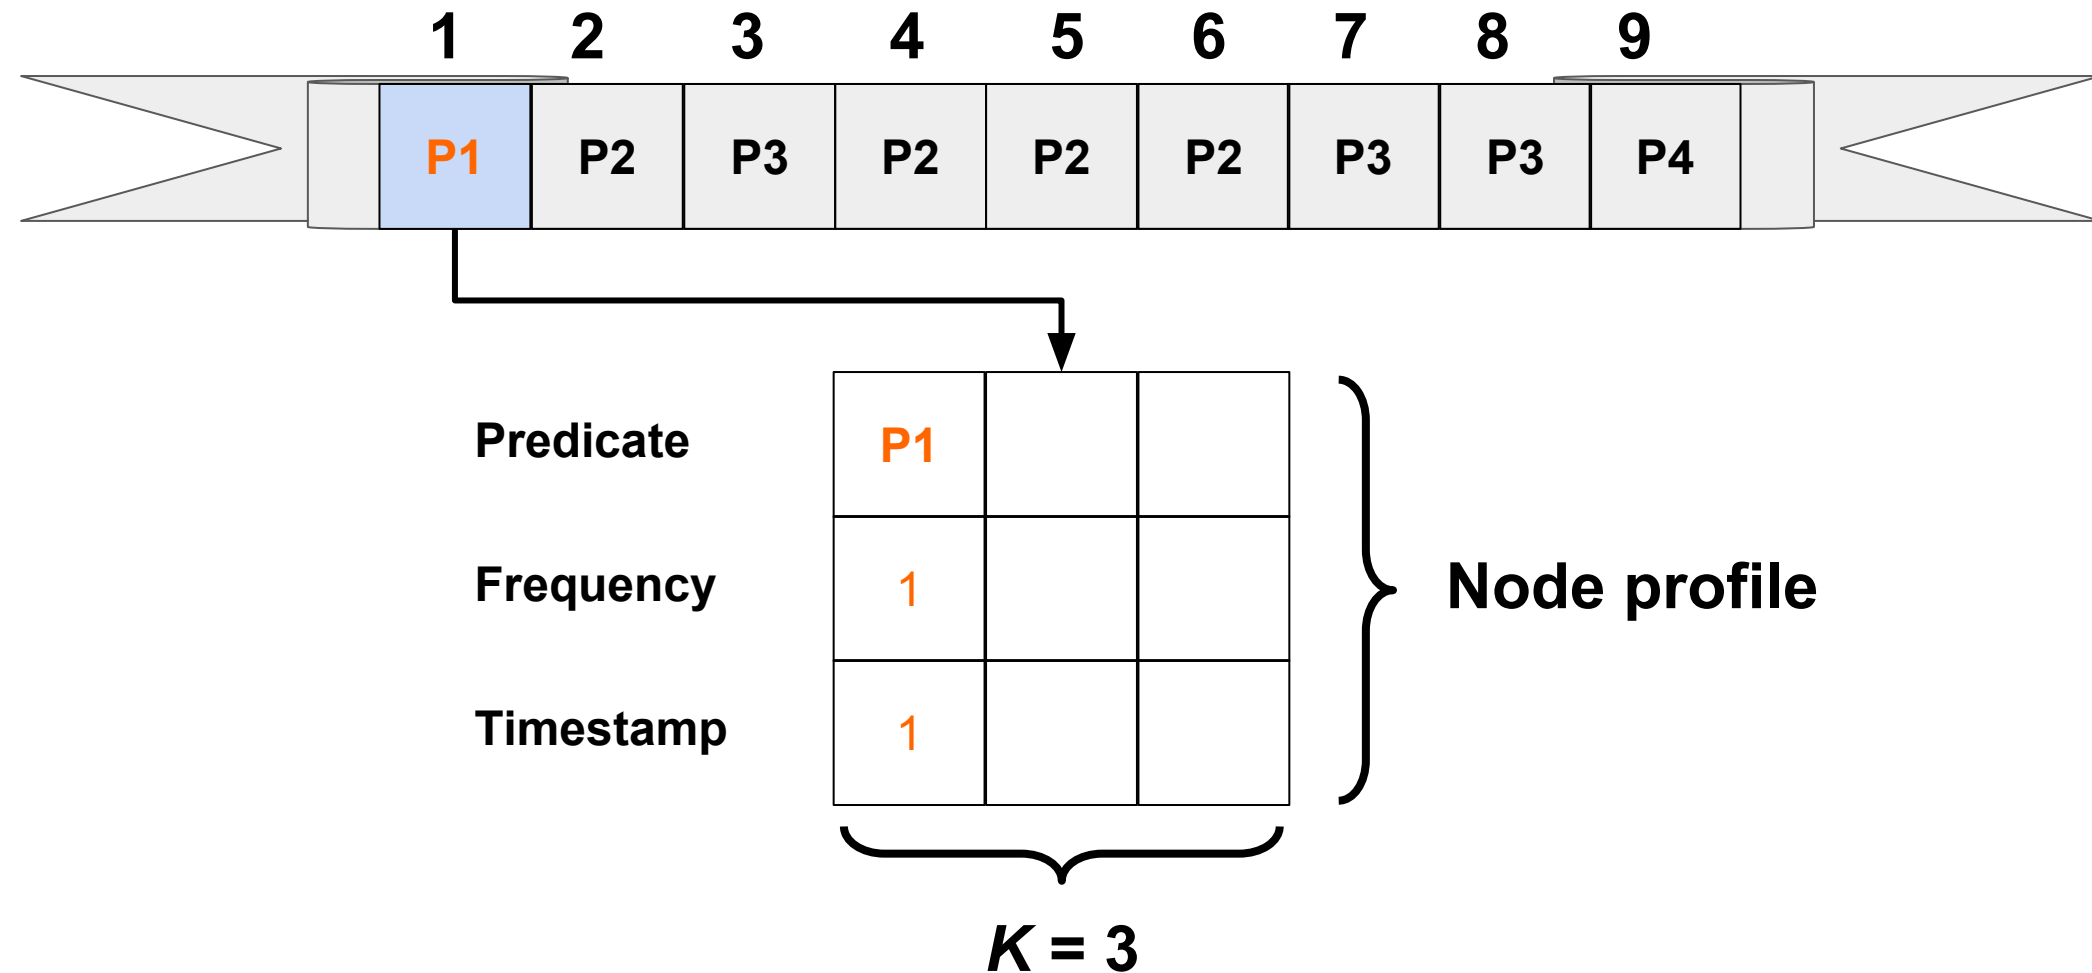
```
1: Pr  $\rightarrow \emptyset$ 
2: while data stream continues do
3:   Receive the next streaming triple  $tp =$ 
4:   if  $(tp.p, f_p, -) \in Pr$  then
5:      $Pr.update(tp.p, f_p + 1, t)$ 
6:   else
7:      $Pr \cup (tp.p, 1, t)$ 
8:   if  $|Pr| > w$  then
9:      $Pr \setminus (p_1, f_{p_1}, t_1) : (p_1, f_{p_1}, t_1) \in Pr \wedge \nexists (p_2, f_{p_2}, t_2) \in Pr : t_2 < t_1$ 
```

If  $|Pr| > w$ :  
Remove the oldest entry

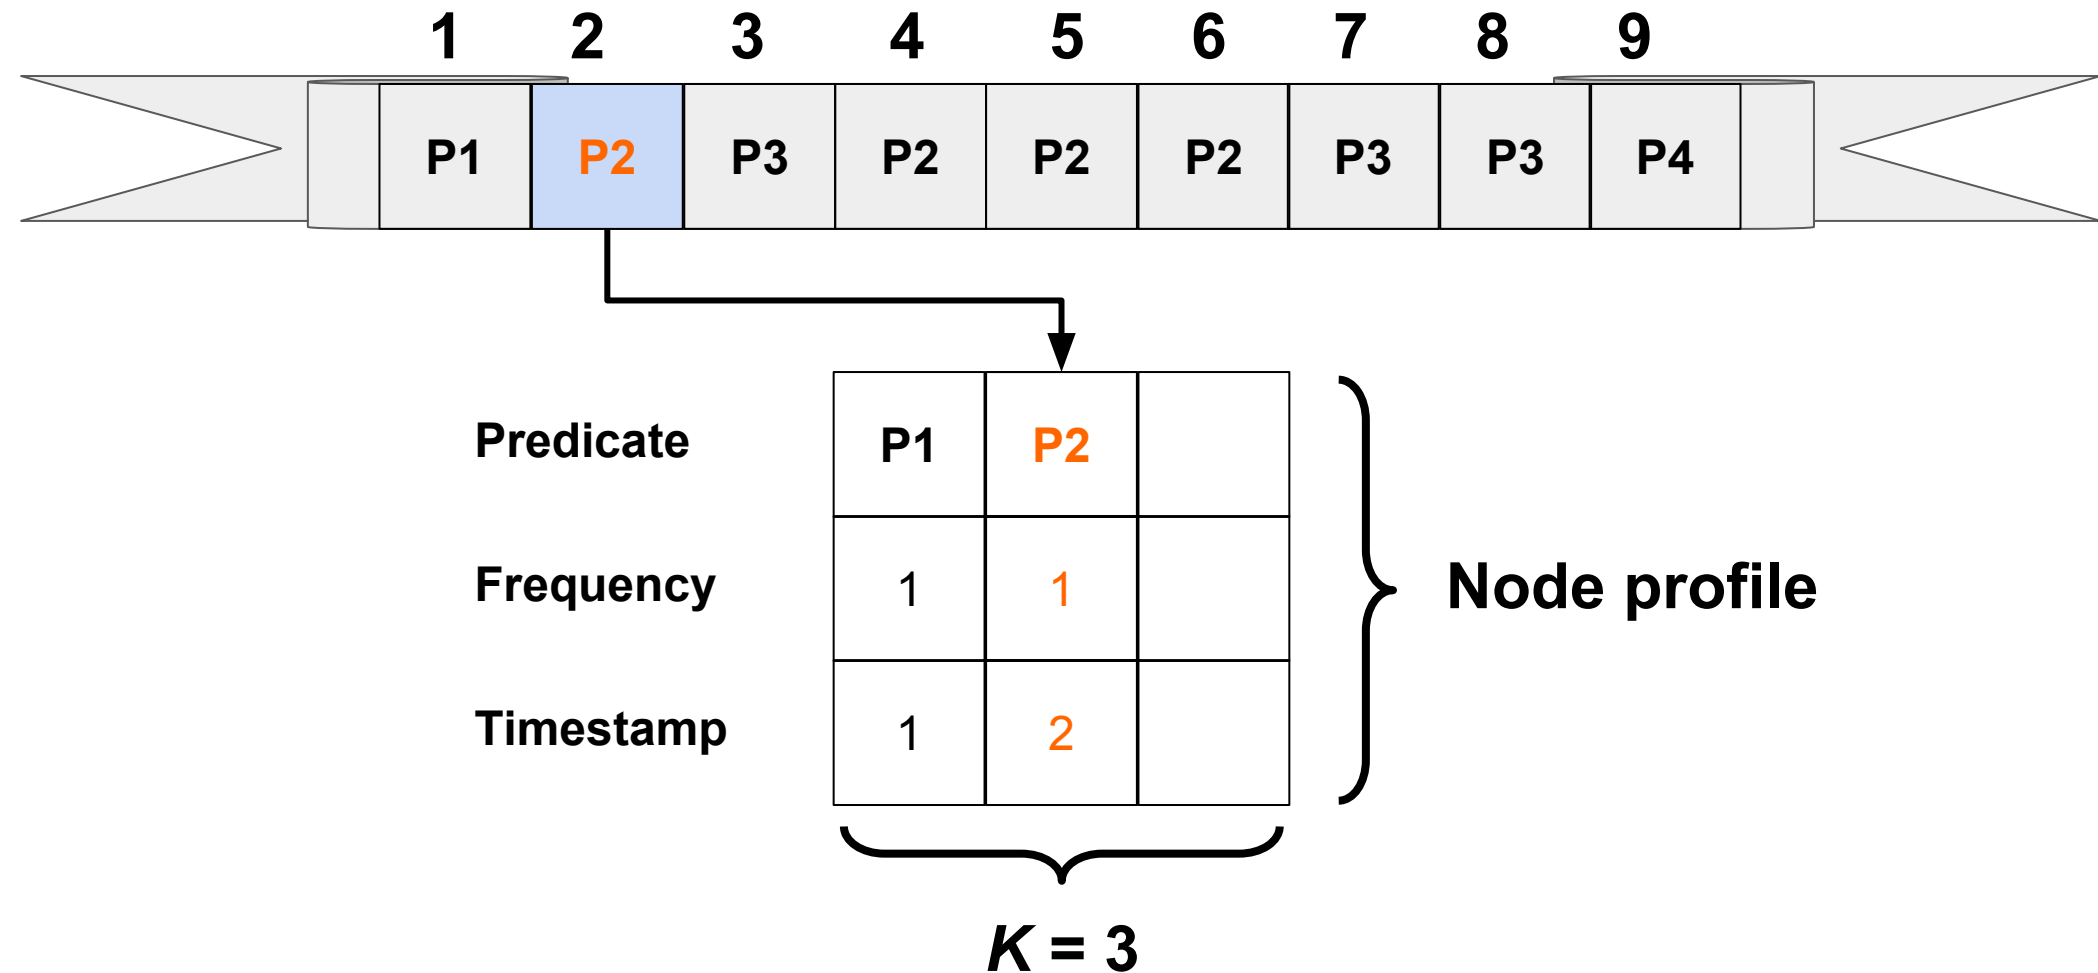




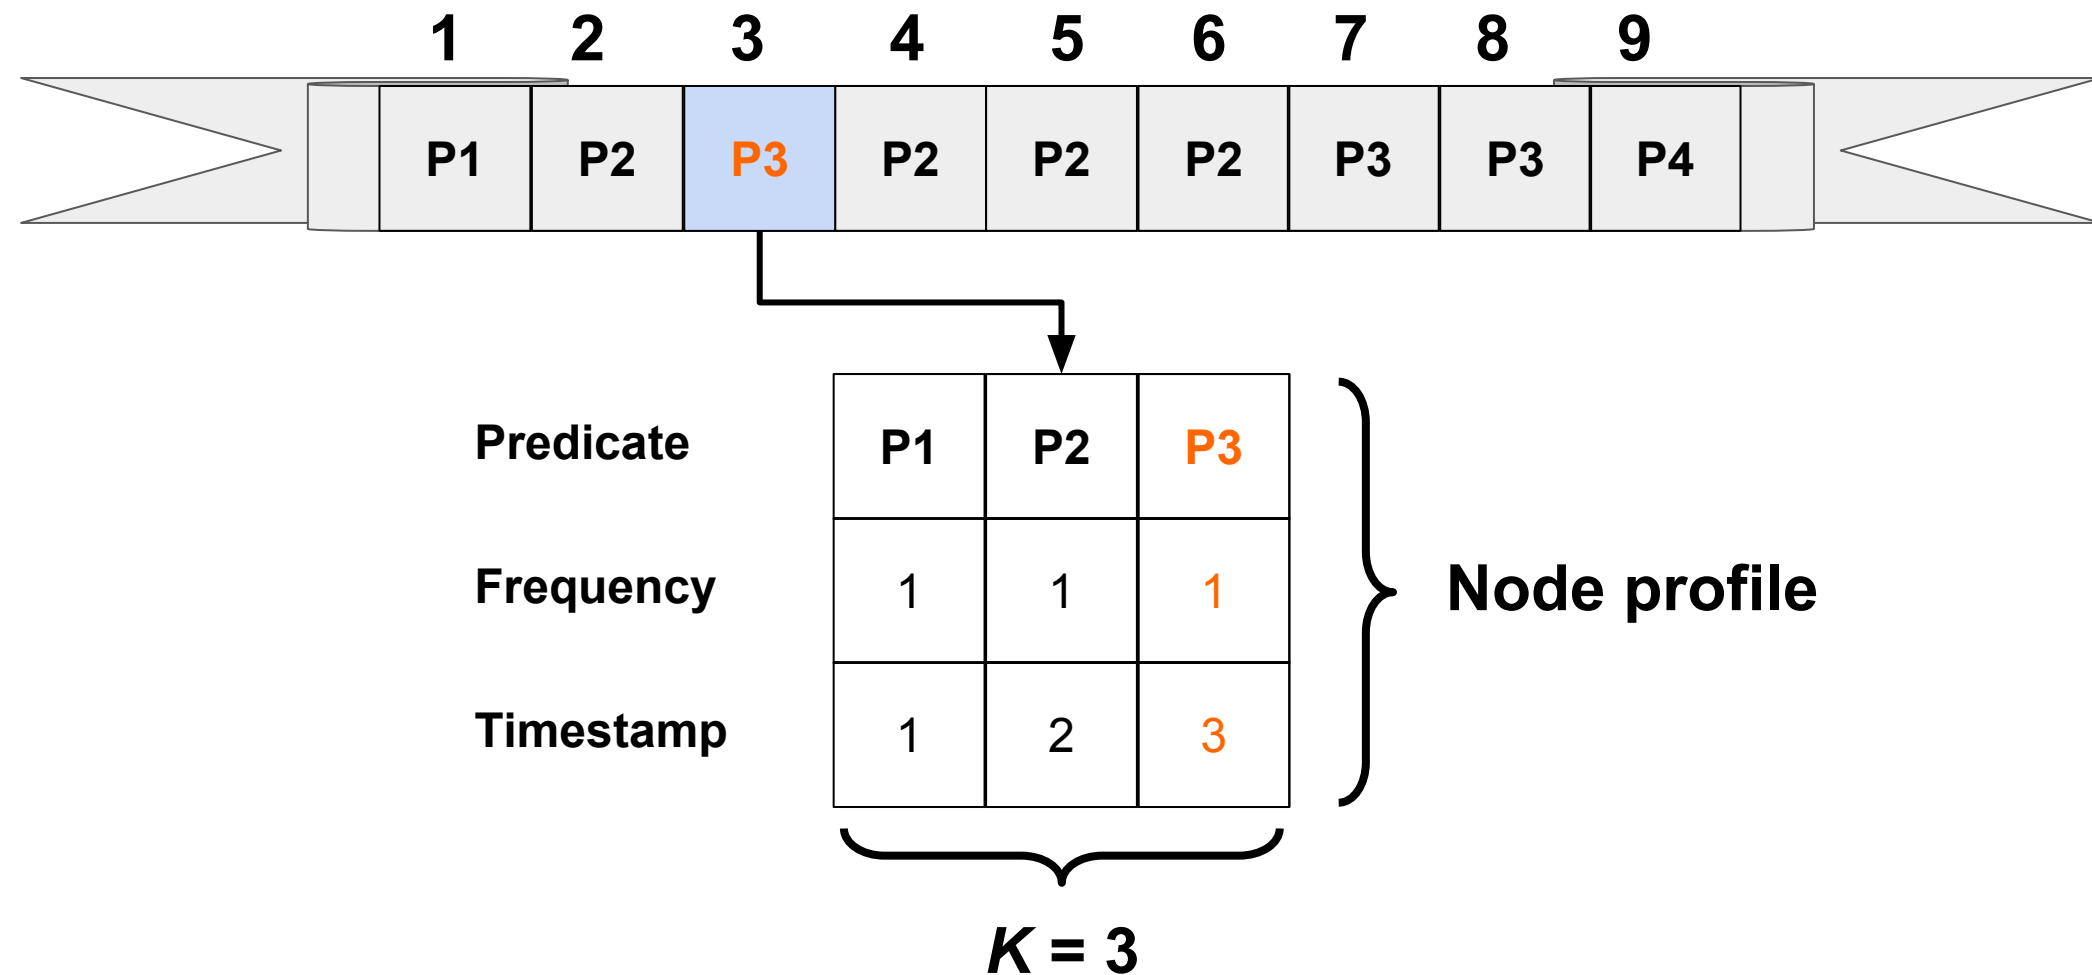
# Computing Profile



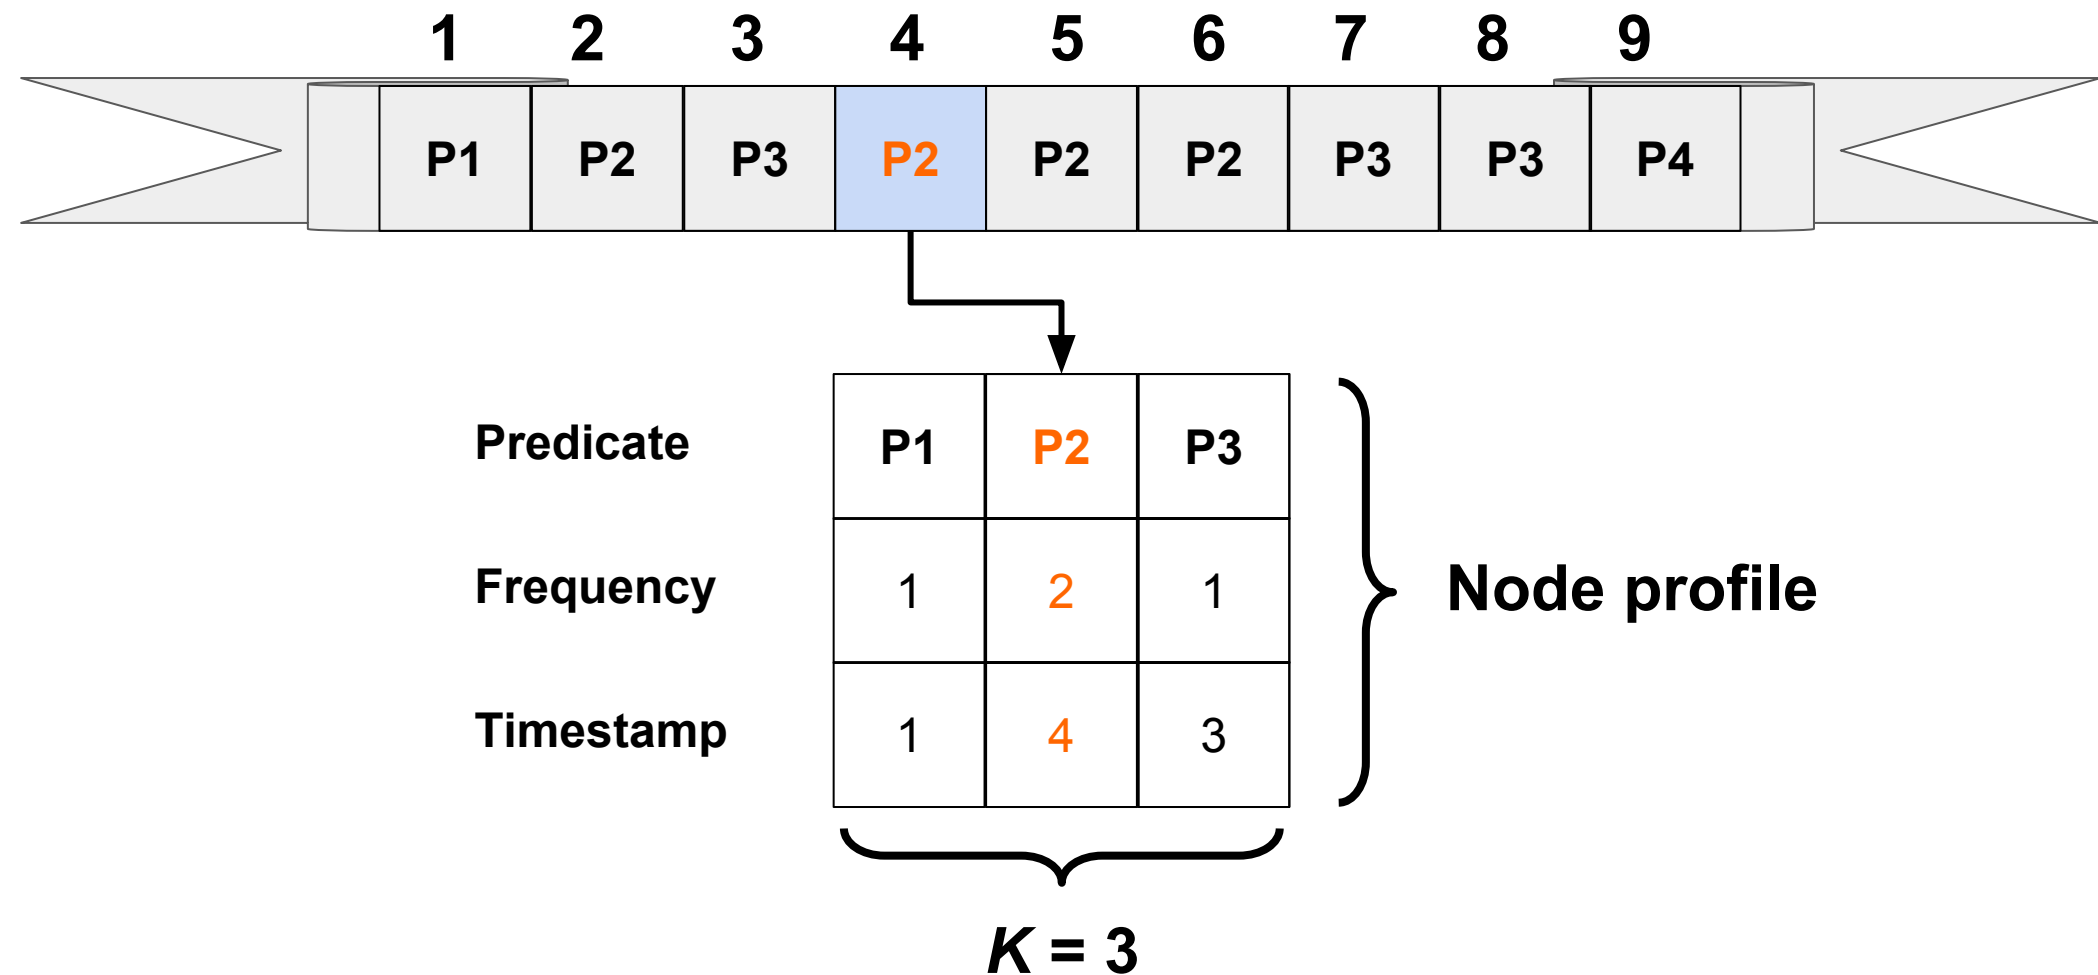
# Computing Profile



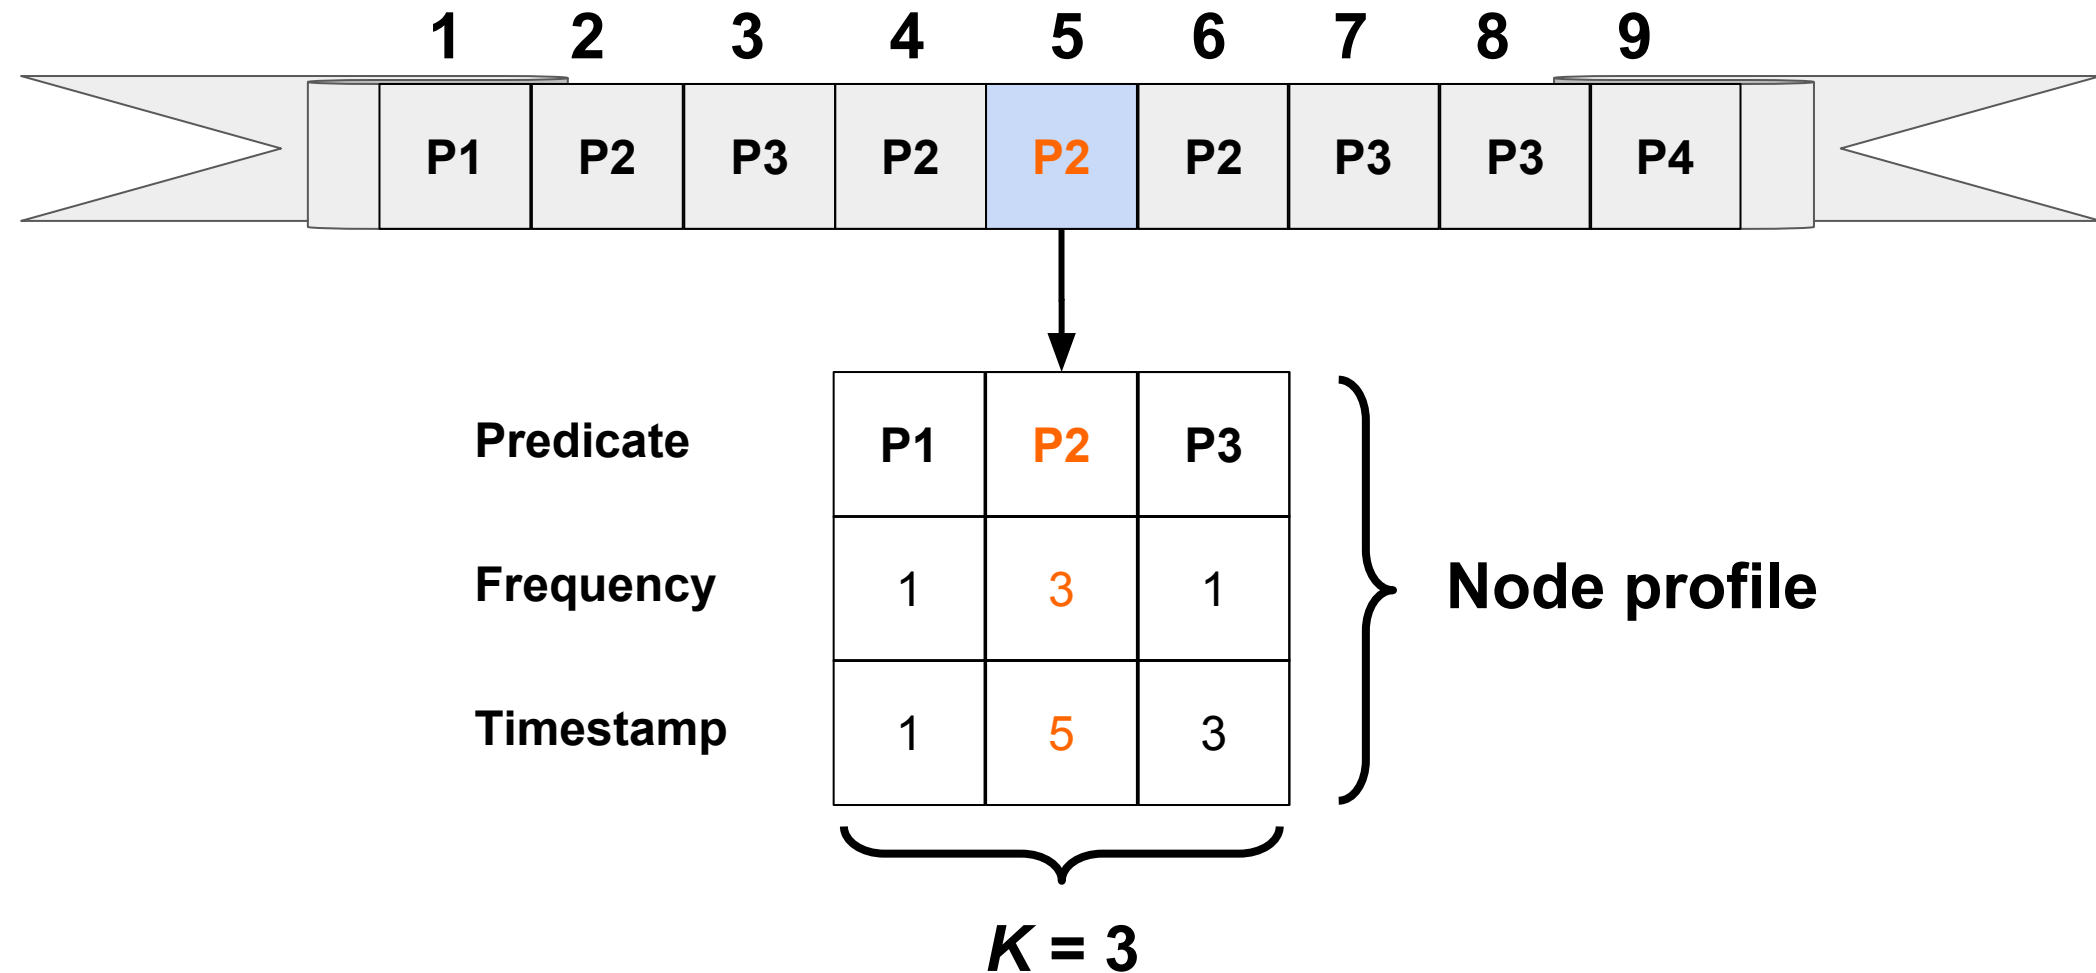
# Computing Profile



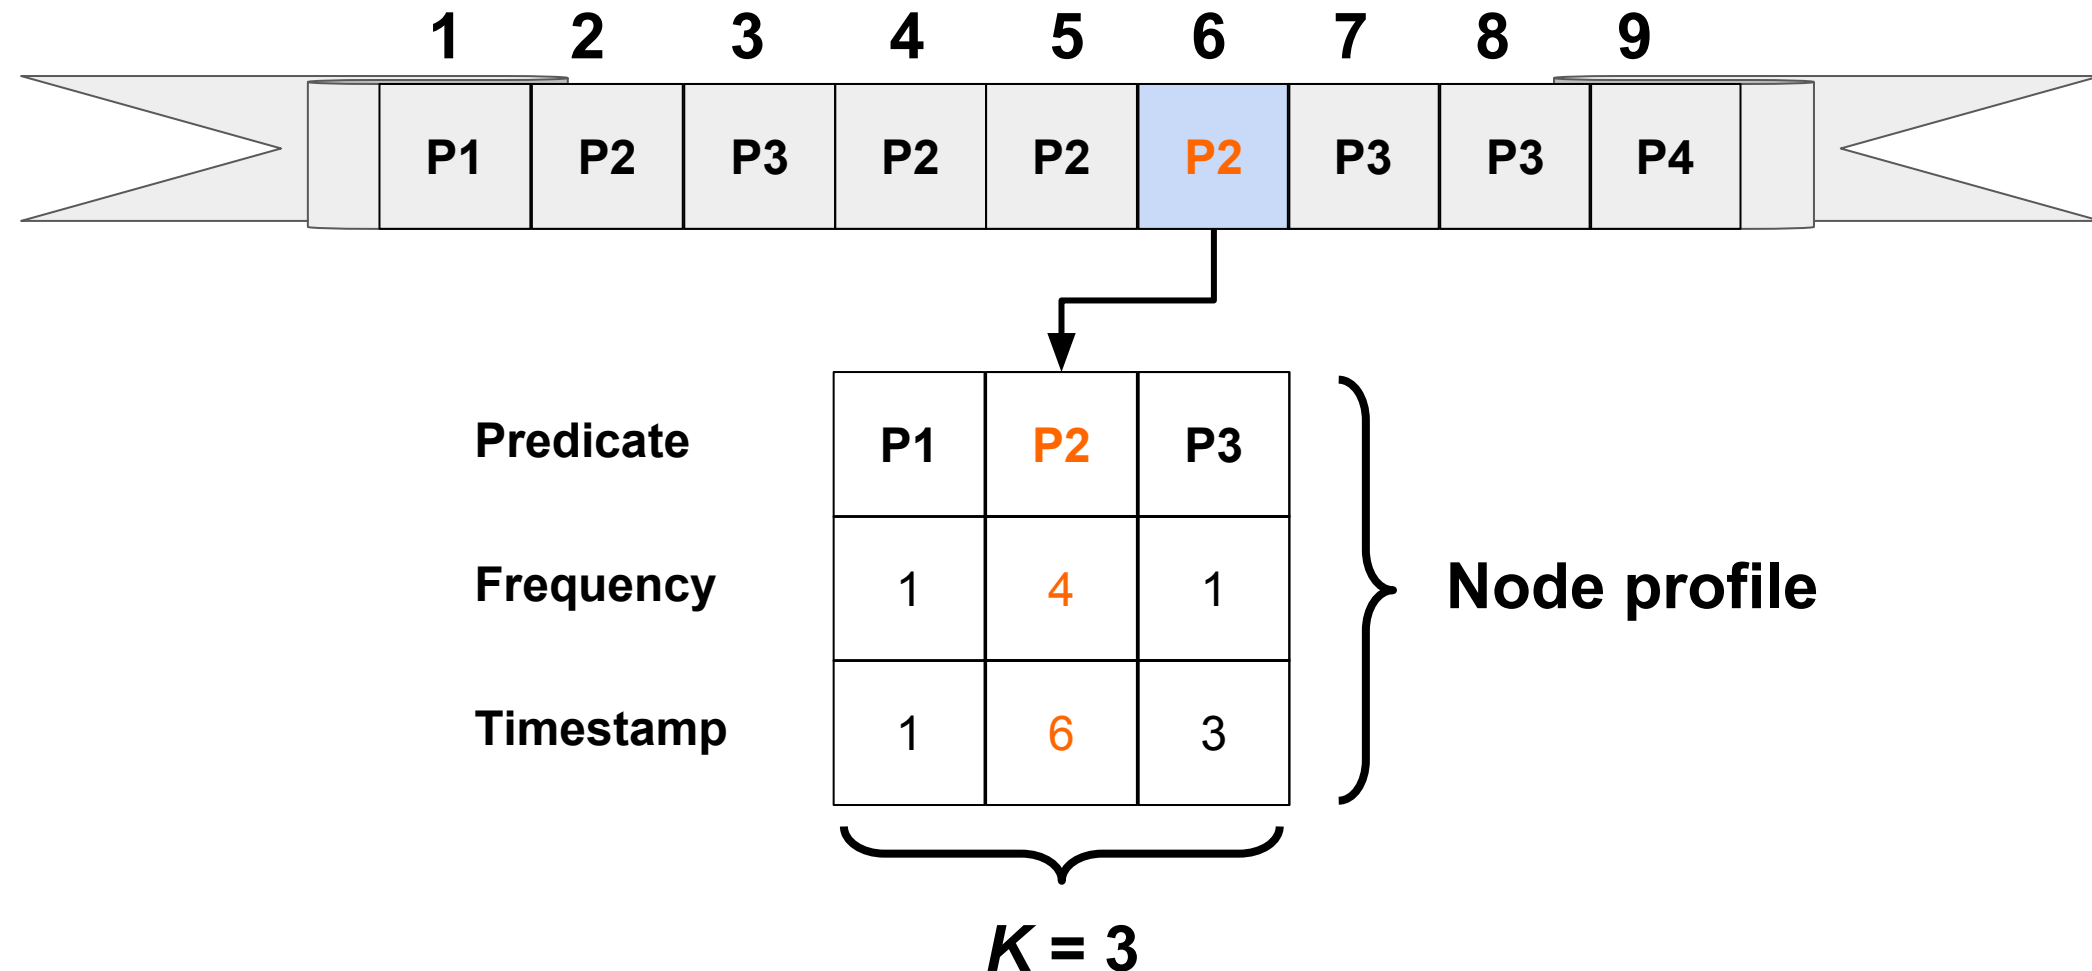
# Computing Profile



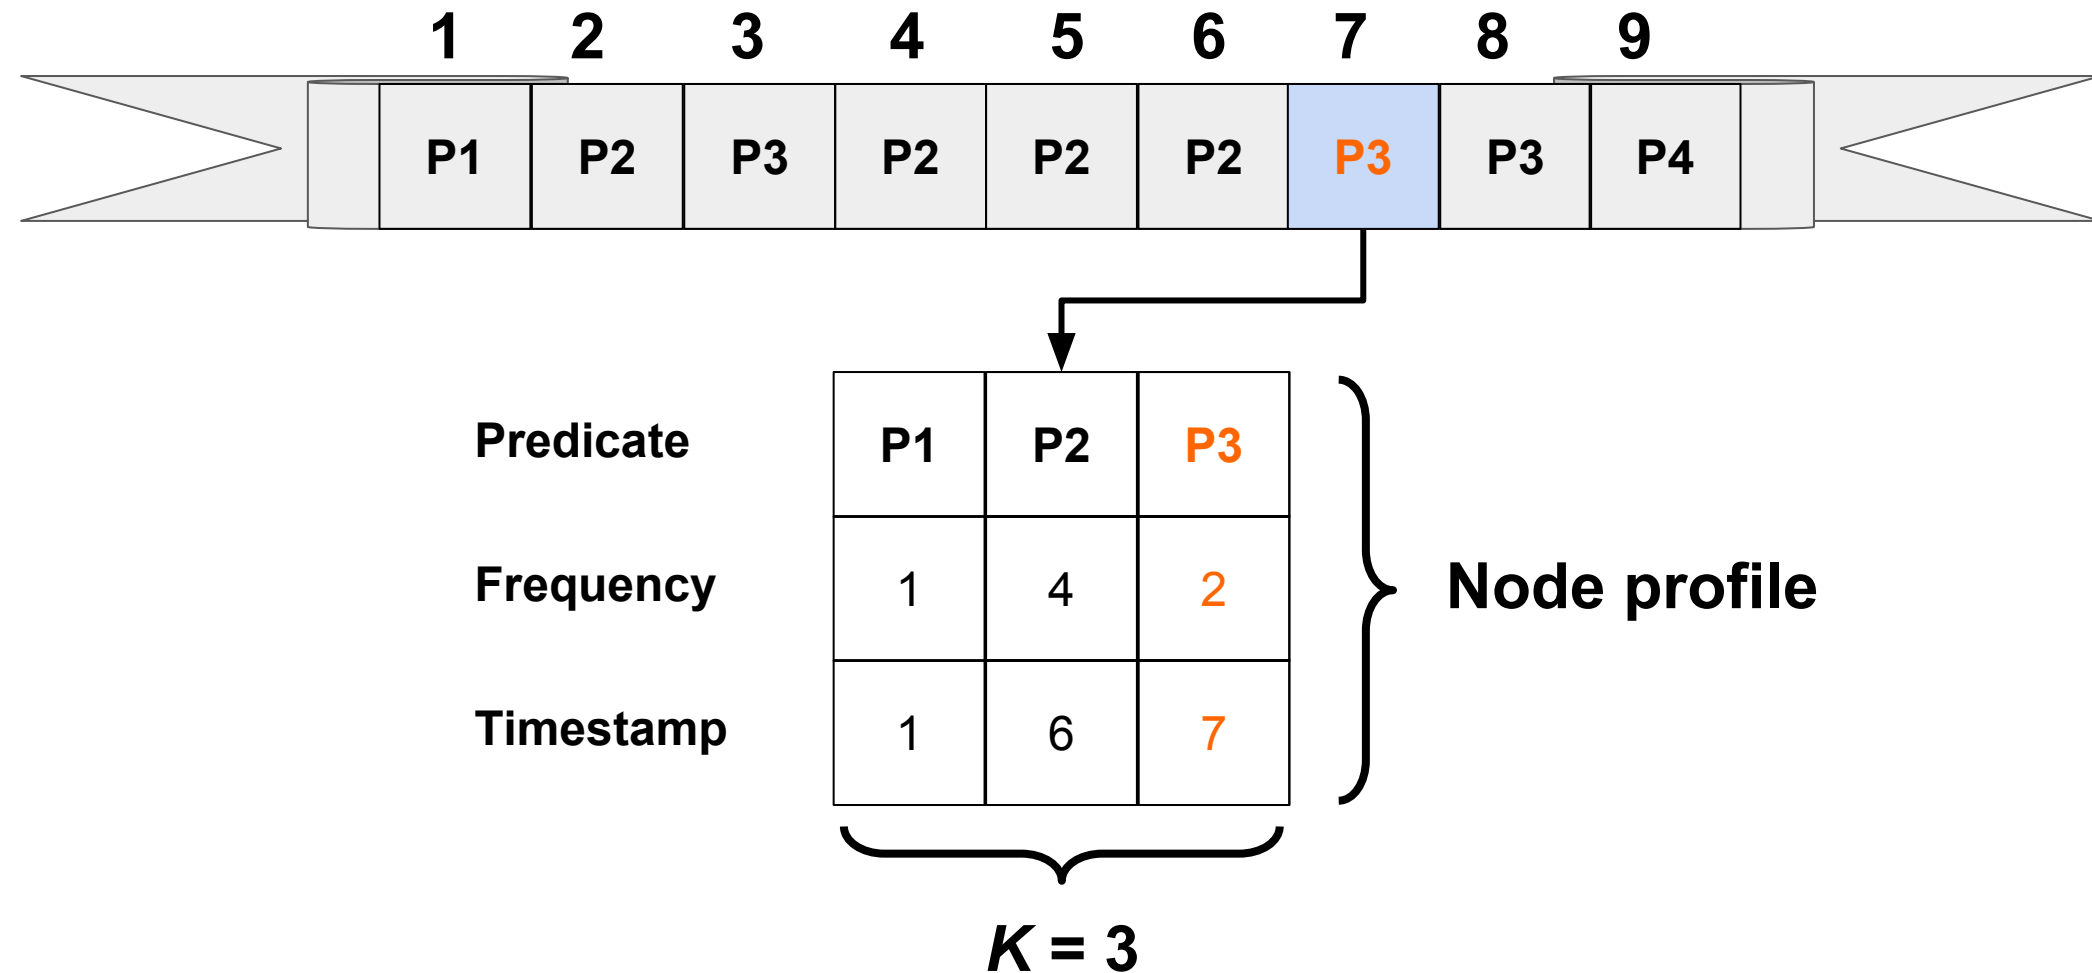
# Computing Profile



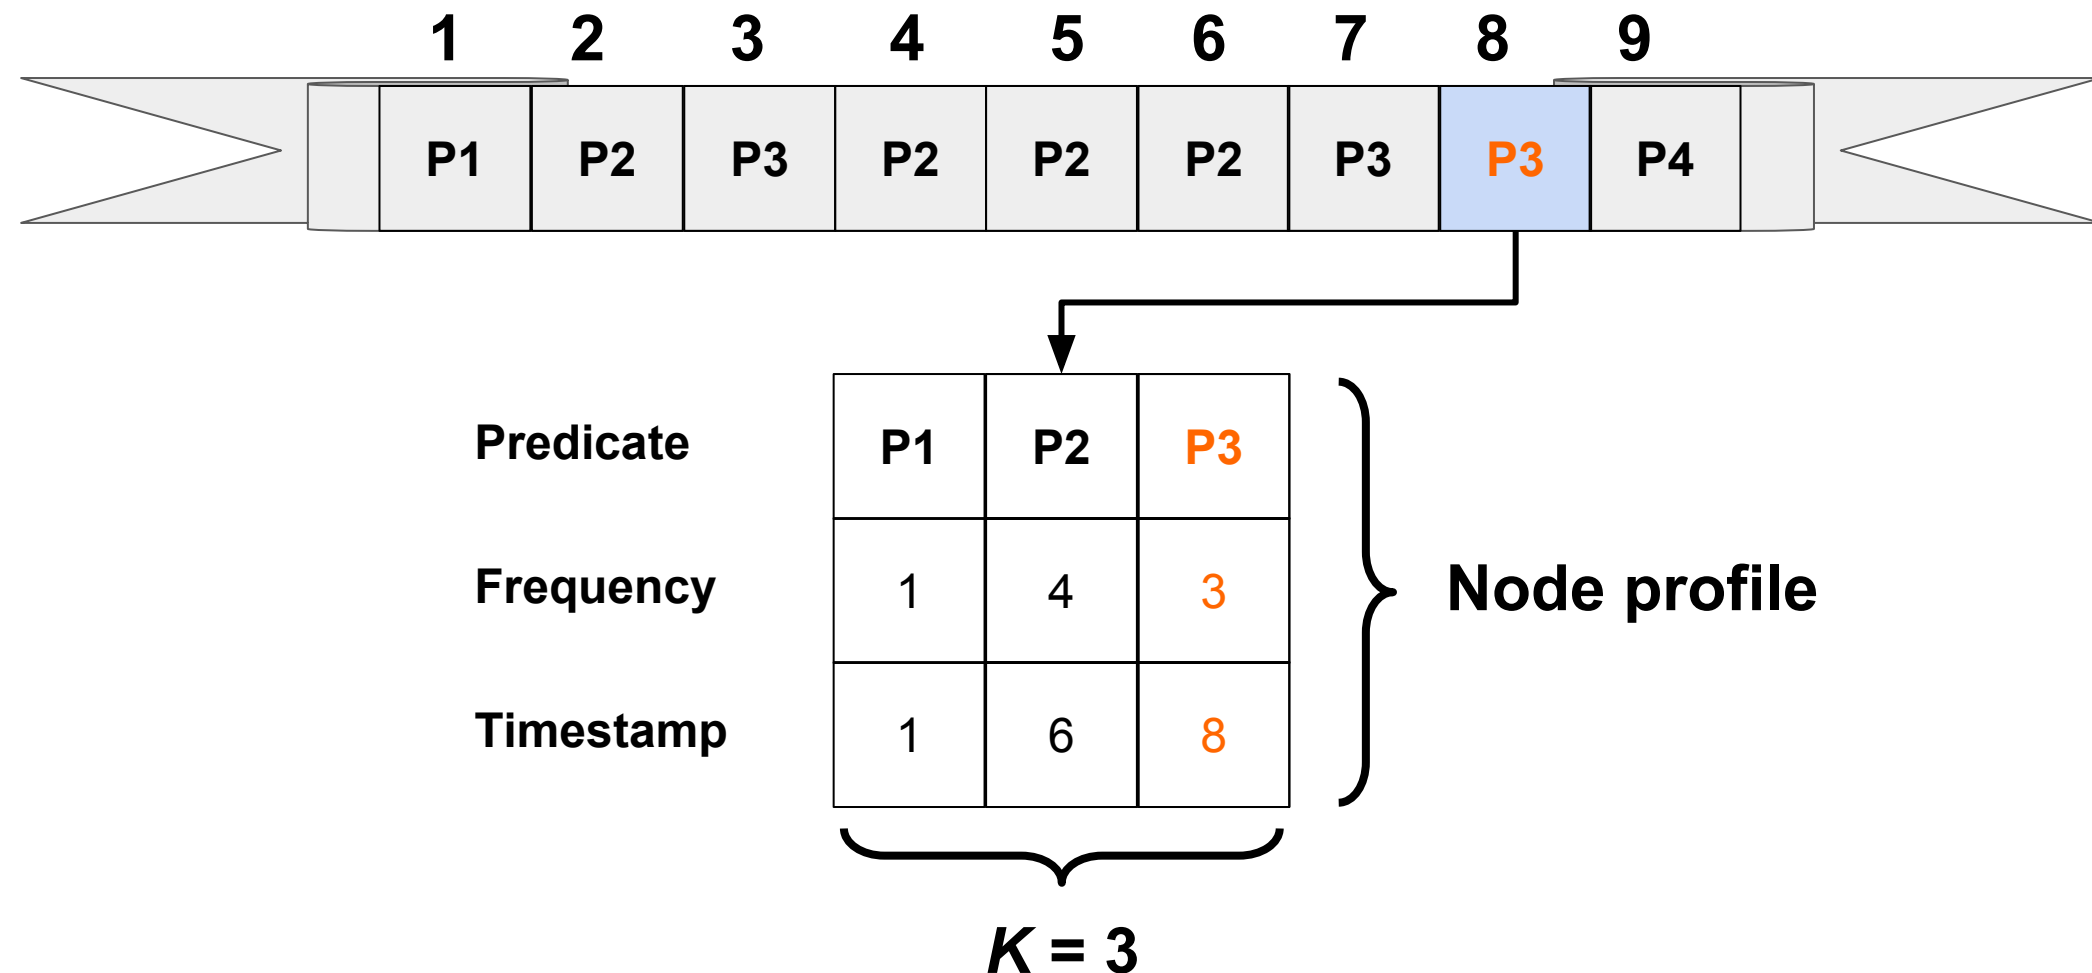
# Computing Profile



# Computing Profile

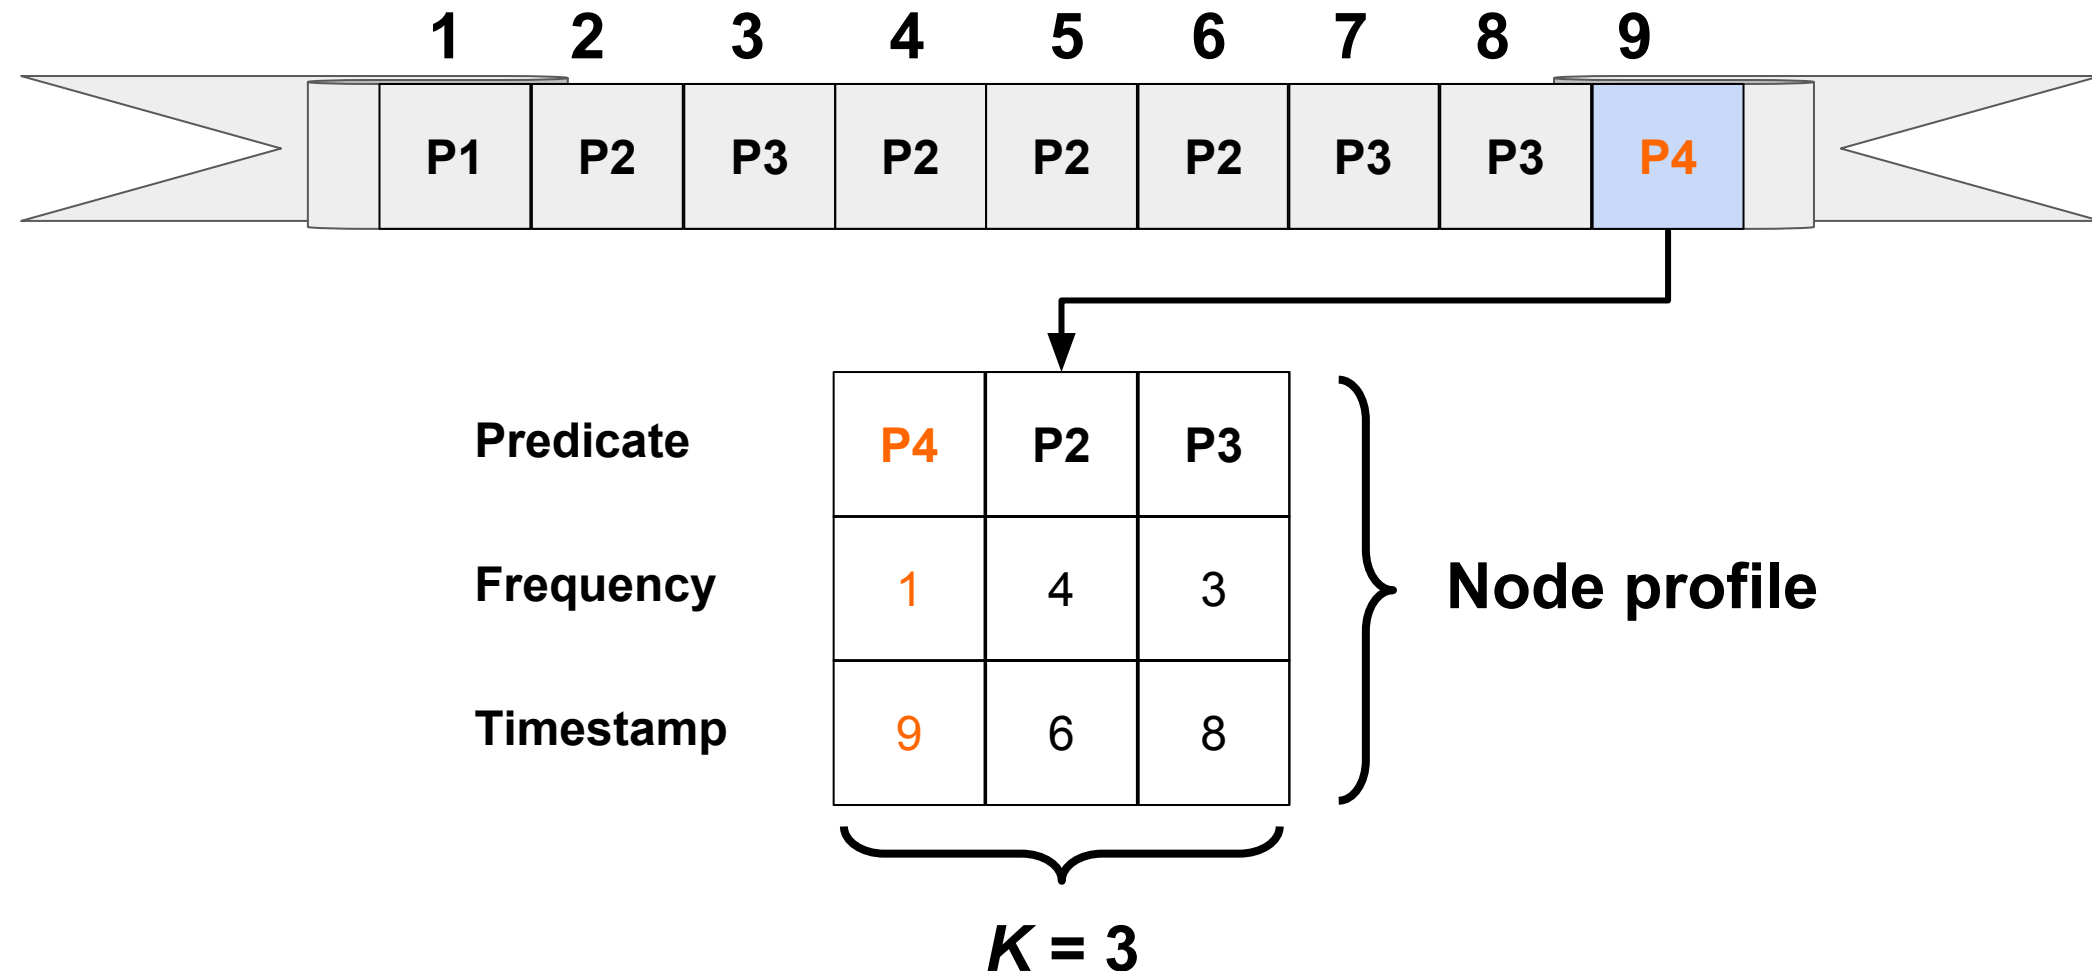


# Computing Profile





# Computing Profile



# How to compare neighbors?

During ranking **nodes** are **compared** thank to their **profile**.

Nodes are compared with the generalized **Jaccard similarity coefficient**.


$$J(x, y) = \frac{\sum_i \min(x_i, y_i)}{\sum_i \max(x_i, y_i)}$$

	P1	P2	P3
C5	10	6	40
C9	10	20	5
C8	10	2	2

- $J(C5, C9) = (10+6+5)/(10+20+40) = \mathbf{0.3}$
- $J(C5, C8) = 14/56 = \mathbf{0.25}$

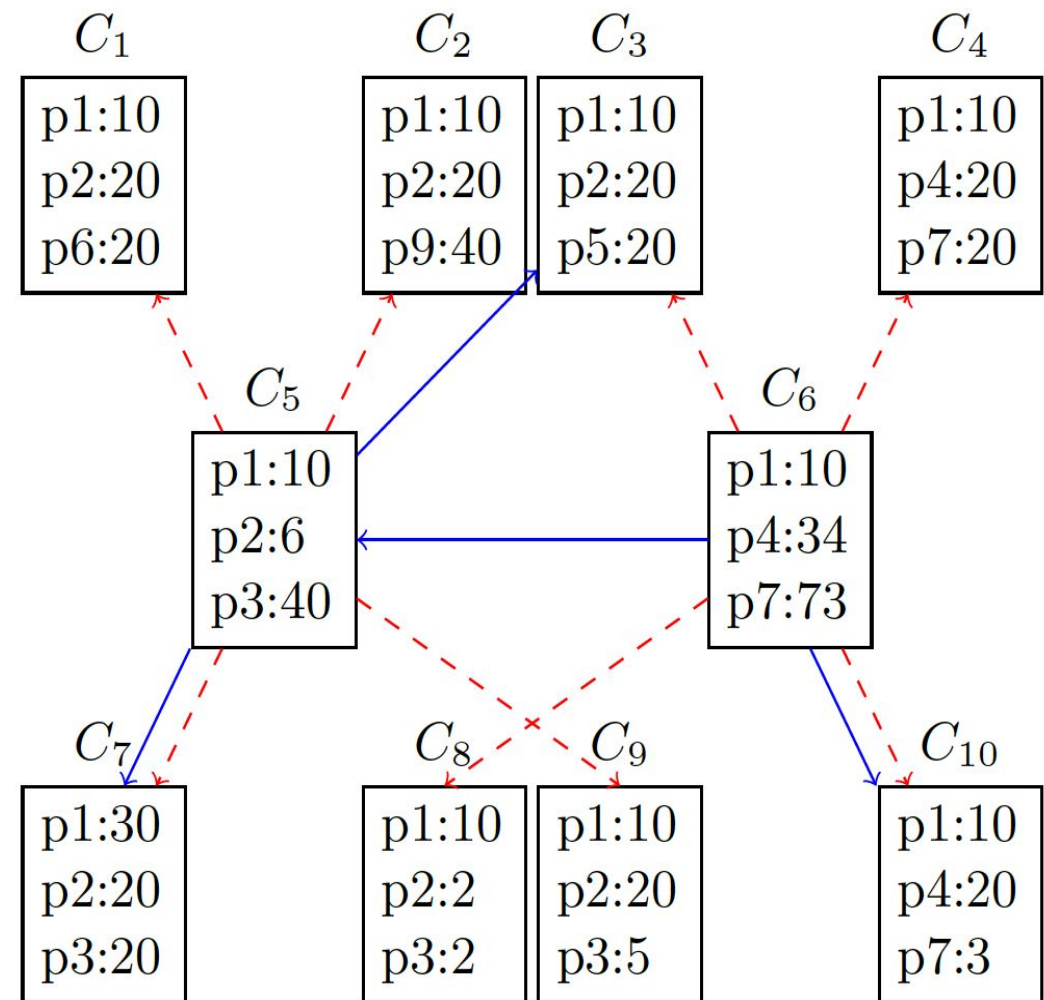
# C5 shuffling with C6 :

#RPS=2, #CON=4



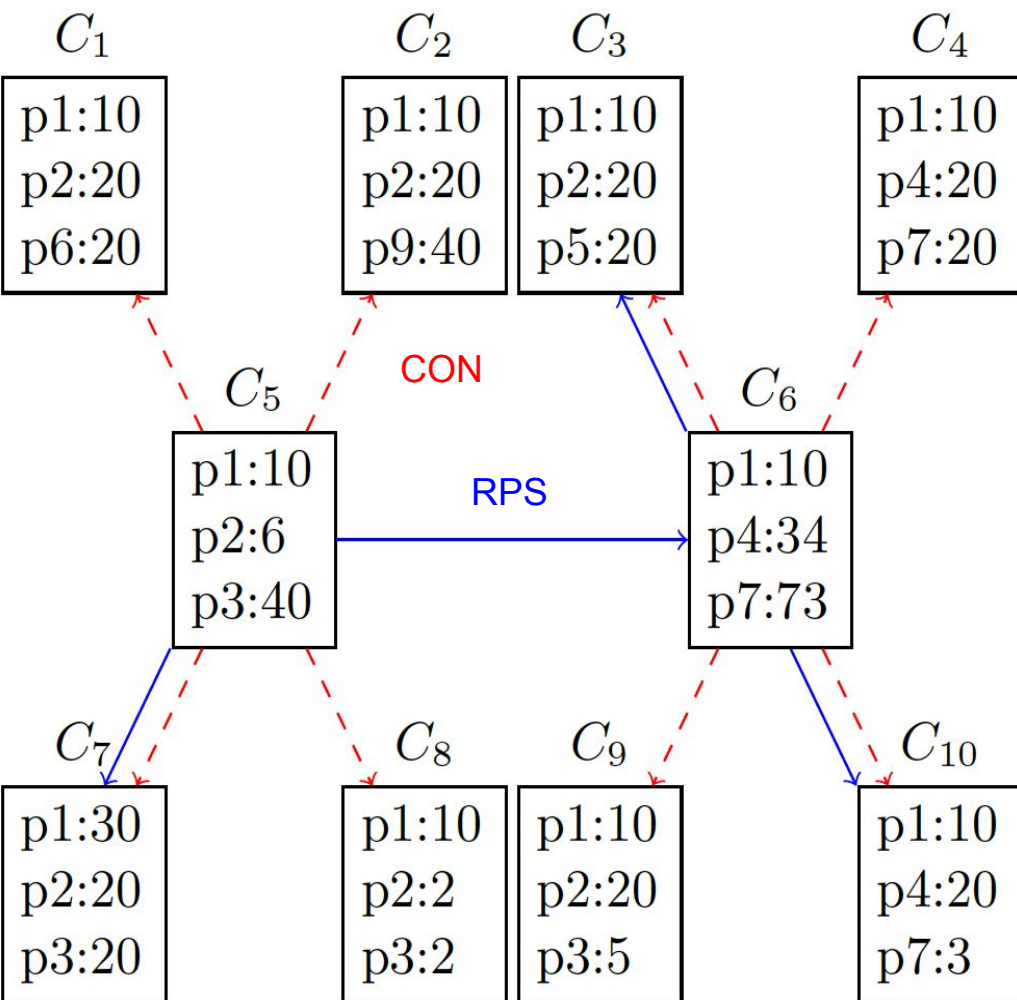
# After Shuffling:

RPS and CON updated




For C5: C9 is more similar than C8  
For C6: C8 is more similar than C9

# C5 shuffling with C6 : #RPS=2, #CON=4



# After Shuffling: RPS and CON updated



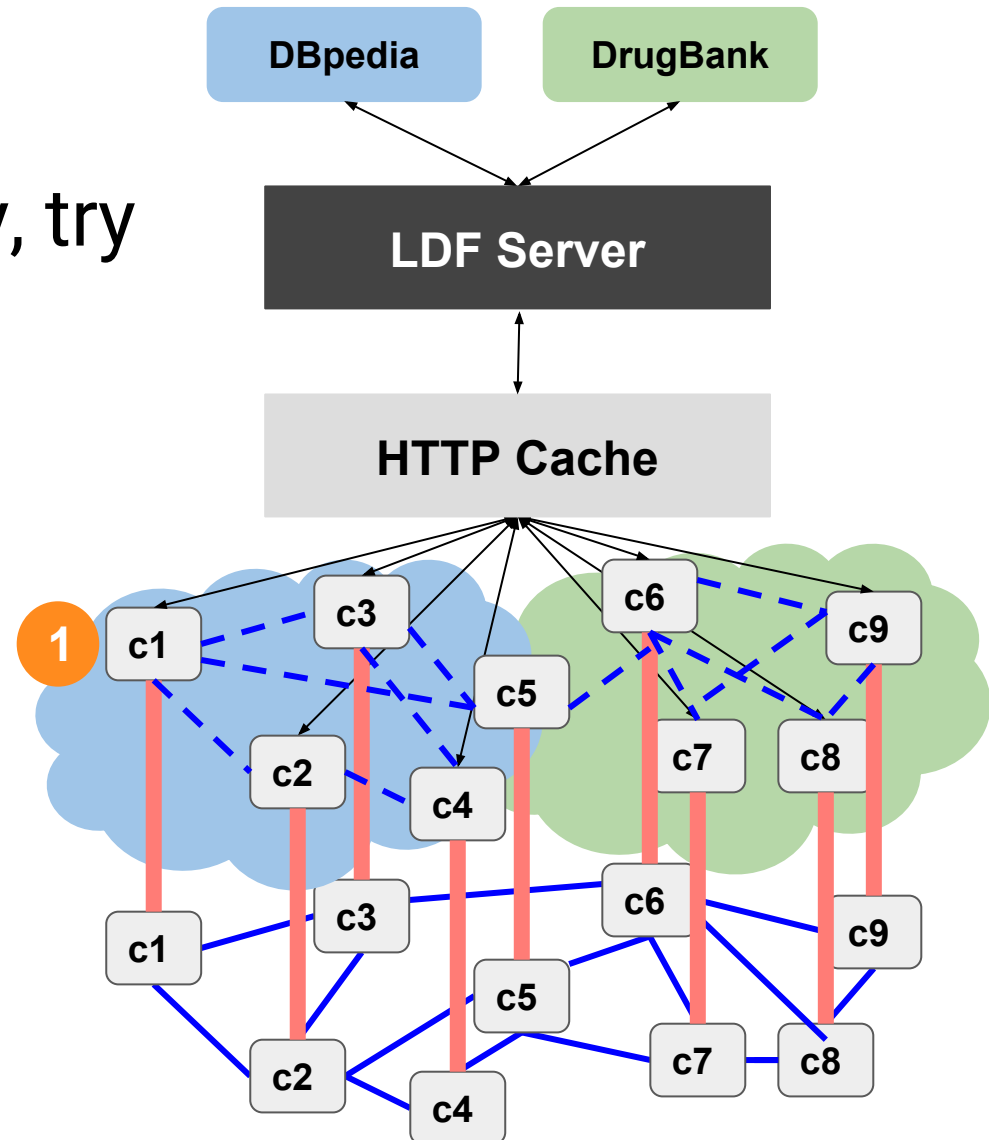
For  $C_5$ :  $C_9$  is more similar than  $C_8$   
For  $C_6$ :  $C_8$  is more similar than  $C_9$

# Queries with CyCLaDEs

C1 executes query Q1.

For each triple pattern query, try to resolve pattern in:

- 1) Local cache
- 2) Neighborhood cache
- 3) HTTP cache
- 4) LDF Server

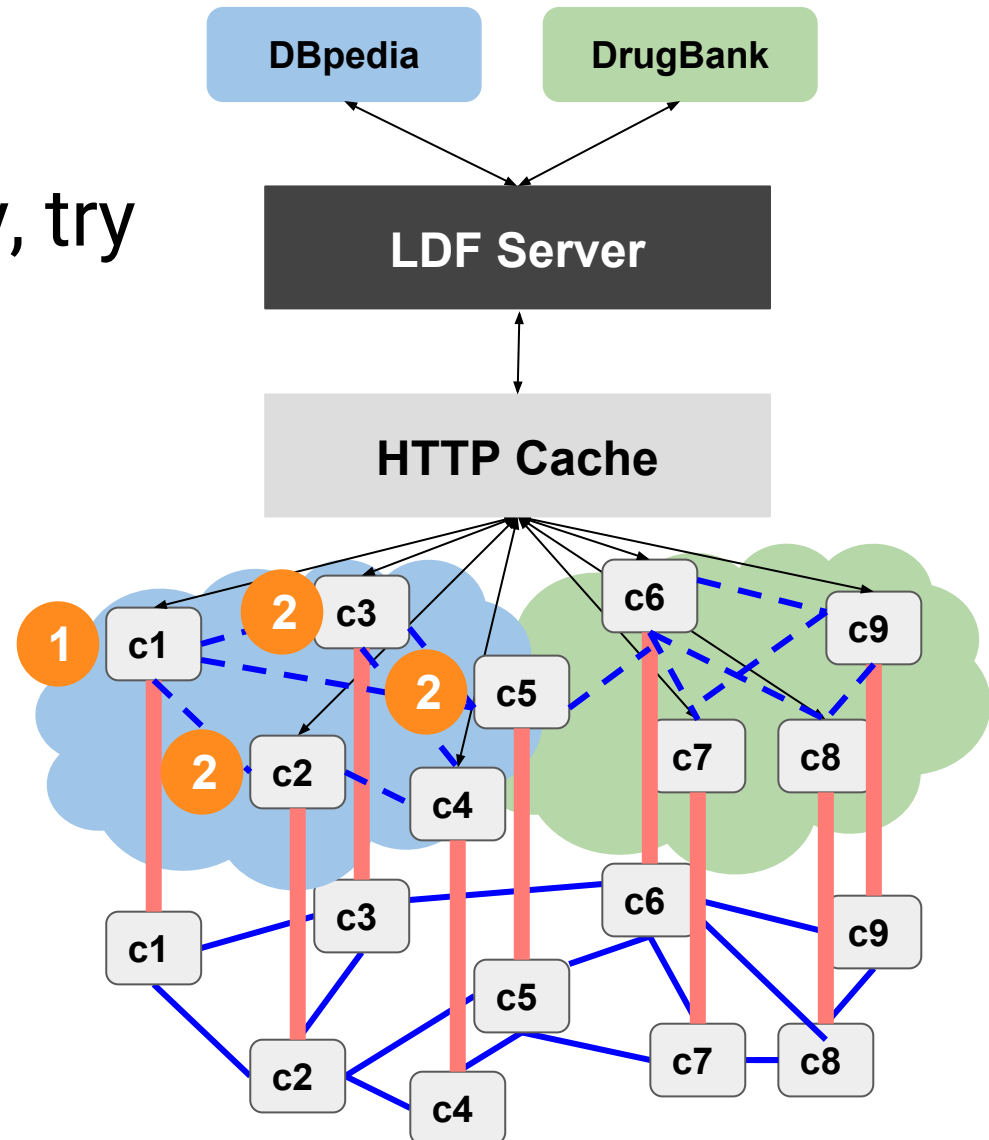


# Queries with CyCLaDEs

C1 executes query Q1.

For each triple pattern query, try to resolve pattern in:

- 1) Local cache
- 2) Neighborhood cache
- 3) HTTP cache
- 4) LDF Server

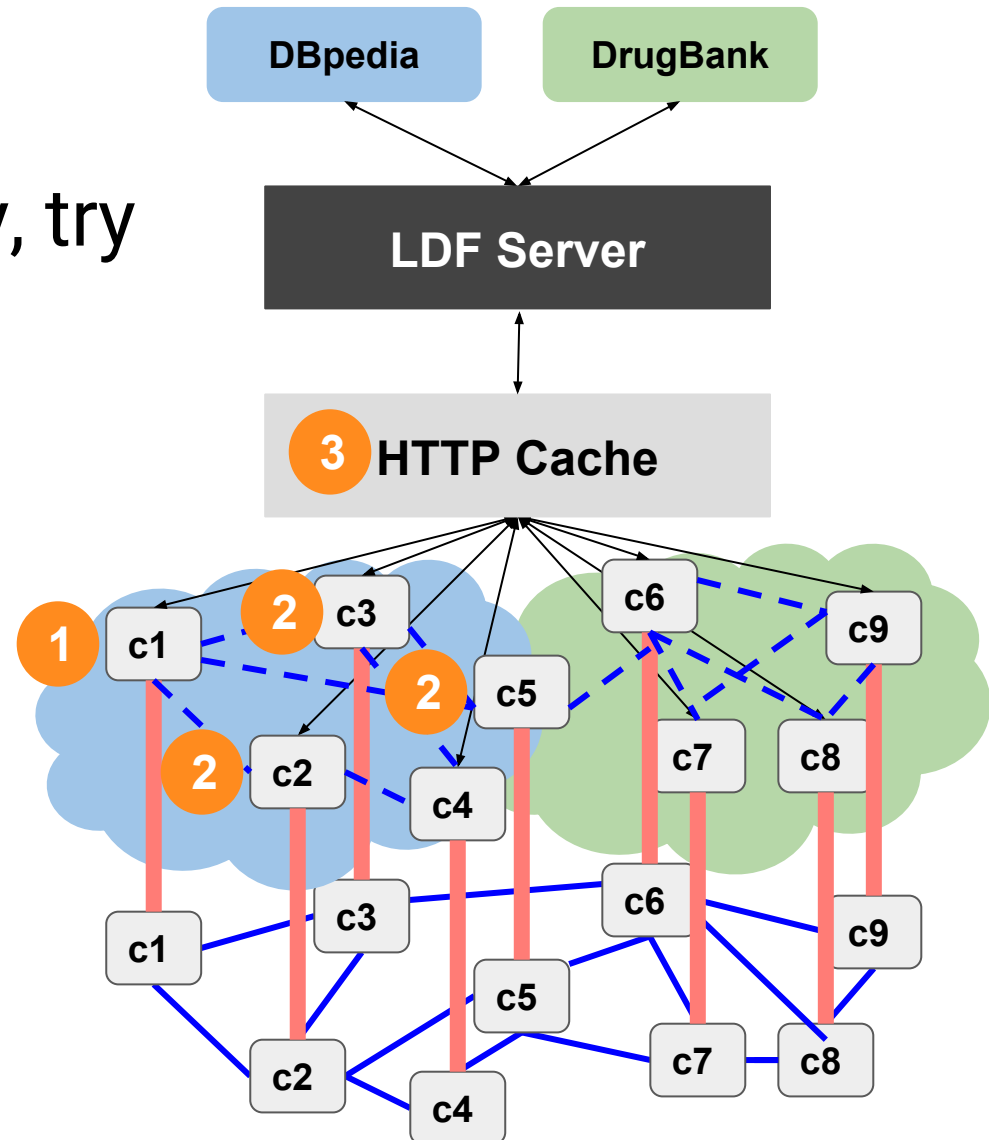


# Queries with CyCLaDEs

C1 executes query Q1.

For each triple pattern query, try to resolve pattern in:

- 1) Local cache
- 2) Neighborhood cache
- 3) HTTP cache
- 4) LDF Server

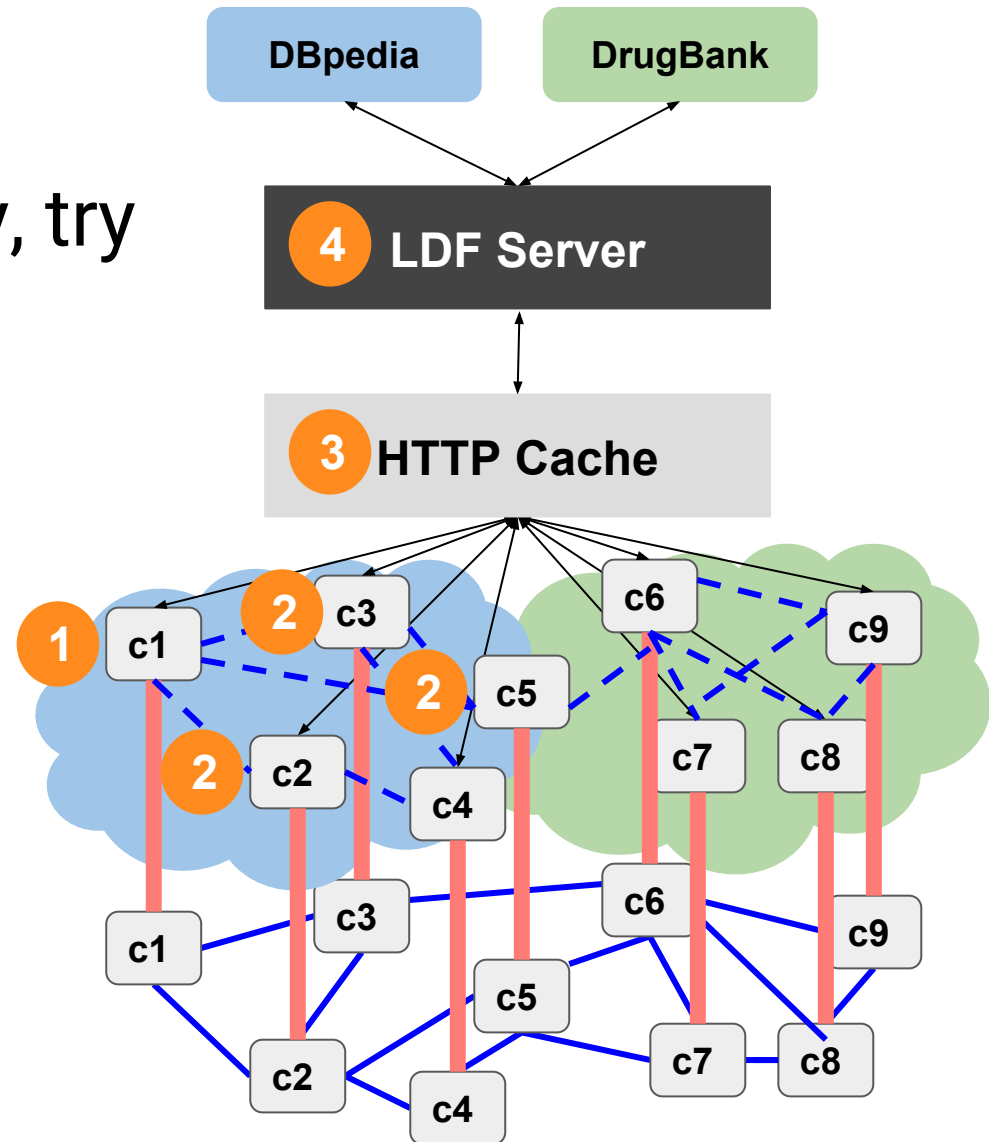


# Queries with CyCLaDEs

C1 executes query Q1.

For each triple pattern query, try to resolve pattern in:

- 1) Local cache
- 2) Neighborhood cache
- 3) HTTP cache
- 4) LDF Server





# Experiments


Berlin SPARQL Benchmark (BSBM) <sup>[7]</sup> chosen for simulating **Web applications**.

One LDF Server with one classic HTTP cache hosting BSBM datasets.

A network of extended LDF clients\* running a query mix of BSBM queries.

\* <https://github.com/pfolz/cyclades>


[7] C. Bizer and A. Schultz. The berlin sparql benchmark. *International Journal on Semantic Web and Information Systems*, 2009.



# Experiments - Parameters

Each **client** has a **query mix** of **25 queries** generated from **12 templates**.

- Shuffling phases occur every 10 seconds
- One **warmup round** followed by one **real round**




**BSBM 1M, cache = 1000, profile size = 10**


**10 clients: RPS = 4, CON = 9**

**50 clients: RPS = 6, CON = 15**

**100 clients: RPS = 7, CON = 20**




Without CyCLaDEs



With CyCLaDEs

**~ 20% neighbors cache hit-rate, whatever the number of clients**

10 clients, RPS = 4, CON = 9, cache = 1000



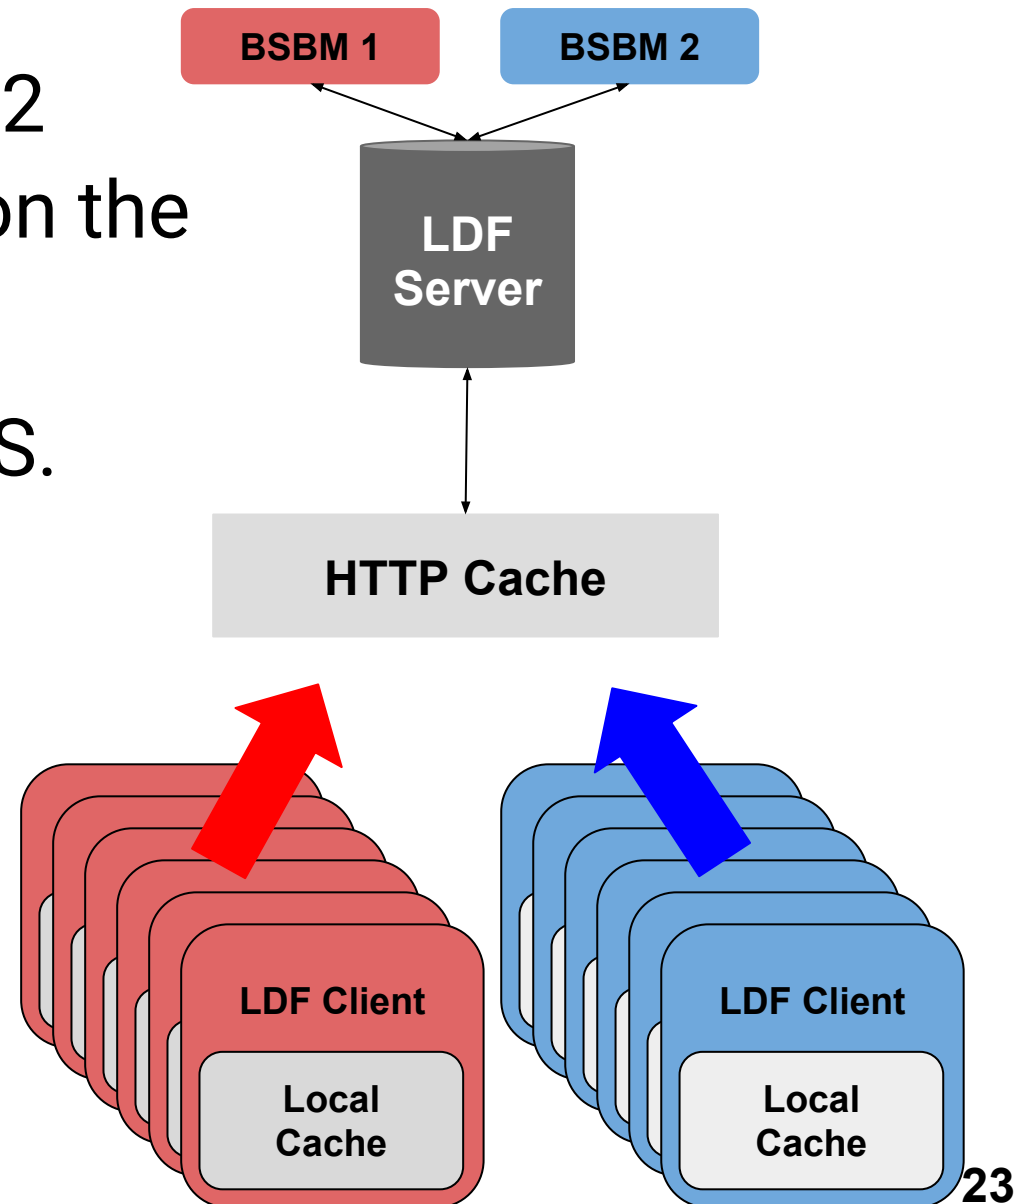
Behavioral cache better resists than local cache

# CyCLaDEs with 2 Communities


Two communities access 2 **different** BSBM datasets on the same server.

Nodes share the same RPS.


**Will CyCLaDEs  
profiles detect  
communities?**



# 2 BSBM 1M datasets, 50 clients per data set, cache = 1000



Profile = 5



Profile = 30

In CON overlay, CyCLaDEs builds two distinct communities BSBM1 and BSBM2

# Conclusion

CyCLaDEs builds a **behavioral decentralized cache** for LDF clients.

CyCLaDEs **reduces calls** to the **server** in the context of **Web applications**.



## Towards a Federation of Data Consumers



# Future Works

Measure the impact on **execution time**.

CyCLaDEs **brings** the **data** to the **queries**.

**Bring queries** to the **data**.



United Federation of Data Consumers



# CyCLaDEs: A Decentralized Cache for Triple Pattern Fragments


**Pauline Folz, Hala Skaf-Molli & Pascal Molli**  
ESWC 2016 - May 2016 - Heraklion



UNIVERSITÉ DE NANTES




## 2 BSBM 1M dataset, 50 clients per data set, cache = 1000




Size of the profile does not impact the number of calls handled by neighborhood, queries use at most 16 different predicates at once

# BSBM 10M, 10 clients, RPS = 4, CON = 9, profile size = 10



Percentage of calls handled by the neighborhood is related to cache size. With a larger cache size we have a bigger decentralized cache

# Query load - Is there a hotspot?



Most of the clients handled around 10,000 cache queries and a few handled more than 100,000 cache queries

# Query Distribution

- Each client has its own query mix
- Each query mix:
  - is generated randomly
  - has 25 queries based on 12 templates
- Query mix:
  - is the same for both rounds
  - is executed in the same order for both rounds

# Shuffle

- Shuffle phases are executed each 10 seconds
- In theory shuffle phases do not impact query processing
  - In implementation, query processing can be interrupted because NodeJs is mono-thread
- Shuffle phase for BSBM 1M:
  - 10 clients  $\approx$  14 shuffles / client
  - 50 clients  $\approx$  80 shuffles / client
  - 100 clients  $\approx$  85 shuffles / client

# Experiment

- 2 rounds:
  - **Warmup** → Bootstrap RPS, CON and HTTP cache
  - **Real** → measures are done
- All clients execute the same query mix for both rounds
- All queries are executed in the same order

# Overlay size

- Random Peer Sampling:
  - RPS view size is **Log(N)**, where **N** is the **number of peers** in the network
- Cluster Overlay Network:
  - CON view size has been chosen following the guidelines of <sup>[5]</sup>

[5] M. Bertier and al. The gossple anonymous social network. In *11th International Middleware Conference Middleware 2010 - ACM/IFIP/USENIX*, volume 6452 of LNCS, pages 191{211. Springer, 2010.