

Optimizing the Performance for Concurrent RDF Stream Processing Queries

Chan Le Van, Feng Gao, and Muhammad Intizar Ali

INSIGHT Centre for Data Analytics,
National University of Ireland, Galway
Emails: {firstname.lastname}@insight-centre.org

Abstract. With the growing popularity of Internet of Things (IoT) and sensing technologies, a large number of data streams are being generated at a very rapid pace. To explore the potentials of the integration of IoT and semantic technologies, a few RDF Stream Processing (RSP) query engines have emerged, which are capable of processing, analyzing and reasoning over semantic data streams in real-time. RSP mitigates data interoperability issues and promotes knowledge discovery and smart decision making for time-sensitive applications. However, a major hurdle in the wide adoption of RSP systems is their query performance. Particularly, the ability of RSP engines to handle a large number of concurrent queries is very limited which refrains large scale stream processing applications (e.g. smart city applications) to adopt RSP.

In this paper, we propose a shared join based approach to improve the performance of an RSP engine for concurrent queries. We also leverage query federation mechanisms to allow distributed query processing over multiple RSP engine instances. We apply load balancing strategies to distribute queries and further optimize the concurrent query performance. We provide a proof of concept implementation by extending CQELS RSP engine and evaluate our approach using existing benchmark datasets for RSP. We also compare the performance of our proposed approach with the state of the art implementation of CQELS RSP engine.

Keywords: Linked Data, RDF Stream Processing, Query Optimization

1 Introduction

A merger of semantic technologies and stream processing has resulted into RDF Stream Processing (RSP). RSP engines refer to the stream processing engines which have the capability to process semantically annotated RDF data streams. Over the past few years, different RSP engines have been proposed, such as CQELS [15], C-SPARQL [3], and SPARQLStream [5] etc. An RSP engine continuously consumes streaming RDF data as input and generates query results as output. An RSP query is registered once and executed constantly. The continuous query execution is a resource intensive and most of the existing RSP engines suffer from scalability and performance issues while processing many concurrent queries. The inability of the RSP engines to handle a large number of concurrent queries is a major performance bottleneck, hence, the main obstacle to the wider adoption of RSP engines. Inability of the RSP engines to share

the resources among different concurrent queries is one of the major causes of the performance bottleneck. For example, in CQELS, each input query is separately compiled and parsed into an execution plan with a separate data buffer. Therefore, its performance may decline drastically whenever an increasing number of queries are registered over a centralized instance of CQELS engine. Many real-time applications consuming continuous data streams often share input data streams for a large number of concurrent queries. We consider that resource sharing capability of RSP engines can help improving the performance of RSP engines. Ideally, any similar and concurrent queries must share the resources required to process common input data streams.

In this paper, we propose resource sharing and load balancing techniques to optimize the performance of concurrent RSP queries. We take the concept of a shared join operator (also known as multiple join operator) introduced in [15], to support sharing of memory and computation resources. Taking the advantage of queries having the shared input streams, the multiple join operator uses a join component to produce shared results for all queries, and a routing component/router is used to route the output items to the corresponding query. We extend the approach of shared join operator and introduce the creation of join sequences using a *reutilization* metric. Our approach creates a network of shared join operators for multiple queries and this network of shared join operators contains various shared join operators which are consuming the same input streams. In order to produce query results, the query evaluation process also involves a traversal of a path in network of shared joins, which is called join sequence. In our approach, intermediate results of a query can be shared with other queries, if they have triple patterns referring to the same stream data buffers. *Reutilization* metric is used to choose optimal join sequences. We extended the existing implementation of CQELS engine and named it as CQELS+. We evaluated CQELS+ in comparison to the existing state-of-the-art CQELS implementation, which showcases that CQELS+ is capable of handling a larger number of concurrent queries. We further evaluated CQELS+ performance for different load balancing techniques for multiple instances of RSP engines deployed in a distributed environment. We evaluated the performance of multiple instances of CQELS+ and compare it with the performance of a single engine instance. The evaluation results revealed that we can have lower query latency by deploying parallel engine instances.

The remainder of this paper is organized as follows;. Section 2 lays the theoretical foundations for our approach. Section 3 elaborates the details of the shared join operations as well as load balancing strategies for parallel RSP engines. Section 4 presents the experimental design and evaluation approach. We also presented results of our empirical analysis in this Section. Section 5 reviews the literature for related works before we conclude and discuss future plans in Section 6.

2 Foundations

In this section, we introduce the basic concepts of RSP, including multi-way join, shared join operator and network of multiple join operators.

Following the principles of Linked Data [4], Linked Stream Data [22] or Linked Streams were introduced to bridge the gap between data streams and Linked Data. Linked Streams not only simplify the data integration process among various streaming sources, but also between streaming and static sources. The following concepts formulate the basis of RDF Stream Processing model, which is also used by CQELS.

- **RDF stream** is a bag of elements $\langle (s, p, o) : [t] \rangle$, where (s, p, o) is an *RDF triple* [20] and t is a timestamp.
- **Window operators** are inspired by the time-based and triple-based window operators of CQL data stream management system [2].
- **Stream Graph Pattern** is supplemented into the *GraphPatternNotTriples* pattern¹ to represent window operators on RDF Stream.

2.1 Multi-way Join

A single multi-way join works over more than two streaming data buffers. It is also known as window buffers [18, 15]. Single multi-way join generates and propagates results in single step rather than creating the join results by passing through a multiple-stage binary joins in the query execution pipeline. Suppose we have n window data buffers W_1, \dots, W_n involved in the join operation. A multi-way join operates as follows:

1. When a new stream data item (tuple) comes to any window buffers, it is used to recursively probe other window buffers to generate the join output, i.e., when a data item M arrives at W_1 , the join processing starts to look for next join candidate in W_2, W_3, \dots, W_n .
2. Let us assume W_2 is the next joining candidate, the algorithm will check data inside W_2 to see if there are existing tuples compatible with M .
3. If any compatible data item is found, the intermediate result sets are created and the algorithm recursively continues until all of the involved window buffers are visited. The algorithm eventually terminates if any further compatible data items can not be found.

2.2 Shared Join Operator and Network of Shared Join Operators

In [15], authors introduced the *shared join operator* as a mechanism to share the computation among multi-way join operations from queries with the same set of window buffers of RDF streams. Let us assume, we have k multi-way join operators and each of them involves m window buffers i.e., we have k multi-way join operators: $(W_1^1 \cdots \bowtie W_1^m), \dots, (W_k^1 \cdots \bowtie W_k^m)$. Additionally, let us assume that with any $1 \leq i \leq m$, we always have $W_1^i, W_2^i, \dots, W_k^i$ referring to the same window buffer, then the equation below (containment property [12]) is always true:

$$W_j^1 \cdots \bowtie W_j^m \subseteq W_{max}^1 \cdots \bowtie W_{max}^m \quad (1)$$

¹ SPARQL 1.1 Grammar: <https://www.w3.org/TR/sparql11-query/#grammar>

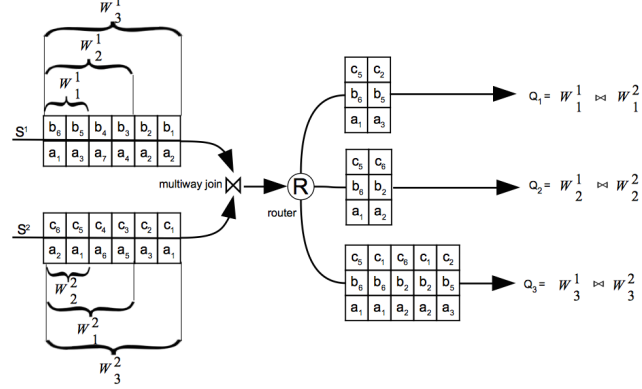


Fig. 1: Shared join operator example

In Equation 1, W_{max}^i refers to the window buffer in the set $\{W_1^i, W_2^i, \dots, W_k^i\}$ with the maximum size. We call W_{max}^i , a shared window buffer.

Intuitively, Equation 1 tells us that the join results of any multi-way join operator for a fixed set of window buffers are included in the join results of the operator for the set of windows with the maximum length (that used by all join operators). For example, Figure 1 shows a shared join operator for three queries, namely Q_1 , Q_2 and Q_3 , which contain 3 multi-way joins $W_1^1 \bowtie W_1^2$, $W_2^1 \bowtie W_2^2$ and $W_3^1 \bowtie W_3^2$, respectively. These queries receive input from two RDF streams: S^1 and S^2 . As depicted in Figure 1, W_{max}^1 will be W_3^1 and W_{max}^2 will be W_3^2 . By applying Equation 1, we have:

$$\forall j \in \{1, 2, 3\}, W_j^1 \bowtie W_j^2 \subseteq W_3^1 \bowtie W_3^2 \quad (2)$$

Leveraging the containment property, an approach discussed in [15] designed the shared join operator from queries having the above mentioned characteristics. It consists of two components: join component and routing component. The join component has the duty of generating the shared results that contain results of all involved queries. The routing component is responsible for filtering the shared results and routing the filtered results to proper queries.

In practice, the queries registered in the processing engine often share only subsets of input streams. The concept of *network of shared join operators* (NSJO) can share the execution for sub-queries (part of the whole queries) within the group of queries. This network includes a set of shared join operators. Each shared join operator consumes the same window buffers. When a new tuple comes to a shared window buffer, it triggers the join components of related shared join operators in the relevant NSJO. The join component then chooses the best cost join sequence (the order of joined window buffers) in all possible join sequences and generates the joined output. During the join process, in case, if the generated results are the output of any shared join operator, the routing component of that shared join operator routes them to proper queries. Otherwise, if the results are further consumed by other shared join operators in the network, the algorithm

recursively continues until all related shared join operators have been visited or no results are created.

3 Optimization for Concurrent CQELS queries

In this section, we elaborate our approach for optimizing CQELS performance under concurrent queries. Firstly, we discuss how NSJO is introduced in CQELS+ (the extended CQELS), including creation of the join sequences (the order of joined window buffers) in a join graph, and calculation of the reutilization metric. We also provide an example of the join graph. Lately, we elaborate our deployment of multiple CQELS+ engines in parallel and distributed fashion to share the workload among different engine instances to obtain better performance.

Algorithm 1: *create_Join_Sequences(curr_Join_Vertex, queries)*

```

1 while coverage < number_of_queries do
2   HAQ ← 0; considering_Patterns ← ∅ ;
3   for  $j \in [1..number\_of\_physical\_windows]$  do
4     cal ← cal_Reuse_Metric(cur_join_Vertex, physical_Windows[j]);
5     patterns = get_Considering_Patterns(cur_join_Vertex,
6     physical_Windows[j]) ;
7     if cal > HAQ then
8       HAQ = cal; next_Buffer = physical_Windows[j];
9       considering_Patterns ← patterns;
10    end
11  end
12  if HAQ > 0 then
13    next_Join_Vertex ← create_Join_Vertex(next_Buffer) ;
14    next_Join_Vertex.set_Considered_Patterns(considering_Patterns) ;
15    curr_Join_Vertex.add_Nexts(next_Join_Vertex) ;
16    prev_Join_Vertex ← curr_Join_Vertex ;
17    curr_Join_Vertex ← next_Join_Vertex ;
18    coverage += count_Covered_Queries(curr_Join_Vertex) ;
19    create_Join_Sequences(current_Join_Vertex) ;
20    curr_Join_Vertex ← prev_Join_Vertex ;
21  end
22  else
23    return;
24  end

```

3.1 CQELS+: Network of Shared Join Operators

According to [15], when a new tuple comes to the shared window buffer (a physical window in our approach) in the NSJO, join sequences in the join graph of this physical window are evaluated to generate the join results for all concerned queries. In our approach, we build join graph at the query registration time based on a metric called *reutilization*. This metric is defined to create the optimal join

sequences in the graph. More specifically, in each step of forming a join sequence, one physical window is selected to be the next join candidate if the generated join results can be consumed by the most queries. The process of creating a join graph has two steps; *create_Join_Sequences* as described in algorithm 1 and supplement join sequences for *special queries* containing multiple stream patterns referring to the same physical window.

Each Shared Window Buffer (SWB) from the involved queries triggers algorithm 1 to create the join graph. The *curr_Join_Vertex* parameter for the first call is the vertex containing the SWB itself and SWB's graph patterns. The *coverage* variable in line 1 holds the number of queries that have been considered. This variable is initialised by the value 0 and it is used to exit the algorithm. Line 2 initializes the *HAQ* and *considering_Patterns* variables. HAQ is intended to hold the *Highest Amount of Queries* reusing the generated join results. The *considering_Patterns* keeps the considering stream patterns in the queries sharing join variables with the current join vertex. From line 3 to line 8, the algorithm iterates over all of the involved physical windows and calculates the *reutilization metric* detailed in the next section. The HAQ and *considering_Patterns* variables are updated whenever a higher calculated value is found. If the condition in line 6 evaluates to *true* (i.e the next join candidate is found), we reset the pointers for the current vertex and previous vertexes based on the newly found candidate as from line 12 to 16. With the chosen join candidate, we count the number of covered queries and update the coverage variable (line 17). At line 18, we recursively call the algorithm with the input parameter of the selected join candidate. Line 19 assigns the previous state to the join candidate. Finally, the algorithm terminates in line 23 if we are not able to find any further join candidate.


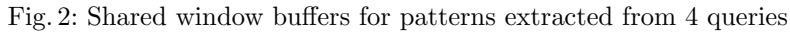
Calculating the reutilization metric: Given n queries (Q_1, \dots, Q_n) , m join sequences (S_1, \dots, S_m) that can share the join results for a set of sub-queries in the current vertex C , t patterns (P_1, \dots, P_t) referring to the physical window in C . We define $J_j^i(k, l, v)$ as a counter for calculating the reutilization for variable v , on sequence S_i and pattern P_j , where $v \in V = \{\text{set of variables in all queries}\}$, $1 \leq i \leq m, 1 \leq j \leq t, l = \text{index of joining variable in } P_j$, e.g., in the first triple pattern in PW1 (Figure 2), the index of variable (*?person.q1*) and (*?loc.q1*) are 1 and 3, respectively². $k = \text{index of the joining variable in } S_i$, $J_j^i(k, l, v)$ is given in the equation below:

$$J_j^i(k, l, v) \begin{cases} = 1, & \text{if } P_j \not\subset S_i \wedge S_i \in Q_y \wedge p_j \in Q_y \\ & \wedge \text{var}(S_i, k) = v \wedge \text{var}(P_j, l) = v \\ = 0, & \text{if otherwise.} \end{cases} \quad (3)$$

The functions $\text{var}(S_i, k)$ and $\text{var}(P_j, l)$ allow to retrieve the variable of a sequence S_i at index k or the variable of a pattern P_j at index l . We calculate the metric based on the equation 4:

$$R = \text{Max} \left(\sum_{v \in V} \sum_{k \in [1, |S_i|]} \sum_{l \in [1, |P_j|]} J_j^i(k, l, v) \right) \quad (4)$$

² The index of join variable in join sequence is similarly defined.



Supplement Join Sequences: Special queries mentioned in section 3.1 require more than one join sequences to generate the join results because any arrived tuple in the shared window buffer is matched with multiple stream patterns inside one query. As Algorithm 1 only considers one join sequence for each query, we need to supplement the rest of join sequences. To do this, we count the number of special queries and rerun Algorithm 1 with the condition that created join sequences are not taken into account.

³ Q4 in this paper is actually Q5 in [6], we do not use Q4 in [6] because its join sequence is too long for the demonstration.

candidates. There are three shared join variables in three join sub-sequences. In this case, the join results are reused by three queries which are Q1, Q3 and Q4. Continuing on this path, PW3 is visited next because it has highest reutilization of 2, and the join results on this path are re-used by Q1 and Q3. At this point, all patterns in Q1 are visited and Q1 is covered. After this buffer, PW1, PW4 and PW4 (again) are respectively chosen and Q3 is also covered. Similar processes are repeated to cover all involved queries, then the algorithm stops. Q2, Q3 and Q4 are covered more than once as shown in the dashed branches in Figure 3. These three special queries require multiple join sequences to generate enough join results. Therefore, we have to run the supplementation step to create join sequences starting from these patterns to make sure, that we do not miss the generated join results. In Figure 3, the red color indicates the queries covered after running Algorithm 1, while the blue color demonstrates the queries covered after the *supplement join sequence* step.

3.2 Load Balancing for Parallel CQELS+ Instances

In Section 3.1 we discussed the means for optimizing RSP performance leveraging shared join operations. However, the described approach is still a centralized one, i.e., the join graphs and all queries are managed by a single machine. While this is feasible for small-scale applications, it cannot satisfy the need for large-scale systems, where multiple servers are deployed, possibly at different geographical locations, to handle the excessively high concurrency. Evidence for the concurrency limitations for RSP can be found in CityBench[1]. Notice that in CityBench only duplicated queries are tested over a single engine instance. To cope with large-scale applications, a distributed way of processing RDF streams is necessary. In [11, 9], a service-oriented approach was proposed to realize RSP federation via service composition. This way, multiple engine instances, even with different engine types, can be deployed in parallel to collaboratively answer an RSP query. In this paper, we follow this principle and develop means for an efficient distributed RSP.

In [10], the problem of multi-modal QoS optimization for the event service composition plan is studied, i.e., it provides means for creating optimal RSP query federation. However, it does not address the problem of the overall system performance. In particular, it does not provide heuristics on which specific engine instance should be used when a new federated query is generated. To determine the workload assigned to multiple RSP engine instances, we implemented a query scheduler for RSP engines, as illustrated in Figure 4.

The architecture shown in Figure 4 consists of a client and multiple servers hosting RSP engines (CQELS+ instances in this paper). The client is a centralized controlling component, which accepts RSP queries from applications and utilizes the data federation component (described in [9]) to create the query federation plan. Then, the scheduler determines which engine instance should be picked for evaluating the plan (which is also a CQELS query) at runtime. The scheduler communicates with the performance monitors implemented on the servers and continuously receives information on the status of the servers. Currently, we implemented three different heuristics (evaluated in Section 4.2) for the load balancing strategy:

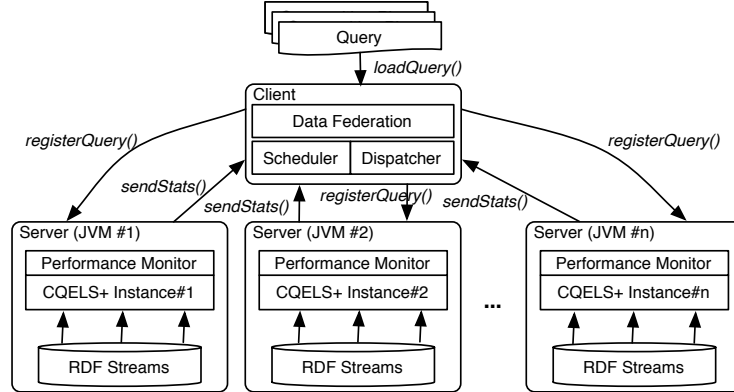


Fig. 4: Architecture for load balancing over parallel CQELS+

- **Rotation:** each instance takes its turn to deploy a new query in a rotation, i.e., we distribute queries evenly among the engine instances;
- **Minimal average latency:** the engine instance with the lowest average query latency is picked to deploy the new query. For the latency monitoring mechanism, we follow the approach provided by CityBench [1], and
- **Minimal data buffer size:** this strategy chooses the current instance which has the minimum total number of elements in data buffers, to register the next query, as the processing time of the join operator depends significantly on the join selectivity of the data buffers. Heuristically, smaller data buffers typically have smaller join selectivity.

4 Evaluation

In this section, we conduct three experiments. First we compare our shared join approach in the CQELS+ engine with the original CQELS. Then, we show the performance of the load balancing over CQELS+ engines. Finally, we evaluate the query registration time for multiple queries. All experiments are deployed on a machine running Ubuntu 12.04 with Intel Quad-Core i7-3520M CPU (2.90 GHz) and 16 GB RAM. The tests are compiled for 64-bit Java Virtual Machine (JVM build 1.7.0_80b15). All experimentation results are reproducible⁴. In the following, we present the detailed design of the experiments and analyze the results.

4.1 Evaluating Shared Joins in CQELS+

In this experiment, we reproduce the experiments in LS-Bench⁵ and compare the performance between the publicly available version of CQELS⁶ with our

⁴ CQELS+: <https://github.com/chanlevan/CqelsplusLoadBalancing>

⁵ LS-Bench: <https://code.google.com/archive/p/lsbench/>

⁶ CQELS engine: https://code.google.com/archive/p/lsbench/wikis/howto_cqels.wiki

implementation of CQELS+. The performance is measured in throughput, i.e., the number of triples processed per unit time. We choose 4 queries: Q2, Q3, Q5 and Q6 from LS-Bench for this evaluation, since they involve the streaming and static data. For each query, we increase the number of duplicated queries registered to the engines.

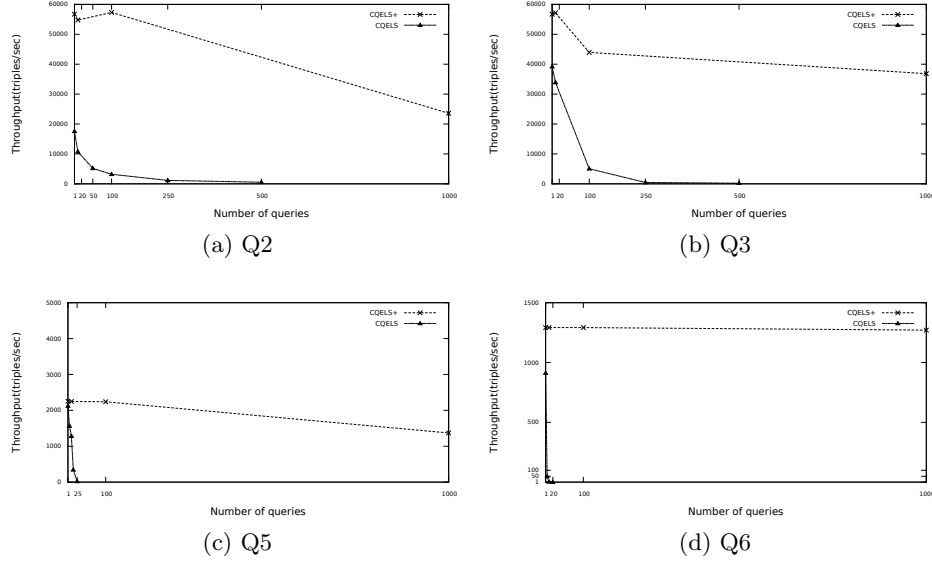


Fig. 5: Throughput comparison between CQELS+ and CQELS

Figure 5 shows the results of this experiment. CQELS+ outperforms CQELS in throughput as well as in handling concurrent queries. The results in Figure 5a and 5b point out that for Q2 and Q3, CQELS is not able to process more than 500 queries and the throughput is about 20,000 and 40,000 triples per second, respectively. As indicated in Figure 5c and 5d, CQELS responses if the number of queries is not higher than 25 and the maximum throughput is about 2,200 triples per second for Q5 and 800 for Q6. The results show that the throughput of CQELS decreases drastically when we increase the number of concurrent queries. Conversely, CQELS+ can handle one thousand queries with higher throughput, about 60,000 triples per second for Q2 and Q3, nearly 2,300 for Q5 and 1,300 for Q6. The performance of CQELS+ reduces slowly when we increase the number of queries.

These results are because of two main factors: the join operator and the scheduling mechanism of the JVM. In the CQELS version we evaluated, the join operators are repeatedly triggered by the arrivals of data elements and the consecutive stream windows and data arriving at those windows are processed independently. Furthermore, if the input data coming to stream windows has the high join selectivity and a large number of involving stream windows, a lot




Fig. 6: Latency with increasing CQELS+ instances




Fig. 7: Latency with increasing number of streams

of intermediate results are generated. The resources used by those intermediate results have to be released frequently, which causes considerable overhead for the JVM. All of these factors have limited the capability of CQELS for handling concurrent queries. CQELS+ instead, on the same data input (same join selectivity and involved stream windows), manipulates the join strategy by using the incremental evaluation over the network of shared join operators. This strategy only processes the changes of data in stream windows rather than processing consecutive stream windows independently. When multiple continuous queries are registered inside the engine, CQELS+ shares the processing for queries whenever possible. Notably in Figure 5c, the throughput of the two engines are almost equal when they evaluate one query. This is because the window size in the query is too long (1 day) and contains very big data buffers, which makes it difficult for CQELS+ to look for compatible join elements.

4.2 Evaluating Load Balancing over CQELS+

In this experiment, we run two different tests. We show the results of the load balancing model presented in Section 3.2 in the first test. Then, we choose the best policy and test the scalability of the system with increasing number of streams. To demonstrate the both tests, we implement the client-server architecture as shown in Figure 4. A new query is registered only after the previous query has been successfully registered to the engine. After all the queries are deployed, we keep the system running for 15 minutes and monitor the latency. For the input data, we use the CityBench[1] as they support end-to-end query latency monitoring. The latency is captured as follows: each query in this test has some joins over sensor data streams, and the queries are slightly adjusted (e.g., removing numerical filters) so that each sensor observation (in the form of a group of triples) can produce at least one result, then, the latency can be derived by comparing the timestamp of the first result produced by each sensor observation and the timestamp of observation entering the triple streams. Notice that all sensor streams produce observations simultaneously with a same frequency.

In the first test, we deploy different CQELS+ engines in different JVMs and apply the aforementioned load balancing strategies, i.e rotation, minimum

average latency and total buffers size. We choose 4 different queries: Q1, Q2, Q3 and Q4 (from CityBench). We register 400 queries (randomly picked from Q1 to Q4) to the system. Figure 6 illustrates the results. With a single instance, the latency for 400 queries is up to 4.5 seconds. The latency decreases about 4 times when we deploy 2 and 3 CQELS+ instances. With this configuration, the buffer size load-balancing strategy is the most efficient one, while the other two strategies also improve the latency for 2 and 3 instances. The reason behind this is perhaps that the number of data buffer size is more accurate in representing the workload, and the latency observed may fluctuate in time. When we use 4 instances, the latency tends to increase. This indicates that the overhead of multiple JVMs and CQELS+ instances starts to outweigh the benefit of load balancing.

The second test aims to check the latency of the multiple CQELS+ instances when we vary the number of streams. Previous experiments (including the LS-Bench results) tested only a limited number of streams. Now we increase the number of streams and monitor the latency over 2 and 3 CQELS+ instances with the data buffer size policy. These streams are picked in turns by 400 randomly generated different queries with Q1 as a template. Figure 7 shows the experiment results. Generally, the latency increases when more streams are used, and with more than 7 streams some unstable processing states can be observed (e.g., latency “spikes” for 7 and 20+ streams). We also observed that with more than 30 streams, the engine often stops responding. The latency increases when we scale the number of streams from 3 to 25 streams in the both configurations (2 and 3 instances). The more streams are involved, the more concurrent threads are created and invoked to stream the data into the system. This makes the system response slower when the number of streams increases. However, the abnormal increase in latency appears with the 7 streams and 25 streams in Figure 7. This abnormal behaviour is perhaps caused by an incorrect engine implementation, which may also be the cause for the system failure when the number of streams is higher than 30.

4.3 Evaluating the Query Registration Time

In the first experiment, we showed that the CQELS+ outperforms CQELS at runtime. However, this is because we build the join probing graphs before query execution. This introduces an overhead for registering new queries: the graph must be updated constantly. Also, the more queries are deployed, the more time it takes to update the graph. Figure 8 shows the time taken for the query registration using 400 random queries from the template of Q1 in the second experiment, using different load balancing strategies. From the results, we can see that for the data buffer strategy, it takes more time than the other two, possibly because this strategy has a higher chance of resulting in the different number of queries deployed on the engine instances.

5 Related work

Sharing query results is not a novel idea in data processing. For example, optimization for multiple queries for static database systems by sharing the query




Fig. 8: Query registration time for 4 CQELS+ instances

operations have been discussed in [21, 7, 8] etc. However, for stream processing systems more complicated query semantics need to be incorporated and the dynamicity of data buffers and input streams have to be considered. Existing RSP engines like C-SPARQL [3] uses a black box approach for handling the semantic queries, thus not able to provide much optimization for multiple queries. Benchmark results like CityBench [1] also showed that with a native join operator implementation, CQELS is better than C-SPARQL at handling multiple queries. Although an adaptive join routing is discussed in [15], where an estimated cost for the join sequences is calculated by monitoring the index of data buffers at run-time. However, it is currently not implemented in CQELS, thus we are not able to compare the performance difference to our approach. Also, we argue that a static join graph has less overhead at run-time while we acknowledge that we are relocating the complexity to query registration time.

Distributed RSP is also discussed recently. In [13] a C-SPARQL version with parallel streaming is developed mainly to optimize RDFS reasoning by splitting and filtering sub-streams. In CQELS Cloud [16], where extensions for CQELS are made for utilizing the elasticity of cloud environments and allow processing nodes to join and leave the network on-demand. However, the load shedding in CQELS Cloud relies on existing DSMS systems (e.g., Storm⁷), whereas in our approach different strategies are designed and tested. Load balancing techniques have been studied extensively in the literature [14, 19, 17]. Various metrics, from basic execution latency and bandwidth usage [14] to sophisticated service correlations [19] and network path analysis [17] have been proposed to evaluate the load and determine the shedding strategy.

6 Conclusions and Future Work

In this paper, we realized shared join operations for CQELS in order to improve its performance when handling multiple concurrent queries. Particularly, we have discussed when and how stream and static inputs can share the data. We provided a solution to share join operators. Our approach pre-processes the queries and builds a join graph before constructing the network of shared join operators. The join graph is constructed based on the heuristics that each vertex should generate join results reusable by as many queries as possible. Our experiments showed that CQELS+ can handle more concurrent queries with higher

⁷ Twitter Storm: <http://storm.apache.org/>

throughputs, compared to the original CQELS. In order to further improve the performance for concurrent RSP queries, we followed the principle of RSP query federation and applied load balancing strategies over distributed RSP engines. Evaluations for the load balancing strategies showed that deploying multiple CQELS+ instances can lower the latency, but the memory overhead for too many parallel instances may outweigh its benefit (e.g., for more than 4 instances on the experiment machine). Also, we found that while the minimal data buffer size strategy performs best at reducing the query latency, it has longer query registration time.

In future work, we plan to investigate on implementing the probing graph based on both the statistics of data in window buffers and join variable. Real-time data from the stream sources come to the system unpredictably, which means the data inside the window buffers is changing and consequently changes the join selectivity. Therefore, the join probing graph built based merely on join variable in patterns is not able to guarantee the optimal join probing sequences. On the other hand, we must consider the overhead when monitoring the dynamics of the data buffers. We also plan to investigate more sophisticated load balancing strategies, e.g., defining a more precise cost model for the query processing pipeline and use it to estimate the total cost after a new query registration. This may involve checking the similarity between queries using metrics like the number of shareable data buffers, the graph edit distance between the join probing graphs before and after the query registration etc.

Acknowledgment

This research has been partially supported by Science Foundation Ireland (SFI) under grant No. SFI/12/RC/2289 and EU FP7 CityPulse Project under grant No.603095. <http://www.ict-citypulse.eu>

References

1. M. I. Ali, F. Gao, and A. Mileo. Citybench: A configurable benchmark to evaluate rsp engines using smart city datasets. In *The Semantic Web-ISWC 2015*, pages 374–389. Springer, 2015.
2. A. Arasu, S. Babu, and J. Widom. The CQL continuous query language: semantic foundations and query execution. *The VLDB Journal*, 15(2):121–142, 2006.
3. D. F. Barbieri, D. Braga, S. Ceri, E. Della Valle, and M. Grossniklaus. C-sparql: Sparql for continuous querying. In *Proc. of WWW*, pages 1061–1062. ACM, 2009.
4. C. Bizer, T. Heath, and T. Berners-lee. Linked data - the story so far. *International Journal on Semantic Web and Information Systems*, 5:1–22, 2009.
5. J.-P. Calbimonte, H. Jeung, O. Corcho, and K. Aberer. Enabling query technologies for the semantic sensor web. *Proc. of IJSWIS*, 8(1):43–63, 2012.
6. M. H. Danh Le-Phuoc, Josiane Xavier Parreira. Linked stream data processing. 2012.
7. S. M. Deen and M. Al-Qasem. A query subsumption technique. In *Proceedings of the 10th International Conference on Database and Expert Systems Applications, DEXA '99*, pages 362–371, London, UK, UK, 1999. Springer-Verlag.
8. Y. Diao and M. J. Franklin. High-performance xml filtering: An overview of yfilter. *IEEE Data Eng. Bull.*, pages 41–48, 2003.

9. F. Gao, M. I. Ali, and A. Mileo. Semantic Discovery and Integration of Urban Data Streams. In *Proceedings of the 13th International Semantic Web Conference (ISWC'14), Workshop on Semantics for Smarter Cities*, 2014.
10. F. Gao, E. Curry, M. Ali, S. Bhiri, and A. Mileo. Qos-aware complex event service composition and optimization using genetic algorithms. In *Proceedings of the 12th International Conference on Service Oriented Computing*, Paris, France, 2014. Springer.
11. F. Gao, E. Curry, and S. Bhiri. Complex Event Service Provision and Composition based on Event Pattern Matchmaking. In *Proceedings of the 8th ACM International Conference on Distributed Event-Based Systems*, Mumbai, India, 2014. ACM.
12. M. A. Hammad, M. J. Franklin, W. G. Aref, and A. K. Elmagarmid. Scheduling for shared window joins over data streams. In *VLDB*. VLDB Endowment, 2003.
13. J. Hoeksema and S. Kotoulas. High-performance distributed stream reasoning using s4. In *Ordring Workshop at ISWC*, 2011.
14. M. Koerner and O. Kao. Multiple service load-balancing with openflow. In *High Performance Switching and Routing (HPSR), 2012 IEEE 13th International Conference on*, pages 210–214. IEEE, 2012.
15. D. Le-Phuoc. *A Native and Adaptive Approach for Linked Stream Data Processing*. PhD thesis, National University of Ireland Galway, IDA Business Park, Lower Dangan, Galway, Ireland, 2012.
16. D. Le-Phuoc, H. N. M. Quoc, C. Le Van, and M. Hauswirth. Elastic and scalable processing of linked stream data in the cloud. In *The Semantic Web–ISWC 2013*, pages 280–297. Springer, 2013.
17. H. Matsuba, K. Joshi, M. Hiltunen, and R. Schlichting. Airfoil: A topology aware distributed load balancing service. In *Cloud Computing (CLOUD), 2015 IEEE 8th International Conference on*, pages 325–332. IEEE, 2015.
18. V. J. F. Naughton and J. Burger. Maximizing the output rate of multi-way join queries over streaming information sources. In *VLDB*. VLDB Endowment, 2003.
19. G. Porter and R. H. Katz. Effective web service load balancing through statistical monitoring. *Communications of the ACM*, 49(3):48–54, 2006.
20. E. Prud'hommeaux and A. Seaborne. SPARQL query language for rdf. *W3C Recommendation*, 4:1–106, 2008.
21. T. K. Sellis. Multiple-query optimization. *ACM Trans. Database Syst.*, 13(1):23–52, Mar. 1988.
22. J. F. Sequeda and O. Corcho. Linked stream data: A position paper. In *SSN'09*, 2009.