


Intelligent Clients for Replicated Triple Pattern Fragments

Thomas Minier, Hala Skaf-Molli, Pascal Molli and
Maria-Esther Vidal



UNIVERSITÉ DE NANTES


ESWC 2018 - Heraklion, Greece
June 6th, 2018



TIB LEIBNIZ INFORMATION CENTRE
FOR SCIENCE AND TECHNOLOGY
UNIVERSITY LIBRARY


Introduction

- Following the **Linked Open Data** principles, data providers made available RDF datasets at low-cost using **TPF servers** [1]
- However, **servers availability** remain an issue:
 - Server down
 - Server heavily loaded



Server Availability

- Data providers **replicate RDF datasets**
 - DBpedia & LANL Linked Data Archive
- **Can we use replicated datasets to improve server availability?**
 - Yes, using **load-balancing**




SPARQL Query load-balancing between Replicated RDF Datasets

- **Good for data providers**
 - Less load -> more available
 - Save €€€ on data hosting
- **Good for data consumers**
 - Tolerance to server failures
 - Tolerance to heavily loaded servers
 - Improve query performance



Problem

How to **balance the load of SPARQL query processing** over replicated **heterogeneous** servers owned by **autonomous** data providers?



Related Work

Triple Pattern Fragments

Existing TPF clients allow to process a federated SPARQL query over a federation of TPF servers [1], *but they **do not support replication nor client-side load balancing***

```
Q1: SELECT DISTINCT ?software ?company
WHERE {
    ?software dbo:developer ?company.      (tp1)
    ?company dbo:locationCountry ?country. (tp2)
    ?country rdfs:label "France"@en.      (tp3)
}
```

DBpedia	11.4s
DBpedia and LANL	28.7s

Linked Data Replication

- Linked Data Replication addressed as a **source-selection** problem [2, 3, 4]
- They **prune** redundant sources != load-balancing

Q1: SELECT DISTINCT ?software ?company
WHERE {
 ?software dbo:developer ?company. (tp_1)
 ?company dbo:locationCountry ?country. (tp_2)
 ?country rdfs:label "France"@en. (tp_3)
}

DBpedia	11.4s
DBpedia or LANL	11.4s or 36s

[2] Montoya, G. et al. "Federated Queries Processing with Replicated Fragments." *ISWC 2015*.

[3] Montoya, G. et al. "Decomposing federated queries in presence of replicated fragments" *Web Semantics: Science, Services and Agents on the World Wide Web (2017)*

[4] Saleem, M. et al. "DAW: duplicate-aware federated query processing over the web of data" *ISWC 2013*

Client-side load-balancing


- **Client-side load-balancing** is well suited for heterogeneous servers [5]
 - + Fit well for intelligent TPF clients
 - + Respect data providers autonomy
 - - Only applied for static files, not for query processing

Ulysses approach

Query evaluation over replicas




Q1: SELECT DISTINCT ?software ?company
WHERE {
 ?software dbo:developer ?company. (tp_1)
 ?company dbo:locationCountry ?country. (tp_2)
 ?country rdfs:label "France"@en. (tp_3)
}

Datasources: DBpedia and a replica from LANL

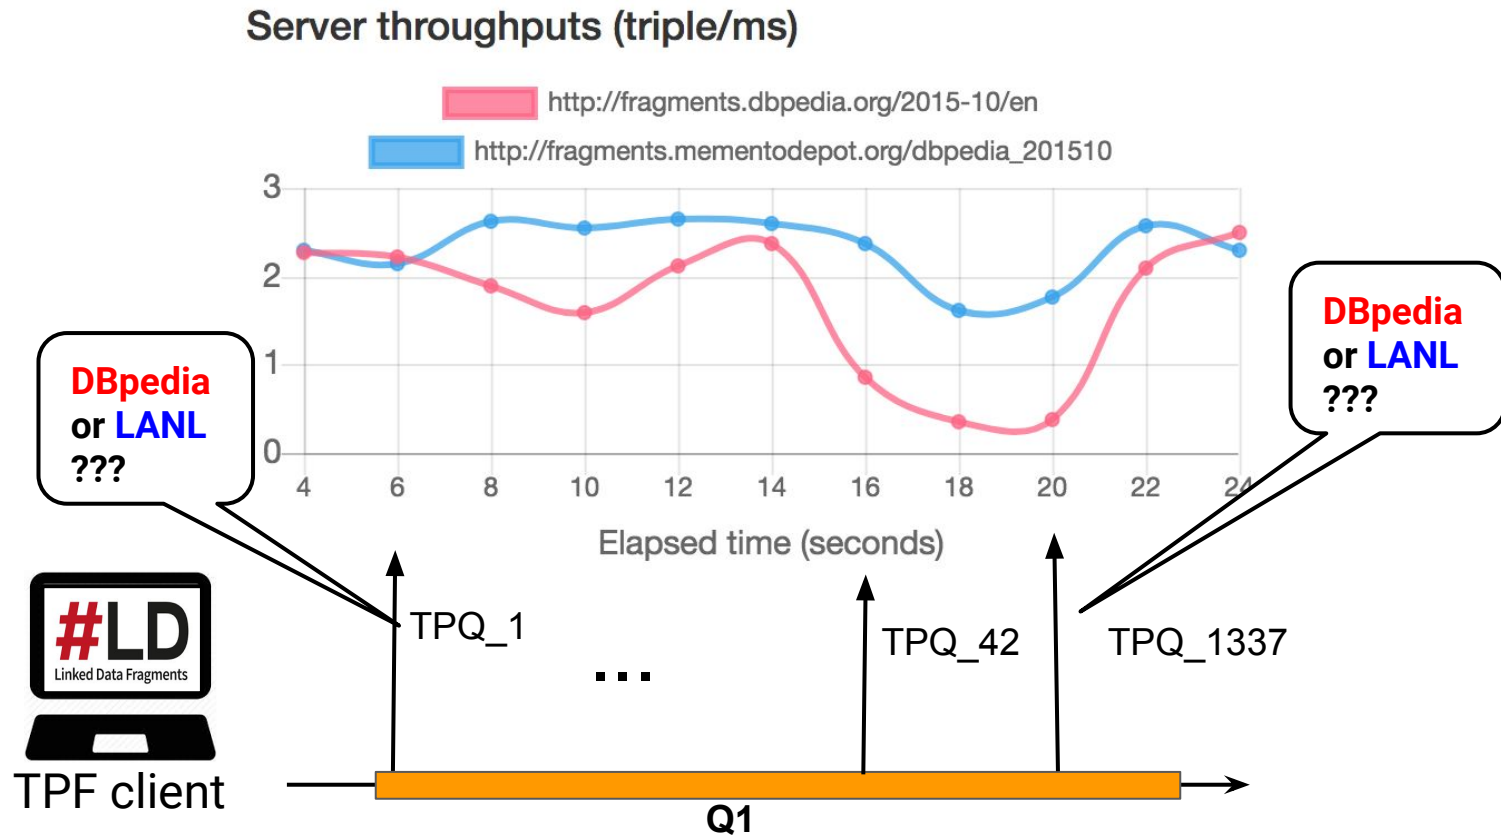


Servers throughputs change over time

Server throughputs (triple/ms)



Where to send Triple Pattern Queries?



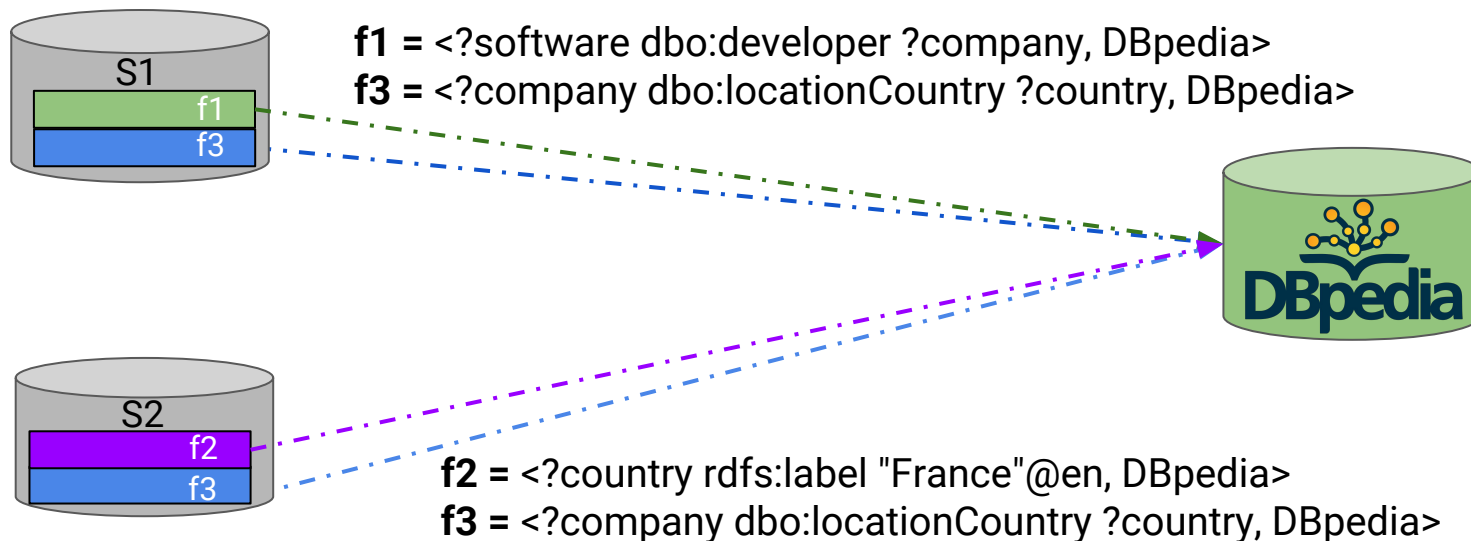
Ulysses: a replication-aware intelligent TPF client

- A **replication-aware source selection**
 - Total/partial replication
- A **light-weighted cost-model**
 - Heterogeneous TPF servers
- A **client-side load balancer**
 - Distributing SPARQL query evaluation



Partial replication model

Fragments of RDF datasets are replicated [2,6]



[2] Montoya, Gabriela, et al. "Federated Queries Processing with Replicated Fragments." *ISWC 2015*.

[6] Ibáñez, Luis-Daniel, et al. "Col-graph: Towards writable and scalable linked open data." *ISWC 2014*.


Ulysses replication-aware source selection

- Replicated fragments are defined using a **catalog** [2]
- Describes which fragment is hosted on which server
- Ulysses loads the catalog when starting

Fragment	Location
f1 = <?software dbo:developer ?company, DBpedia>	S1
f2 = <?country rdfs:label "France"@en, DBpedia>	S2
f3 = <?company dbo:locationCountry ?country, DBpedia>	S1, S2

How to get server throughput?

Server throughputs (triple/ms)



Computing Server throughput

- A **server throughput** is deduced from its **access time**
 - Triple patterns can be evaluated in approximate **constant time** [7] (with HDT backend)
- During query processing, a TPF client executes many triple pattern queries
 - A lot of free probes!

Definition 5 (Server throughput). *Given a set of TPF servers $S = \{S_1, \dots, S_n\}$, $\Delta = \{\delta_1, \dots, \delta_n\}$ where δ_i is the access time of S_i , and $P = \{p_1, \dots, p_n\}$ where p_i is the number of results served per access to S_i .*

$\forall S_i \in S$, the server throughput w_i of S_i is $w_i = \frac{p_i}{\delta_i}$

Computing Server throughput



Access time $\delta_1 = 100\text{ms}$

Page size $p_1 = 100$ triples



Access time $\delta_2 = 100\text{ms}$

Page size $p_2 = 400$ triples



Access time $\delta_3 = 500\text{ms}$

Page size $p_3 = 400$ triples

Computing Server throughput



Access time $\delta_1 = 100\text{ms}$

Page size $p_1 = 100$ triples

\Rightarrow **Server throughput $\omega_1 = 1$ triples/ms**



Access time $\delta_2 = 100\text{ms}$

Page size $p_2 = 400$ triples

\Rightarrow **Server throughput $\omega_2 = 4$ triples/ms**




Access time $\delta_3 = 500\text{ms}$

Page size $p_3 = 400$ triples

\Rightarrow **Server throughput $\omega_3 = 0.8$ triples/ms**


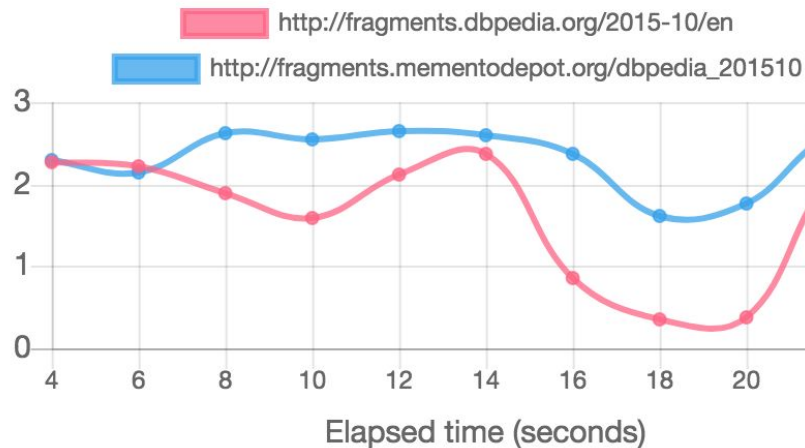
Hard to compare: normalize!

Server throughputs (triple/ms)




Hard to compare: normalize!

Server throughputs (triple/ms)



Servers capability factors



Computing TPF servers capabilities

Definition 6 (Server capability). *Given a set of TPF servers $S = \{S_1, \dots, S_n\}$ and $W = \{w_1, \dots, w_n\}$ where w_i is the throughput of S_i .*

$$\forall S_i \in S, \text{ the capability } \phi_i \text{ of } S_i \text{ is } \phi_i = \frac{w_i}{\min W}$$

Computing TPF servers capabilities



Access time $\delta_1 = 100\text{ms}$

Server throughput $\omega_1 = 1$ triples/ms

Page size $p_1 = 100$ triples



Access time $\delta_2 = 100\text{ms}$

Server throughput $\omega_2 = 4$ triples/ms

Page size $p_2 = 400$ triples



Access time $\delta_3 = 500\text{ms}$

Server throughput $\omega_3 = 0.8$ triples/ms

Page size $p_3 = 400$ triples

Computing TPF servers capabilities



Access time $\delta_1 = 100\text{ms}$

Page size $p_1 = 100$ triples

Server throughput $\omega_1 = 1$ triples/ms



Capability factor $\phi_1 = 1.25$



Access time $\delta_2 = 100\text{ms}$

Page size $p_2 = 400$ triples

Server throughput $\omega_2 = 4$ triples/ms



Capability factor $\phi_2 = 6.25$



Access time $\delta_3 = 500\text{ms}$

Page size $p_3 = 400$ triples

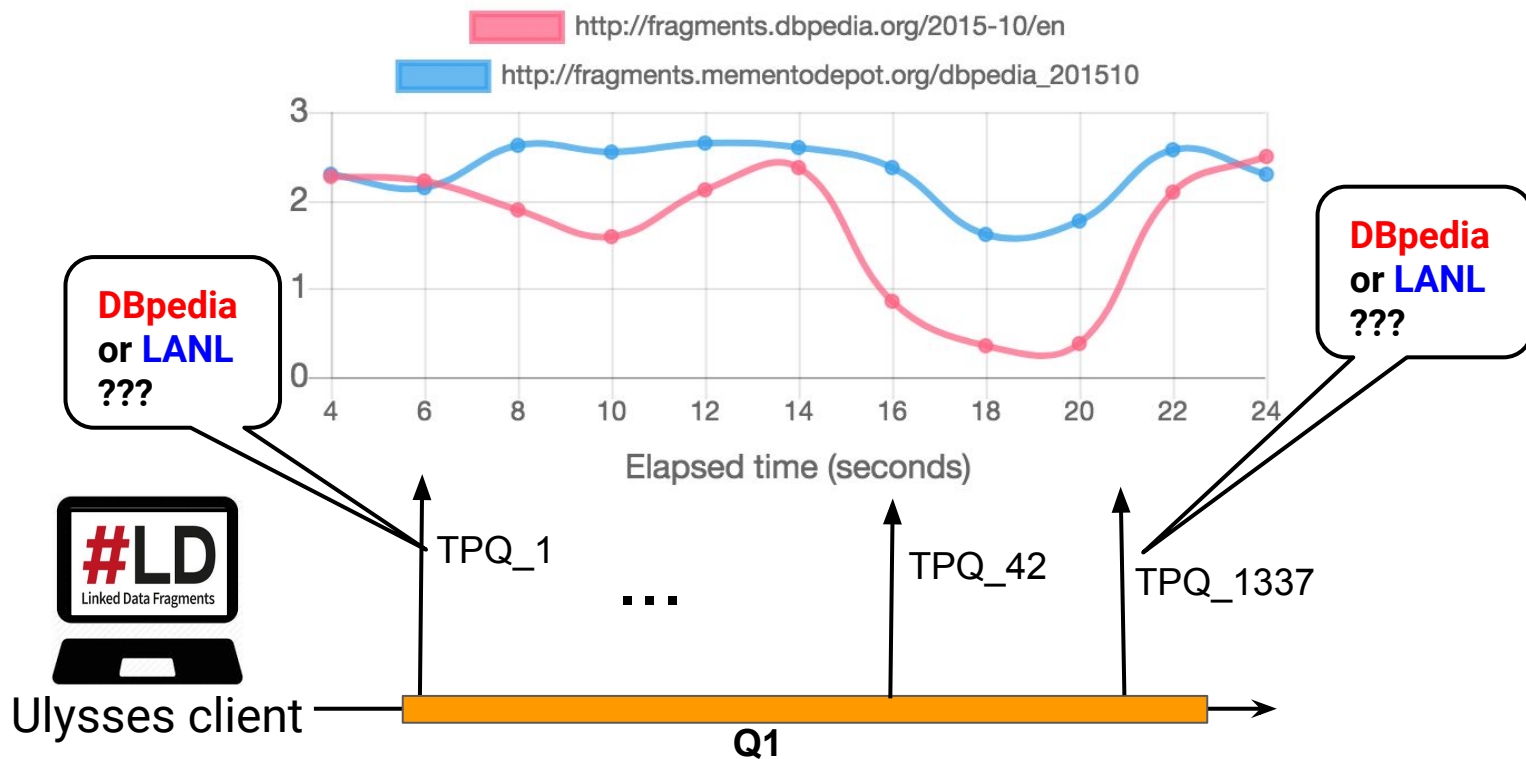
Server throughput $\omega_3 = 0.8$ triples/ms



Capability factor $\phi_3 = 1$

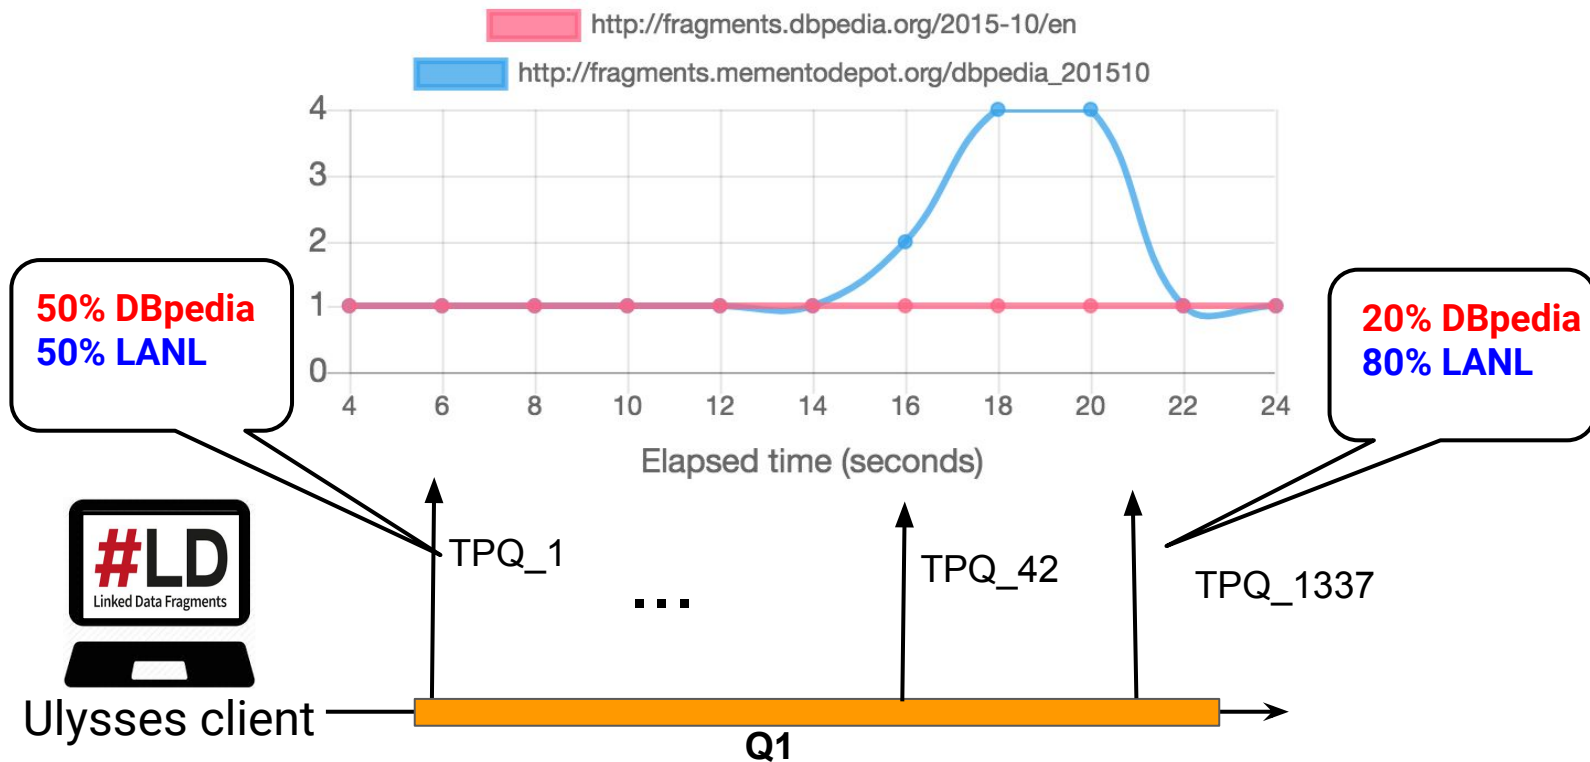
Ulysses in action

Server throughputs (triple/ms)



Ulysses in action

Servers capability factors



Weighted random access of TPF servers

Definition 7 (Weighted random access). *Given a set of TPF servers $S = \{S_1, \dots, S_n\}$ and $\Phi = \{\phi_1, \dots, \phi_n\}$ where ϕ_i is the capability of S_i .*

When selecting a TPF server $S_i \in S$ to evaluate a triple pattern tp , the probability of selecting S_i is: $\mathcal{A}(S_i) = \frac{\phi_i}{\sum_{j=1}^n \phi_j}$, such as: (i) $\sum_{S_i \in S} \mathcal{A}(S_i) = 1$; (ii) $\forall S_i \in S, 0 \leq \mathcal{A}(S_i) \leq 1$.

Weighted random access of TPF servers



Capability factor $\phi_1 = 1.25$



Capability factor $\phi_2 = 6.25$



Capability factor $\phi_3 = 1$

Weighted random access of TPF servers



Capability factor $\phi_1 = 1.25 \Rightarrow$ **Access probability $A_1 = 14.7\%$**



Capability factor $\phi_2 = 6.25 \Rightarrow$ **Access probability $A_2 = 73.5\%$**



Capability factor $\phi_3 = 1 \Rightarrow$ **Access probability $A_3 = 11.7\%$**

Experimental Study


Experimental setup

- **Dataset:** Waterloo SPARQL Diversity Test Suite [8] (WatDiv) synthetic dataset with 10^7 triples
- **Queries:** 100 random WatDiv queries (STAR, PATH and SNOWFLAKE shaped SPARQL queries)
- **Replication configurations:**
 - ***Total replication:*** each server replicates the whole dataset
 - ***Partial replication:*** fragments are created from the 100 random queries and are replicated up to two times.

Experimental setup


- **Servers:** hosted on Amazon EC2 cloud using *t2.micro* instances
- **Network configurations:**
 - HTTP proxies are used to simulate network latencies and special conditions
 - ***Homogeneous:*** all servers have access latencies of 300ms.
 - ***Heterogeneous:*** The first server has an access latency of 900ms, and other servers have access latencies of 300ms.

Ulysses balance the load according to servers processing capabilities





Homogeneous servers and total replication

Ulysses balance the load according to servers processing capabilities




Heterogeneous servers and total replication

Ulysses balance the load according to servers processing capabilities




Homogeneous servers and partial replication

Ulysses improves query execution time under the load



Ulysses tolerates server failures




S1, S2, S3 **homogeneous**: S1 fails at 5s and S3 fails at 20s

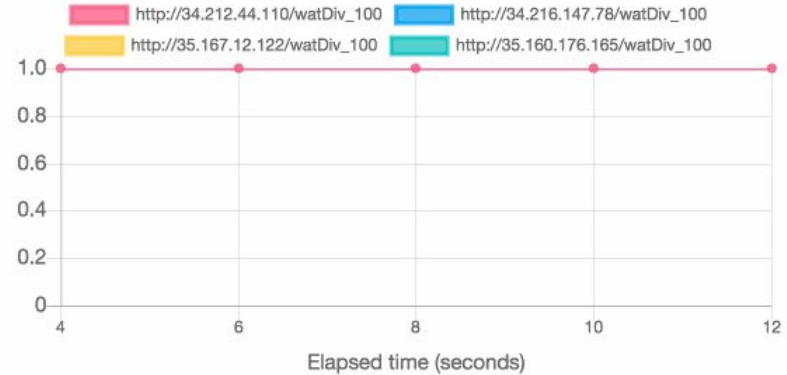
Ulysses in real-life

<http://ulysses-demo.herokuapp.com>


Number of HTTP calls per server



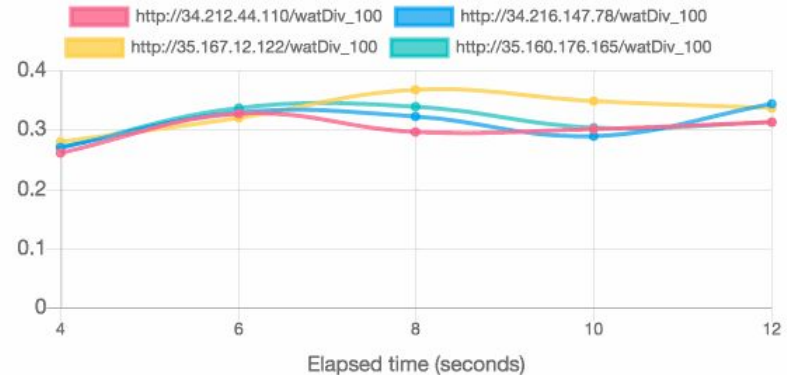
Servers capability factors



Servers access times (ms)



Server throughputs (triple/ms)




Conclusion

- How to **balance the load of SPARQL query processing** over replicated **heterogeneous** servers owned by **autonomous** data providers?
 - Using a **client-side load-balancer** based on Ulysses **cost-model**
 - Require **no changes** from data providers!



Future Works

- **How to build the catalog** of replicated fragments?
 - Provided by data providers as **metadata**
 - A central **index** of replicated RDF datasets
- Consider **divergence** over replicated data
 - Load-balance only if datasets are *k-atomic* [9] or *delta-consistent* [10]



[9] A. Aiyer et al. “On the availability of non strict quorum systems”. In Proceedings of the 19th International Symposium on Distributed Computing (DISC) (2005)

[10] Cao, J. et al. “Data consistency for cooperative caching in mobile environments.” Computer (2007)

Intelligent Clients for Replicated Triple Pattern Fragments

Come to see the demo tomorrow! (290)

<http://ulysses-demo.herokuapp.com>

ESWC 2018 - Heraklion, Greece

June 6th, 2018

