Optimizing the computation of overriding

P.A. Bonatti, I.M. Petrova, and L. Sauro

Dip. Ing. Elet. e Tecnologie dell'Informazione, Università di Napoli Federico II

Abstract. We introduce optimization techniques for reasoning in \mathcal{DL}^N —a recently introduced family of nonmonotonic description logics whose characterizing features appear well-suited to model the applicative examples naturally arising in biomedical domains and semantic web access control policies. Such optimizations are validated experimentally on large KBs with more than 30K axioms. Speedups exceed 1 order of magnitude. For the first time, response times compatible with real-time reasoning are obtained with nonmonotonic KBs of this size.

1 Introduction

Recently, a new family of nonmonotonic Description Logics (DLs), called \mathcal{DL}^N , has been introduced [7]. It supports normality concepts NC to denote the normal/standard/prototypical instances of a concept C, and prioritized defeasible inclusions (DIs) $C \sqsubseteq_n D$ with the following meaning: "by default, the instances of C satisfy D, unless stated otherwise", that is, unless some higher priority axioms entail $C \sqcap \neg D$; in that case, $C \sqsubseteq_n D$ is overridden. The normal/standard/prototypical instances of C are required to satisfy all the DIs that are not overridden in C.

Given the negligible number of applications based on nonmonotonic logics deployed so far, \mathcal{DL}^N has been designed to address real-world problems and concrete knowledge engineering needs. In this regard, the literature provides clear and articulated discussions of how nonmonotonic reasoning can be of help in important contexts related to the semantic web, such as biomedical ontologies [24,27] (with several applications, such as literature search) and (semantic web) policy formulation [28]. These and other applications are extensively discussed in [7].

The distinguishing features in \mathcal{DL}^N 's design are: (i) \mathcal{DL}^N adopts the simplest possible criterion for overriding, that is, inconsistency with higher priority axioms; (ii) all the normal instances of a concept C conform to the same set of default properties, sometimes called *prototype* in the following; (iii) the conflicts between DIs that cannot be resolved with priorities are regarded as knowledge representation errors and are to be fixed by the knowledge engineer (typically, by adding specific DIs). No traditional non-monotonic logic satisfies (i), and very few satisfy (ii) or (iii). \mathcal{DL}^N behaves very well on applicative examples due to the following consequences of (i)–(iii) (a comparison with other nonmonotonic DLs with respect to these features is summarized in Table 1):

No inheritance blocking: In several nonmonotonic logics a concept with exceptional properties inherits *none* of the default properties of its superclasses. This undesirable phenomenon is known as *inheritance blocking*.

No undesired closed-world assumption (CWA) effects: In some nonmonotonic DLs, an exceptional concept is shrinked to the individuals that explicitly belong to it, if any; hence, it may become inconsistent.

Control on role ranges: Unlike most nonmonotonic DLs, \mathcal{DL}^N axioms can specify whether a role should range only over normal individuals or not.

Detect inconsistent prototypes: \mathcal{DL}^{N} facilitates the identification of all conflicts that cannot be resolved with priorities (via consistency checks over normality concepts), because their correct resolution is application dependent and should require human intervention (cf. [7, Sec. 1] and Example 1 below).

Tractability: \mathcal{DL}^N is currently the only nonmonotonic DL known to preserve the tractability of all low-complexity DLs, including \mathcal{EL}^{++} and DL-lite (that underly the OWL2-EL and OWL2-QL profiles). This opens the way to processing very large nonmonotonic KBs within these fragments.

	CIRC				RA		PR	
Features	[6,5]	[1,2]	[12]	[17,16]	[10,9]	[11]	[21]	$\mathcal{DL}^{\mathrm{N}}$
no inheritance blocking	✓	√	√			√	✓	✓
no CWA effects		√	√		√	√		√
fine-grained control on role ranges				sometimes				√
detects inconsistent prototypes				sometimes			√	√
preserves tractability								√ (*)

^(*) It holds for subsumption, assertion checking, concept consistency, KB consistency.

Table 1. Partial comparison with other nonmonotonic DLs, cf.[7], where CIRC, DEF, AEL, TYP, RAT, PR stand, respectively, for Circumscribed DLs, Default DLs, Autoepistemic DLs, DLs with Typicality, DLs with Rational Closure, and Probabilistic DLs.

The performance of \mathcal{DL}^N inference has been experimentally analyzed on large KBs (with more than 20K concept names and over 30K inclusions). The results are promising; still, as defeasible inclusions approach 25% of the KB, query response time slows down enough to call for improvements. In this paper, we study two optimization techniques to improve \mathcal{DL}^N query response time:

- 1. Many of the axioms in a large KB are expected to be irrelevant to the given query. We investigate the use of *module extractors* [26,23] to focus reasoning on relevant axioms only. The approach is not trivial (module extractors are unsound for most nonmonotonic logics, including circumscription, default and autoepistemic logics) and requires an articulated correctness proof.
- 2. We introduce a new algorithm for query answering, that is expected to exploit incremental reasoners at their best. Incremental reasoning is crucial as \mathcal{DL}^N 's reasoning method iterates consistency tests on a set of KBs with large intersections. While the assertion of new axioms is processed very efficiently, the computational cost of axiom deletion is generally not negligible. We introduce an *optimistic reasoning method* that is expected to reduce the number of deletions.

Both optimizations are validated experimentally. Speedups exceed 1 order of magnitude. To the best of our knowledge, this is the first time that response times compatible with real-time reasoning are obtained with nonmonotonic KBs of this size.

The paper is organized as follows: Sec. 2 provides the basics of \mathcal{DL}^{N} and illustrates its inferences with examples. Sections 3 and 4 introduce the two optimization methods,

respectively, and prove their correctness. Their experimental assessment is in Sec. 5. We assume the reader to be familiar with description logics, see [14] for all details. The code and test suites are available at: http://goo.gl/KnMO91.

2 Preliminaries

Let \mathcal{DL} be any classical description logic language (see [14] for definitions), and let \mathcal{DL}^{N} be the extension of \mathcal{DL} with a new concept name NC for each \mathcal{DL} concept C. The new concepts are called *normality concepts*.

A \mathcal{DL}^{N} knowledge base is a disjoint union $\mathcal{KB} = \mathcal{S} \cup \mathcal{D}$ where \mathcal{S} is a finite set of \mathcal{DL}^{N} inclusions and assertions (called *strong* or classical axioms) and \mathcal{D} is a finite set of *defeasible inclusions* (DIs, for short) that are expressions $C \sqsubseteq_{n} D$ where C is a \mathcal{DL} concept and D a \mathcal{DL}^{N} concept. If $\delta = (C \sqsubseteq_{n} D)$, then $\operatorname{pre}(\delta)$ and $\operatorname{con}(\delta)$ denote C and D, respectively. Informally speaking, the set of DIs satisfied by all the instances of a normality concept NC constitute the *prototype* associated to C.

DIs are prioritized by a strict partial order \prec . If $\delta_1 \prec \delta_2$, then δ_1 has higher priority than δ_2 . \mathcal{DL}^N solves automatically only the conflicts that can be settled using \prec ; any other conflict shall be resolved by the knowledge engineer (typically by adding suitable DIs). Two priority relations have been investigated so far. Both are based on *specificity*: the specific default properties of a concept C have higher priority than the more generic properties of its superconcepts (i.e. those that subsume C). The priority relation used in most of [7]'s examples identifies those superconcepts with strong axioms only:

$$\delta_1 \prec \delta_2 \text{ iff } \mathsf{pre}(\delta_1) \sqsubseteq_{\mathcal{S}} \mathsf{pre}(\delta_2) \text{ and } \mathsf{pre}(\delta_2) \not\sqsubseteq_{\mathcal{S}} \mathsf{pre}(\delta_1).^1$$
 (1)

The second priority relation investigated in [7] is

$$\delta_1 \prec \delta_2 \text{ iff } rank(\delta_1) > rank(\delta_2),$$
 (2)

where $rank(\cdot)$ is shown in Algorithm 1 and corresponds to the ranking function of rational closure [10,11]. This relation uses also DIs to determine superconcepts, so (roughly speaking) a DI $C \sqsubseteq_n D$ —besides defining a default property for C—gives the specific default properties of C higher priority than those of D. The advantage of this priority relation is that it resolves more conflicts than (1); the main advantage of (1) is predictability; e.g. the effects of adding default properties to an existing, classical KB are more predictable, as the hierarchy used for determining specificity and resolving conflicts is the original, validated one, and is not affected by the new DIs (see also the related discussion in [3,4], that adopt (1)).

The expression $\mathcal{KB} \bowtie \alpha$ means that α is a \mathcal{DL}^N consequence of \mathcal{KB} . Due to space limitations, we do not report the model-theoretic definition of \bowtie and present only its reduction to classical reasoning [7]. For all subsumptions and assertions α , $\mathcal{KB} \bowtie \alpha$ holds iff $\mathcal{KB}^{\Sigma} \models \alpha$, where Σ is the set of normality concepts that explicitly occur in $\mathcal{KB} \cup \{\alpha\}$, and \mathcal{KB}^{Σ} is a classical knowledge base obtained as follows (recall that $\mathcal{KB} = \mathcal{S} \cup \mathcal{D}$):

¹ As usual, $C \sqsubseteq_{\mathcal{S}} D$ means that $\mathcal{S} \models C \sqsubseteq D$.

Algorithm 1: Ranking function

```
Input: Ontology \mathcal{KB} = \mathcal{S} \cup \mathcal{D}
      Output: the function rank(\cdot)
\mathbf{1} \ i := -1; \ \mathcal{E}_0 := \{ C \sqsubseteq D \mid C \sqsubseteq_n D \in \mathcal{D} \}
2 repeat
                \begin{array}{l} \mathcal{E}_{i+1} := \{ C \sqsubseteq D \in \mathcal{E}_i \mid \mathcal{S} \cup \mathcal{E}_i \models C \sqsubseteq \bot \} \\ \text{forall } C \sqsubseteq_n D \text{ s.t. } C \sqsubseteq D \in \mathcal{E}_i \setminus \mathcal{E}_{i+1} \text{ do} \\ & \sqsubseteq \operatorname{assign} \operatorname{rank}(C \sqsubseteq_n D) := i \end{array}
7 until \mathcal{E}_{i+1} = \mathcal{E}_i
8 forall C \sqsubseteq_n D \in \mathcal{E}_{i+1} do assign rank(C \sqsubseteq_n D) := \infty
9 return rank(\cdot)
```

First, for all DIs $\delta \in \mathcal{D}$ and all $NC \in \Sigma$, let:

$$\delta^{\mathrm{N}C} = \left(\mathrm{N}C \sqcap \mathsf{pre}(\delta) \sqsubseteq \mathsf{con}(\delta) \right). \tag{3}$$

The informal meaning of δ^{NC} is: "NC's instances satisfy δ ".

Second, let $S' \downarrow_{\prec \delta}$ denote the result of removing from the axiom set S' all the δ_0^{NC} such that $\delta_0 \not\prec \delta$:

$$\mathcal{S}' \downarrow_{\prec \delta} = \mathcal{S}' \setminus \{ \delta_0^{NC} \mid NC \in \Sigma \land \delta_0 \not\prec \delta \}.$$

Third, let $\delta_1, \ldots, \delta_{|\mathcal{D}|}$ be any *linearization* of (\mathcal{D}, \prec) .² Finally, let $\mathcal{KB}^{\Sigma} = \mathcal{KB}^{\Sigma}_{|\mathcal{D}|}$, where the sequence \mathcal{KB}^{Σ}_i $(i = 1, 2, \ldots, |\mathcal{D}|)$ is inductively defined as follows:

$$\begin{split} \mathcal{KB}_{0}^{\Sigma} &= \mathcal{S} \cup \left\{ \mathrm{N}C \sqsubseteq C \mid \mathrm{N}C \in \Sigma \right\} \\ \mathcal{KB}_{i}^{\Sigma} &= \mathcal{KB}_{i-1}^{\Sigma} \cup \left\{ \delta_{i}^{\mathrm{N}C} \mid \delta_{i} \in \mathcal{KB}, \mathrm{N}C \in \Sigma, \text{ and} \right. \\ \mathcal{KB}_{i-1}^{\Sigma} \downarrow_{\prec \delta_{i}} \cup \left\{ \delta_{i}^{\mathrm{N}C} \right\} \not\models \mathrm{N}C \sqsubseteq \bot \right\}. \end{split} \tag{5}$$

In other words, the above sequence starts with KB's strong axioms extended with the inclusions NC $\sqsubseteq C$, then processes the DIs δ_i in non-increasing priority order. If δ_i can be consistently added to C's prototype, given all higher priority DIs selected so far (which is verified by checking that NC $\not\sqsubseteq \bot$ in line (5)), then its translation $\delta_i^{\rm NC}$ is included in \mathcal{KB}^{Σ} (i.e. δ_i enters C's prototype), otherwise δ_i is discarded, and we say that δ_i is overridden in NC.

Examples 2.1

We start with a brief discussion of \mathcal{DL}^{N} 's conflict handling. Most other logics silently neutralize the conflicts between nonmonotonic axioms with the same (or incomparable) priorities by computing the inferences that are invariant across all possible ways

That is, $\{\delta_1, \dots, \delta_{|\mathcal{D}|}\} = \mathcal{D}$ and for all $i, j = 1, \dots, |\mathcal{D}|$, if $\delta_i \prec \delta_j$ then i < j.

of resolving the conflict. A knowledge engineer might solve it in favor of *some* of its possible resolutions, instead; however, if the logic silently neutralizes the conflict, then missing knowledge may remain undetected and unfixed. This approach may cause serious problems in the policy domain:

Example 1. Suppose that project coordinators are both administrative staff and research staff. By default, administrative staff are allowed to sign payments, while research staff are not. A conflict arises since both of these default policies apply to project coordinators. Formally, \mathcal{KB} can be formalized with:

Admin
$$\sqsubseteq_n \exists has_right.Sign$$
 (6)
Research $\sqsubseteq_n \neg \exists has_right.Sign$ (7)

PrjCrd \sqsubseteq Admin \sqcap Research (8)

Leaving the conflict unresolved may cause a variety of security problems. If project coordinators should not sign payments, and the default policy is open (authorizations are granted by default), then failing to infer $\neg\exists \text{has_right.Sign}$ would improperly authorize the signing operation. Conversely, if the authorization is to be granted, then failing to prove $\exists \text{has_right.Sign}$ causes a denial of service (the user is unable to complete a legal operation). To prevent these problems, \mathcal{DL}^{N} makes the conflict visible by inferring $\mathcal{KB} \models N \text{PrjCrd} \sqsubseteq \bot$ (showing that PrjCrd's prototype is inconsistent). This can be proved by checking that $\mathcal{KB}^{\varSigma} \models N \text{PrjCrd} \sqsubseteq \bot$, where $\varSigma = \{N \text{PrjCrd}\}$. Then \mathcal{KB}^{\varSigma} consists of (8), $N \text{PrjCrd} \sqsubseteq \text{PrjCrd}$, and the translation of (6) and (7) (none overrides the other because none is more specific under any of the two priorities):

$$\label{eq:normalized_normalized} $$ N \PrjCrd \cap Admin \sqsubseteq \exists has_right.Sign, \\ N \PrjCrd \cap Research \sqsubseteq \neg \exists has_right.Sign. $$$ $$ $$$$

Here is another applicative example from the semantic policy domain, showing \mathcal{DL}^N 's behavior on multiple exception levels.

Example 2. We are going to axiomatize the following natural language policy: "In general, users cannot access confidential files; Staff can read confidential files; Blacklisted users are not granted any access. This directive cannot be overridden." Note that each of the above directives contradicts (and is supposed to override) its predecessor in some particular case. Authorizations can be reified as objects with attributes subject (the access requestor), target (the file to be accessed), and privilege (such as read and write). Then the above policy can be encoded as follows:

Staff
$$\sqsubseteq$$
 User (9)

Blklst
$$\Box$$
 Staff (10)

UserReqst
$$\sqsubseteq_n \neg \exists privilege$$
 (11)

$$StaffReqst \sqsubseteq_n \exists privilege.Read$$
 (12)

where BlkReq $\doteq \exists$ subj.Blklst, StaffReqst $\doteq \exists$ subj.Staff, and UserReqst $\doteq \exists$ subj.User. By (9), both the specifity relations (1) and (2) yield (12) \prec (11), that is,

(12) has higher priority than (11). Let $\Sigma = \{\text{NStaffReqst}\}; (12) \text{ overrides (11) in NStaffReqst (under (1) as well as (2)), so } \mathcal{KB}^{\Sigma} \text{ consists of: (9), (10), (13), plus}$

$$NS \texttt{taffReqst} \sqsubseteq S \texttt{taffReqst}$$

$$NS \texttt{taffReqst} \sqsubseteq \exists \texttt{privilege.Read}$$

Consequently, $\mathcal{KB} \approx \mathtt{NStaffReqst} \sqsubseteq \exists \mathtt{privilege.Read}.$ Similarly, it can be verified that:

1. Normally, access requests involving confidential files are rejected, if they come from generic users: KB № NUserReqst ⊑ ¬∃privilege;

2. Blacklisted users cannot do anything by (13), so, in particular: $\mathcal{KB} \approx \text{NBlkReq} \sqsubseteq \neg \exists \text{privilege}.$

Some applicative examples from the biomedical domain can be found in [7] (see Examples 3, 4, 10, 12, and the drug contraindication example in Appendix C). Like the above examples, they are all correctly solved by \mathcal{DL}^N with both priority notions. Applicative examples hardly exhibit the complicated networks of dependencies between conflicting defaults that occur in artificial examples. Nonetheless, we briefly discuss the artificial examples, too, as a mean of comparing \mathcal{DL}^N with other logics such as [25,11,5].

In several cases, e.g. examples B.4 and B.5 in [25], \mathcal{DL}^{N} agrees with [25,11,5] under both priority relations. Due to space limitations, we illustrate only B.4.

Example 3 (Juvenile offender). Let \mathcal{KB} consist of axioms (14)-(18) where J, G, M, P abbreviate JuvenileOffender, GuiltyOfCrime, IsMinor and ToBePunished, respectively.

${\tt J} \;\sqsubseteq\; {\tt G}$	(14)	$\mathtt{J}\sqsubseteq\mathtt{G}$	(19)
$J \;\sqsubseteq\; M$	(15)	$\mathtt{J}\sqsubseteq \mathtt{M}$	(20)
$\mathtt{M} \sqcap \mathtt{G} \sqsubseteq_n \neg \mathtt{P}$	(16)	$NJ \sqsubseteq J$	(21)
$\mathtt{M} \sqsubseteq_n \neg \mathtt{P}$	(17)	$\mathrm{N}\mathtt{J} \sqcap \mathtt{M} \sqcap \mathtt{G} \sqsubseteq \neg \mathtt{P}$	(22)
$\mathtt{G} \sqsubseteq_n \mathtt{P}$	(18)	$ ext{NJ} \sqcap ext{M} \sqsubseteq \neg ext{P}$	(23)

On one hand, criminals have to be punished and, on the other hand, minors cannot be punished. So, what about juvenile offenders? The defeasible inclusion (16) breaks the tie in favor of their being underage, hence not punishable. By setting $\Sigma = \{NJ\}$, priorities (1) and (2) both return axioms (19)–(23) as \mathcal{KB}^{Σ} . Then, clearly, $\mathcal{KB}^{\Sigma} \models NJ \sqsubseteq \neg P$ which is \mathcal{DL}^{N} 's analogue of the inferences of [25,11,5].

In other cases (e.g. example B.1 in [25]) \mathcal{DL}^N finds the same conflicts as [25,11,5]. However, \mathcal{DL}^N 's semantics signals these conflicts to the knowledge engineer whereas in [25,11,5] they are silently neutralized.

Example 4 (Double Diamond). Let KB be the following set of axioms:

$\mathtt{A} \sqsubseteq_n \mathtt{T}$	(24)	a = _ p	(20)
$A \sqsubseteq_n P$	(25)	$\mathtt{S} \sqsubseteq_n \mathtt{R}$	(28)
		$\mathtt{P} \sqsubseteq_n \mathtt{Q}$	(29)
$\mathtt{T}\sqsubseteq_n \mathtt{S}$	(26)	$\mathtt{Q} \sqsubseteq_n \neg \mathtt{R}$	(30)
$\mathtt{P} \sqsubseteq_n \lnot \mathtt{S}$	(27)	$\mathbf{q} \equiv n$	(50)

DIs (26) and (27) have incomparable priority under (1) and (2). Consequently, it is easy to see that NA \sqsubseteq S and NA \sqsubseteq ¬S are both implied by \mathcal{KB}^{Σ} and hence the knowledge engineer is warned that NA is inconsistent. The same conflict is silently neutralized in [11,25,5] (A's instances are subsumed by neither S nor ¬S and no inconsistency arises). Similarly for the incomparable DIs (28) and (30) and the related conflict.

The third category of examples (e.g. B.2 and B.3 in [25]) presents a more variegated behavior. In particular, [11] and \mathcal{DL}^N with priority (2) solve all conflicts and infer the same consequences; [25] solves only some conflicts; [5] is not able to solve any conflict and yet it does not raise any inconsistency warning; \mathcal{DL}^N with priority (1) cannot solve the conflicts but raises an inconsistency warning. Here, for the sake of simplicity, we discuss in detail a shorter example which has all relevant ingredients.

Example 5. Let KB be the following defeasible knowledge base:

$$A \sqsubseteq_n B$$
 (31) $A \sqsubseteq_n C$ (32) $B \sqsubseteq_n \neg C$ (33)

According to priority (1) all DIs are incomparable. Therefore, \mathcal{DL}^N warns (by inferring NA $\sqsubseteq \bot$) that the conflict between NA $\sqsubseteq C$ and NA $\sqsubseteq \neg C$ cannot be solved. Note that [5] adopts priority (1), too, however according to circumscription, any interpretation where A's instances are either in $\neg C \sqcap B$ or in C is a model, so A is satisfiable (the conflict is silently neutralized). Under priority (2), instead, axiom (31) gives (31) and (32) higher priority than (33). Consequently, NA $\sqsubseteq C$ prevails over NA $\sqsubseteq \neg C$. In this case, \mathcal{DL}^N and rational closure infer the same consequences.

3 Relevance and modularity

The naive construction of \mathcal{KB}^{Σ} must process all the axioms in $\mathcal{KB}^{\Sigma}_{all} = \mathcal{KB}^{\Sigma}_{0} \cup \{\delta^{NC} \mid \delta \in \mathcal{D}, \ NC \in \Sigma\}$. Here we optimize \mathcal{DL}^{N} inference by quickly discarding some of the irrelevant axioms in $\mathcal{KB}^{\Sigma}_{all}$ using modularization techniques.

Roughly speaking, the problem of module extraction can be expressed as follows: given a reference vocabulary Sig, a module is a (possibly minimal) subset $\mathcal{M} \subseteq \mathcal{KB}$ that is relevant for Sig in the sense that it preserves the consequences of \mathcal{KB} that contain only terms in Sig.

The interest in module extraction techniques is motivated by several ontology engineering needs. We are interested in modularization as an optimization technique for querying large ontologies: the query is evaluated on a (hopefully much smaller) module of the ontology that preserves the query result (as well as any inference whose signature is contained in the query's signature).

However, the problem of deciding whether two knowledge bases entail the same axioms over a given signature is usually harder than standard reasoning tasks. Consequently deciding whether \mathcal{KB}' is a module of \mathcal{KB} (for Sig) is computationally expensive in general. For example, $DL-Lite_{horn}$ complexity grows from PTIME to coNP-TIME-complete [20]; for \mathcal{ALC} , complexity is one exponential harder [15], while for \mathcal{ALCQIO} the problem becomes even undecidable [22].

In order to achieve a practical solution, a syntactic approximation has been adopted in [26,18]. The corrisponding algorithm $\top \bot^*$ -Mod (Sig, \mathcal{KB}) is defined in [26, Def. 4]

and reported in Algorithm 2 below. It is based on the property of \bot -locality and \top -locality of single axioms (line 2). An axiom is local w.r.t. Sig if the substitution of all non-Sig terms with \bot (resp. \top) turns it into a tautology.

The module extractor identifies a subset $\mathcal{M} \subseteq \mathcal{KB}$ of the knowledge base and a signature Sig (containing all symbols of interest) such that all axioms in $\mathcal{KB} \setminus \mathcal{M}$ are local w.r.t. Sig. This guarantees that every model of \mathcal{M} can be extended to a model of \mathcal{KB} by setting each non-Sig term to either \bot or \top . In turn, this property guarantees that any query whose signature is contained in Sig has the same answer in \mathcal{M} and \mathcal{KB} .

The function $x\operatorname{-Mod}(Sig,\mathcal{KB})$ (lines 9-19), where x stands for \top or \bot , describes the procedure for constructing modules of a knowledge base \mathcal{KB} for each notion of locality. Starting with an empty set of axioms (line 11), iteratively, the axioms α that are non-local are added to the module (line 2) and, in order to preserve soundness, the signature against which locality is checked is extended with the terms in α (line 2). Iteration stops when a fixpoint is reached.

Modules based on a single syntactic locality can be further shrinked by iteratively nesting \top -extraction into \bot -extraction, thus obtaining $\top \bot^*$ -Mod (Sig, \mathcal{KB}) modules (lines 2-2).

The notions of module and locality must be extended to handle DIs, before we can apply them to $\mathcal{DL}^{\mathbb{N}}$. Definition 1 generalizes the substitutions operated by the module extraction algorithm, abstracting away procedural details. As in [26], both \widetilde{X} and $\operatorname{sig}(X)$ denote the signature of X.

Definition 1. (Module, locality) $A \perp \bot^*$ -substitution for KB and a signature Sig is a substitution σ over $KB \setminus Sig$ that maps each concept name on \top or \bot , and every role name on the universal role or the empty role. A strong axiom α is σ -local iff $\sigma(\alpha)$ is a tautology. A DI $C \sqsubseteq_n D$ is σ -local iff $C \sqsubseteq D$ is σ -local. A set of axioms is σ -local if all of its members are. We say that an axiom α is \top -local (resp. \bot -local) if α is σ -local where the substitution σ uniformly maps concept names to \top (resp. \bot).

A (syntactic) module of KB with respect to Sig is a set $\mathcal{M} \subseteq KB$ such that $KB \setminus \mathcal{M}$ is σ -local for some $\top \bot^*$ -substitution σ for KB and $\widetilde{\mathcal{M}} \cup Sig$.

Let $\mathrm{Mod}_{\mathrm{DI}}(Sig,\mathcal{KB})$ be the variant of $\top\bot^*\mathrm{-Mod}(Sig,\mathcal{KB})$ where the test in line 2 is replaced with (the complement of) the \top or $\bot\mathrm{-locality}$ condition of Def. 1 (that covers DIs, too). Using the original correctness argument for $\top\bot^*\mathrm{-Mod}(Sig,\mathcal{KB})$ cf. [18, Prop. 42], it is easy to see that $\mathrm{Mod}_{\mathrm{DI}}(Sig,\mathcal{KB})$ returns a syntactic module of \mathcal{KB} w.r.t. Sig according to Def. 1. If \mathcal{KB} contains no DIs (i.e. it is classical), then Def. 1 is essentially a rephrasing of standard syntactic notions of modules and locality, 3 so

for all queries
$$\alpha$$
 such that $\widetilde{\alpha} \subseteq Sig$, $\mathcal{M} \models \alpha$ iff $\mathcal{KB} \models \alpha$. (34)

However, proving that $\top \bot^*$ -Mod_{DI}(Sig, \mathcal{KB}) is correct for *full* \mathcal{DL}^N is far from obvious: removing axioms from \mathcal{KB} using module extractors is incorrect under most nonmonotonic semantics (including circumscription, default logic and autoepistemic logic). The reason is that nonmonotonic inferences are more powerful than classical

² For efficiency, this test is approximated by a matching with a small set of templates.

³ Informally, $\top \perp^*$ -Mod's greedy strategy tends to find small Def. 1's modules.

Algorithm 2: $\top \bot^*$ -Mod (Sig, \mathcal{KB})

```
Input: Ontology KB, signature Siq
     Output: \top \perp^*-module \mathcal{M} of \mathcal{KB} w.r.t. Sig
     // main
 1 begin
           \mathcal{M} := \mathcal{KB}
 2
           repeat
 3
                  \mathcal{M}' := \mathcal{M}
 4
                  \mathcal{M} := \top \text{-Mod}(\bot \text{-Mod}(\mathcal{M}, Sig), Sig)
 5
           until \mathcal{M} \neq \mathcal{M}'
 6
           return \mathcal{M}
 8 end
 9 function x-Mod(\mathcal{KB}, Sig) // x \in \{\bot, \top\}
10 begin
           \mathcal{M}:=\emptyset, \quad \mathcal{T}:=\mathcal{KB}
11
12
           repeat
13
                  changed = false
                  forall \alpha \in \mathcal{T} do
14
                        if \alpha is not x-local w.r.t. Sig \cup \widetilde{\mathcal{M}} then
15
                               \mathcal{M} := \mathcal{M} \cup \{\alpha\}
16
17
                               \mathcal{T} := \mathcal{T} \setminus \{\alpha\}
                               changed = true
18
           until changed = false
19
20
           return \mathcal{M}
21 end
```

inferences, and the syntactic locality criterions illustrated above fail to capture some of the dependencies between different symbols.

Example 6. Given the knowledge base $\{\top \sqsubseteq A \sqcup B\}$ and $Sig = \{A\}$, the module extractor returns an empty module (because by setting $B = \top$ the only axiom in the KB becomes a tautology). The circumscription of this KB, assuming that both A and B are minimized, does not entail $A \sqsubseteq \bot$, while the circumscription of the empty module entails it.

Example 7. Under the stable model semantics, the normal logic program

$$p \leftarrow \neg p, q$$
$$q \leftarrow \neg r$$
$$r \leftarrow \neg q$$

entails r, both credulously and skeptically. The module extractor, given $Sig = \{r\}$, removes the first rule (that becomes a tautology by setting p =true). The module does not entail r skeptically anymore, and erroneously entails q credulously. Analogues of this example apply to default and autoepistemic logic, using the usual translations. They can be adapted to default DL [2] and the extension of DL based on MKNF [12].

Now we illustrate the correct way of applying $\top \bot^*$ - $\mathrm{Mod}_{\mathrm{DI}}$ to a $\mathcal{DL}^{\mathrm{N}}$ $\mathcal{KB} = \mathcal{S} \cup \mathcal{D}$ and a query α (subsumption or assertion). Let Σ be the union of $\widetilde{\alpha}$ and the set of normality concepts occurring in \mathcal{KB} . Let

$$\mathcal{M}_0 = \mathrm{Mod}_{\mathrm{DI}}(\Sigma, \mathcal{KB} \cup \mathrm{N}\Sigma)$$
,

where $N\Sigma$ abbreviates $\{NC \subseteq C \mid NC \in \Sigma\}$.

Example 8. Let KB be the knowledge base:

$$A \sqsubseteq B$$
 (35) $B \sqcap C \sqsubseteq A$ (37)

$$A \sqsubseteq_n D \sqcap E \tag{36}$$

and α the query $NA \subseteq D$. \mathcal{M}_0 is calculated as follows: first, since no normality concept occurs in \mathcal{KB} , Σ is equal to the signature $\widetilde{\alpha} = \{NA, D\}$.

Algorithm 2 calls first the function \bot -Mod($\mathcal{KB} \cup N\Sigma, \Sigma$). Notice that by replacing C and F with \bot , axioms (37) and (38) becomes a tautology. Consequently, it is easy to see that the returned knowledge base is $\mathcal{KB}' = \{(35), (36), NA \sqsubseteq A\}$.

Then, \top -Mod is called on \mathcal{KB}' and Σ . Now, replacing B with \top makes $A \sqsubseteq B$ a tautology, so the resulting knowledge base is $\mathcal{KB}'' = \{(36), NA \sqsubseteq A\}$. It is easy to see that a fix point is reached and hence \mathcal{KB}'' is returned.

We shall prove that $(\mathcal{KB} \cap \mathcal{M}_0)^{\Sigma}$ can be used in place of \mathcal{KB}^{Σ} to answer query α . This saves the cost of processing $\mathcal{KB}^{\Sigma}_{all} \setminus \mathcal{M}$, where

$$\mathcal{M} = (\mathcal{KB}_0^{\Sigma} \cap \mathcal{M}_0) \cup \{\delta^{NC} \mid \delta \in \mathcal{D} \cap \mathcal{M}_0, \ NC \in \Sigma\}.$$

Note that $\mathcal{KB}^{\Sigma}_{all} \setminus \mathcal{M}$ is usually even larger than $\mathcal{KB} \setminus \mathcal{M}_0$ because for each DI $\delta \notin \mathcal{M}_0$, all its translations δ^{NC} (NC $\in \Sigma$) are removed from \mathcal{M} .

Lemma 1. \mathcal{M} is a module of $\mathcal{KB}_{all}^{\Sigma}$ w.r.t. Σ .

Proof. Since Mod_{DI}(·,·) returns modules, $(\mathcal{KB} \cup N\Sigma) \setminus \mathcal{M}_0$ is σ-local, for some $\top \bot^*$ -substitution σ for $\mathcal{KB} \cup N\Sigma$ and $\widetilde{\mathcal{M}}_0 \cup \Sigma$. So, for all axioms β in $\mathcal{KB}_0^\Sigma \setminus \mathcal{M}$, β is σ-local (as $\mathcal{KB}_0^\Sigma \setminus \mathcal{M} \subseteq (\mathcal{KB} \cup N\Sigma) \setminus \mathcal{M}_0$). Moreover, for all $\beta = \delta^{ND} \in \{\delta^{NC} \mid \delta \in \mathcal{D}, \ NC \in \Sigma\} \setminus \mathcal{M}$, it must hold $\delta \in \mathcal{D} \setminus \mathcal{M}_0$ (by construction of \mathcal{M}), and hence δ is σ-local. Now note that if $\sigma(E \sqsubseteq F)$ is a tautology, then also $\sigma(ND \sqcap E \sqsubseteq F)$ is a tautology, therefore the σ-locality of δ implies the σ-locality of δ^{ND} . It follows that all β in $\mathcal{KB}_{all}^\Sigma \setminus \mathcal{M}$ are σ-local.

Finally, note that $\operatorname{sig}(\mathcal{KB}^{\Sigma}_{all}) = \operatorname{sig}(\mathcal{KB} \cup \operatorname{N}\Sigma)$ and $\widetilde{\mathcal{M}} \cup \Sigma = \widetilde{\mathcal{M}}_0 \cup \Sigma$, therefore σ is also a $\top \bot^*$ -substitution for $\mathcal{KB}^{\Sigma}_{all}$ and $\mathcal{M} \cup \Sigma$. It follows immediately that \mathcal{M} is a module of $\mathcal{KB}^{\Sigma}_{all}$ w.r.t. Σ .

Lemma 2. If \mathcal{M} is a module of \mathcal{KB} w.r.t. a signature Sig and $\mathcal{KB}' \subseteq \mathcal{KB}$, then $\mathcal{KB}' \cap \mathcal{M}$ is a module of \mathcal{KB}' w.r.t. Sig.

Proof. If \mathcal{M} is a module of \mathcal{KB} w.r.t. Sig, then $\mathcal{KB} \setminus \mathcal{M}$ is σ -local for some $\top \bot^*$ -substitution σ for \mathcal{KB} and $\widetilde{\mathcal{M}} \cup Sig$. Let σ' be the restriction of σ to the symbols in

$$\widetilde{\mathcal{KB}}' \setminus (\widetilde{\mathcal{M}} \cup Sig) = \widetilde{\mathcal{KB}}' \setminus (\operatorname{sig}(\mathcal{KB}' \cap \mathcal{M}) \cup Sig).$$

Clearly, σ' is a $\top \bot^*$ -substitution for \mathcal{KB}' and $\widetilde{\mathcal{KB}}' \setminus (\operatorname{sig}(\mathcal{KB}' \cap \mathcal{M}) \cup Sig)$. Moreover, for all $\beta \in \mathcal{KB}' \setminus \mathcal{M}$, $\sigma(\beta) = \sigma'(\beta)$, by construction, so $\mathcal{KB}' \setminus \mathcal{M}$ is σ' -local. Then $\mathcal{KB}' \cap \mathcal{M}$ is a syntactic module of \mathcal{KB}' w.r.t. Sig.

The relationship between $(\mathcal{KB} \cap \mathcal{M}_0)^{\Sigma}$ and \mathcal{KB}^{Σ} is:

Lemma 3.
$$\mathcal{KB}^{\Sigma} \cap \mathcal{M} \subseteq (\mathcal{KB} \cap \mathcal{M}_0)^{\Sigma} \subseteq \mathcal{KB}^{\Sigma}$$
.

Proof. It suffices to prove by induction that for all $i = 0, 1, ..., |\mathcal{D}|,^4$

$$\mathcal{KB}_i^{\Sigma} \cap \mathcal{M} \subseteq (\mathcal{KB} \cap \mathcal{M}_0)_i^{\Sigma} \subseteq \mathcal{KB}_i^{\Sigma}$$
.

Base case (i = 0):

$$\frac{\mathcal{K}\mathcal{B}_{0}^{\Sigma}\cap\mathcal{M}}{\mathcal{E}(\mathcal{S}\cap\mathcal{M}_{0})\cup\mathcal{N}\mathcal{D}} = (\mathcal{S}\cap\mathcal{M}_{0})\cup\mathcal{N}\mathcal{D} = (\mathcal{S}\cap\mathcal{M}_{0})\cup\mathcal{N}\mathcal{D} = (\mathcal{S}\cap\mathcal{M}_{0})\cup\mathcal{N}\mathcal{D} = (\mathcal{S}\mathcal{D})\cup\mathcal{N}\mathcal{D} = (\mathcal{S}\mathcal{D})\cup\mathcal{N$$

Induction step (i > 0): By induction hypothesis (IH)

$$\mathcal{KB}_{i-1}^{\Sigma} \cap \mathcal{M} \subseteq (\mathcal{KB} \cap \mathcal{M}_0)_{i-1}^{\Sigma} \subseteq \mathcal{KB}_{i-1}^{\Sigma}$$
.

First suppose that $\delta_i \notin \mathcal{M}_0$ (and hence for all NC, $\delta_i^{\mathrm{NC}} \notin \mathcal{M}$). Then $\mathcal{KB}_i^{\Sigma} \cap \mathcal{M} = \mathcal{KB}_{i-1}^{\Sigma} \cap \mathcal{M}$, $(\mathcal{KB} \cap \mathcal{M}_0)_i^{\Sigma} = (\mathcal{KB} \cap \mathcal{M}_0)_{i-1}^{\Sigma}$, and (by def.) $\mathcal{KB}_{i-1}^{\Sigma} \subseteq \mathcal{KB}_i^{\Sigma}$. The Lemma follows by IH.

Next assume that $\delta_i \in \mathcal{M}_0$ and let NC be any normality concept in Σ . Note that $\delta_i^{NC} \in \mathcal{M}$. By IH,

$$(\mathcal{KB}_{i-1}^{\Sigma}\cap\mathcal{M})\downarrow_{\prec\delta_{i}}\cup\{\delta_{i}^{NC}\}\subseteq(\mathcal{KB}\cap\mathcal{M}_{0})_{i-1}^{\Sigma}\downarrow_{\prec\delta_{i}}\cup\{\delta_{i}^{NC}\}\subseteq\mathcal{KB}_{i-1}^{\Sigma}\downarrow_{\prec\delta_{i}}\cup\{\delta_{i}^{NC}\}\,.$$

The leftmost term equals $(\mathcal{KB}_{i-1}^{\Sigma}\downarrow_{\prec\delta_i}\cup\{\delta_i^{\mathrm{N}C}\})\cap\mathcal{M}\subseteq\mathcal{KB}_{all}^{\Sigma}$, so by Lemmas 1 and 2 and (34), the leftmost term entails $\mathrm{N}C\sqsubseteq\bot$ iff the rightmost does. It follows that the middle term $(\mathcal{KB}\cap\mathcal{M}_0)_{i-1}^{\Sigma}\downarrow_{\prec\delta_i}\cup\{\delta_i^{\mathrm{N}C}\}$ entails $\mathrm{N}C\sqsubseteq\bot$ iff the other two terms do. Then, $\delta_i^{\mathrm{N}C}$ is added to $(\mathcal{KB}\cap\mathcal{M}_0)_i^{\Sigma}$ iff $\delta_i^{\mathrm{N}C}$ belongs to $\mathcal{KB}_i^{\Sigma}\cap\mathcal{M}$ and \mathcal{KB}_i^{Σ} , and the Lemma follows using the IH.

As a consequence, the modularized construction is correct:

Theorem 1.
$$(\mathcal{KB} \cap \mathcal{M}_0)^{\Sigma} \models \alpha \text{ iff } \mathcal{KB}^{\Sigma} \models \alpha.$$

Proof. By Lemmas 1 and 2, and (34), $\mathcal{KB}^{\Sigma} \models \alpha$ iff $\mathcal{KB}^{\Sigma} \cap \mathcal{M} \models \alpha$. The Theorem then follows by Lemma 3.

⁴ Here the sets $(\mathcal{KB} \cap \mathcal{M}_0)_i^{\Sigma}$ are defined by replacing \mathcal{S} with $\mathcal{S} \cap \mathcal{M}_0$, and \mathcal{KB} with $\mathcal{KB} \cap \mathcal{M}_0$ in (4) and (5), while δ_i ranges over *all* the DIs in \mathcal{KB} , not just $\mathcal{D} \cap \mathcal{M}_0$. This formulation facilitates the comparison with \mathcal{KB}^{Σ} . By the condition $\delta_i \in \mathcal{KB}$ in (5) for all $\delta_i \notin \mathcal{M}_0$, $(\mathcal{KB} \cap \mathcal{M}_0)_i^{\Sigma} = (\mathcal{KB} \cap \mathcal{M}_0)_{i-1}^{\Sigma}$, so this definition is equivalent to building $(\mathcal{KB} \cap \mathcal{M}_0)^{\Sigma}$ using only $\mathcal{D} \cap \mathcal{M}_0$.

4 Optimistic computation

The construction of \mathcal{KB}^{Σ} repeats the concept consistency check (5) over a sequence of knowledge bases $(\mathcal{KB}^{\Sigma}_{i-1}\downarrow_{\prec\delta_i}\cup\{\delta_i^{NC}\})$ that share a (possibly large) common part \mathcal{KB}^{Σ}_0 , so incremental reasoning mechanisms help by avoiding multiple computations of the consequences of \mathcal{KB}^{Σ}_0 . On the contrary, the set of δ_j^{NC} may change significantly at each step due to the filtering $\downarrow_{\prec\delta_i}$. This operation requires many axiom deletions, which as already highlighted in [19], are less efficient than monotonically increasing changes. The optimistic algorithm introduced here (Algorithm 3) computes a knowledge base \mathcal{KB}^* equivalent to \mathcal{KB}^{Σ} in a way that tends to reduce the number of deletions, as it will be assessed in Sec. 5.

Phase 1 optimistically assumes that the DIs with the same priority as δ_i^{NC} do not contribute to entailing $NC \sqsubseteq \bot$ in (5), so they are not filtered with \downarrow_{δ_i} in line 3. Phase 2 checks whether the DIs discarded during Phase 1 are actually overridden by applying \downarrow_{δ_i} (lines 3 and 3). DIs are processed in non-increasing priority order as much as possible (cf. line 3) so as to exploit monotonic incremental classifications.

The following theorem shows the correctness of Alghorithm 3 in case the normality concepts do not occur in \mathcal{KB} , but only in the queries. We call such knowledge bases *N-free*. It is worth noting that the optimistic method is not generally correct when \mathcal{KB} is not N-free and $|\mathcal{L}| > 1$, yet it may still be applicable after the module extractor if the latter removes all normality concepts from \mathcal{KB} .

Theorem 2. If KB is N-free, then Algorithm 3's output is equivalent to KB^{Σ} .

Proof. First assume that Σ is a singleton $\{NC\}$. We start by proving some invariants of lines 6-10.

Claim 1: $\mathcal{KB}^{\Sigma} \models \Pi$.

 $\textbf{Claim 2} \text{: If, for some } j < i, \delta_j^{\mathrm{N}C} \in \mathcal{KB}^{\varSigma} \setminus \varPi \text{, then } \mathcal{KB}^{\varSigma} \models \mathrm{N}C \sqsubseteq \bot.$

We prove these two claims simultaneously. Both claims hold vacuously at the first execution of line 6. Next, assume by induction hypothesis that they hold at line 6 in some iteration; we have to prove that they still hold at the next iteration. There are two possibilities: First suppose that for some j < i, $\delta_j^{NC} \in \mathcal{KB}^{\Sigma} \setminus \Pi$. By Claim 2, $\mathcal{KB}^{\Sigma} \models NC \sqsubseteq \bot$. This immediately implies that Claim 2 holds also at the next iteration. Moreover, it implies Claim 1 because all members of Π have an occurrence of NC in the left-hand side, cf. (3).

We are left with the case in which

for all
$$j < i$$
, if $\delta_j^{\text{N}C} \in \mathcal{KB}^{\Sigma}$ then $\delta_j^{\text{N}C} \in \Pi$. (39)

If the condition in line 3 is true, then Π is not changed, so Claim 1 must hold at the next iteration. Otherwise, by (39),

$$\mathcal{KB}_0^{\Sigma} \cup \Pi' \supseteq \mathcal{KB}_{i-1}^{\Sigma} \downarrow_{\delta_i} \cup \{\delta_i^{NC}\}, \tag{40}$$

and hence $NC \sqsubseteq \bot$ is not provable in (5), either. It follows that δ_i^{NC} belongs to both Π (by line 3) and \mathcal{KB}^{Σ} (by (5)). This proves Claim 1 for iteration i.

Concerning Claim 2, first suppose that the condition in line 3 is true; then either Claim 2 remains vacuously satisfied, or $\delta_i^{\mathrm{N}C} \in \mathcal{KB}^{\Sigma} \setminus \Pi$. The latter (plus the def. of \mathcal{KB}^{Σ} and the ind. hyp. of Claim 1) implies that $\mathcal{KB}_0^{\Sigma} \cup \Pi'$ is entailed by \mathcal{KB}^{Σ} . It follows that $\mathcal{KB}^{\Sigma} \models \mathrm{N}C \sqsubseteq \bot$ as well, which proves Claim 2 in this case. Finally, if the condition in line 3 is false, then at line 3 $\delta_i^{\mathrm{N}C} \in \Pi$. Together with (39), this implies that Claim 2 holds vacuously.

Claim 3: If $\mathcal{KB}^{\Sigma} \models NC \sqsubseteq \bot$ then $\mathcal{KB}^* \models NC \sqsubseteq \bot$.

Suppose not (we shall derive a contradiction). The assumption and Claim 1 imply that there must be some $\delta_k^{NC} \in \mathcal{KB}^\Sigma \setminus \mathcal{KB}^*$. Let δ_i^{NC} be the one with minimal k. Using minimality, it can be proved that $\mathcal{KB}_0^\Sigma \cup \Pi \downarrow_{\delta_i} = \mathcal{KB}_{i-1}^\Sigma \downarrow_{\delta_i}$, so the concept consistency tests in lines 3 and 3 (the latter instantiated with j=i and D=C) are equivalent to the one in (5). But then $\delta_i^{NC} \in \mathcal{KB}^\Sigma$ iff $\delta_i^{NC} \in \mathcal{KB}^*$, which contradicts the assumption, so

Claim 4: If $\mathcal{KB}^{\Sigma} \models \mathcal{NC} \sqsubseteq \bot$ then $\mathcal{KB}^* \equiv \mathcal{KB}^{\Sigma}$.

Note that $\mathcal{KB}^+ \models \mathcal{NC} \sqsubseteq \bot$ then $\mathcal{KB}^- \models \mathcal{KB}^-$. Note that $\mathcal{KB}^* \subseteq \mathcal{KB}_0^{\Sigma} \cup \Pi \cup \{\mathcal{NC} \sqsubseteq \bot\}$ (cf. lines 11, 15, and 22). Clearly, $\mathcal{KB}^{\Sigma} \models \mathcal{KB}_0^{\Sigma} \cup \Pi \cup \{\mathcal{NC} \sqsubseteq \bot\}$ (by def., Claim 1 and the assumption), so $\mathcal{KB}^{\Sigma} \models \mathcal{KB}^*$. We are left to prove $\mathcal{KB}^* \models \mathcal{KB}^{\Sigma}$. By Claim 3, $\mathcal{KB}^* \models \mathcal{NC} \sqsubseteq \bot$, and this inclusion in turn entails all $\delta_i^{\mathcal{NC}} \in \mathcal{KB}^{\Sigma}$ (cf. (4)). The other members of \mathcal{KB}^{Σ} are those in \mathcal{KB}_0^{Σ} , by definition, and $\mathcal{KB}^* \supseteq \mathcal{KB}_0^{\Sigma}$ (line 11). It follows that $\mathcal{KB}^* \models \mathcal{KB}^{\Sigma}$, which completes the proof of Claim 4.

Claim 5: If $\mathcal{KB}^{\Sigma} \not\models NC \sqsubseteq \bot$ then $\mathcal{KB}^* \equiv \mathcal{KB}^{\Sigma}$. Suppose that $\mathcal{KB}^{\Sigma} \not\models NC \sqsubseteq \bot$. Then, by the contrapositive of Claim 2 and Claim $1, \mathcal{KB}_0^{\Sigma} \cup \Pi \equiv \mathcal{KB}^{\Sigma}$, which further implies that the concept consistency tests in lines 3 and 3 are equivalent to the corresponding test in (5). Then it can be proved that if any of these tests were true, then also $\mathcal{KB}^{\Sigma} \models \mathrm{N}C \sqsubseteq \bot$ because, by Π 's construction, in that case δ_i must be in conflict with some other DI with the same priority. However, $\mathcal{KB}^{\Sigma} \models NC \sqsubseteq \bot$ contradicts the assumption. It follows that all tests in lines 3 and 3 are false, so $\mathcal{KB}^* = \mathcal{KB}_0^{\Sigma} \cup \Pi$, and we have already argued that this knowledge base is equivalent to \mathcal{KB}^{Σ} . This completes the proof for $|\Sigma| = 1$.

For $|\Sigma| > 1$, note that the tests in lines 3 and 3 (resp. 3) do not depend on any $\delta_k^{\rm NE}$ such that $E \neq C$ (resp. $E \neq D$). Indeed, by \perp -locality, all such δ_k^{NE} are local w.r.t. the signature of $\mathcal{KB} \cup \{NC \sqsubseteq \bot\}$, so they can be removed without changing the result of the concept consistency test [26]. However, after their removal, the concept consistency tests correpond to the ones for the singleton case $\Sigma = \{NC\}$, which we have already proved correct.

Experimental assessment

Currently there are no "real" KBs encoded in a nonmonotonic DL, because standard DL technology does not support nonmonotonic reasoning. The nonmonotonic KBs encoded in the hybrid rule+DL system DLV-Hex [13] are not suited to our purposes because they do not feature default inheritance due to a restriction of the language: DL predicates cannot occur in rule heads, so rules cannot be used for encoding default inheritance. Consequently, synthetic test cases are the only choice for evaluating our algorithms. We

Algorithm 3: Optimistic-Method

```
Input: \mathcal{KB} = \mathcal{S} \cup \mathcal{D}, \Sigma
       Output: a knowledge base \mathcal{KB}^* such that \mathcal{KB}^* \equiv \mathcal{KB}^{\Sigma}
       // Phase 1
  1 compute a linearization \delta_1, \dots, \delta_{|\mathcal{D}|} of \mathcal{D}
  2 \Pi:=\emptyset // \Pi collects the prototypes
  3 \varDelta := \emptyset // ordered list of all discarded \delta_i^{\mathrm NC}
  4 for i=1,2,\ldots,|\mathcal{D}| do
               for NC \in \mathcal{D} do
  5
                         \Pi' := \Pi \cup \{\delta_i^{NC}\}
                         if \mathcal{KB}_0^{\Sigma} \cup \Pi' \not\models \mathrm{N}C \sqsubseteq \bot then
  7
                          \Pi := \Pi'
  8
  9
                         else
                            append \delta_i^{\mathrm NC} to \Delta
10
        // Phase 2
11 \mathcal{KB}^* = \mathcal{KB}_0^{\Sigma} \cup \Pi
12 while \Delta \neq \emptyset do
                extract from \varDelta its first element \delta_i^{{\rm N}C}
13
               if (\mathcal{KB}_0^{\Sigma} \cup \Pi) \downarrow_{\prec \delta_i} \cup \{\delta_i^{NC}\} \not\models NC \sqsubseteq \bot then \mathcal{KB}^* := \mathcal{KB}^* \cup \{NC \sqsubseteq \bot\} extract all \delta_k^{NE} with E = C from \Delta
14
15
16
17
               else
                         // \delta_i^{{\rm N}C} is actually overridden
18
                          \begin{array}{ll} \textbf{while } \Delta \ contains \ some \ \delta^{\mathrm{N}D}_j \ \ such \ that \ \delta \prec \delta_j \ \ \textbf{do} \\ & | \ \ \text{extract from } \Delta \ \text{the first such } \delta^{\mathrm{N}D}_j \end{array} 
19
20
                                  \begin{array}{l} \mathbf{if} \; (\mathcal{KB}_0^{\varSigma} \cup \varPi) \downarrow_{\prec \delta_j} \cup \{\delta_j^{\mathrm{N}D}\} \not\models \mathrm{N}D \sqsubseteq \bot \; \mathbf{then} \\ \big| \; \; \mathcal{KB}^* := \mathcal{KB}^* \cup \{\mathrm{N}D \sqsubseteq \bot\} \end{array}
21
22
                                            extract all \delta_k^{{
m N}E} with E=D from \varDelta
23
24
                                           \delta := \delta_i
```

start with the two test suites introduced in [7] as they have been proved to be nontrivial w.r.t. a number of structural parameters, including nonclassical features like exception levels and the amount of overriding. The two test suites are obtained by modifying the popular Gene Ontology $(GO)^5$, which contains 20465 atomic concepts and 28896 concept inclusions. In one test suite, randomly selected axioms of GO are turned into DIs, while in the second suite random synthetic DIs are injected in GO. The amount of strong axioms transformed into DIs is controlled by CI-to-DI-rate, expressed as the percentage of transformed axioms w.r.t. |GO| while the amount of additional synthetic DIs is controlled by Synthetic-DI-rate, i.e. the ratio $|\mathcal{D}|/|GO|$. The number of conflicts between DIs can be increased by adding an amount of random disjointness axioms specified by parameter DA-rate (see [7] for further details).

The experiments were performed on an Intel Core i7 2,5GHz laptop with 16 GB RAM and OS X 10.10.1, using Java 1.7 configured with 8 GB RAM and 3 GB stack space. Each reported value is the average execution time over ten nonmonotonic ontologies and fifty queries on each ontology. For each parameter setting, we report the execution time of: (i) the naive $\mathcal{DL}^{\mathbb{N}}$ reasoner of [7]; (ii) the optimistic method introduced in Sec. 4 (Opt); (iii) the module extraction method of Sec. 3 (Mod) using the module extraction facility of the OWLAPI; (iv) the sequential execution of Mod and Opt, i.e. Algorithm 3 is applied to $\mathcal{KB} \cap \mathcal{M}_0$. This combined method is correct by Theorem 2 and Theorem 1.

CI-to-DI	naive	opt	mod	mod+opt
	12.91			
	22.37			
	31.50			
	42.97			
25%	55.22	25.17	00.44	00.36

ı	Synth DIs	naive	opt	mod	mod+opt
Ī	05%	11.64			0.42
		21.66			0.67
		32.80			1.64
		41.51			1.42
L	25%	51.85	22.33	3.05	2.09

Table 2. Impact of $|\mathcal{D}|$ on performance (sec) – DA rate = 15% – priority (1)

Table 2 shows the impact of the number of DIs on response time for the two test suites, as DI rate ranges from 5% to 25%. The methods Mod and Mod+Opt are slightly less effective in the second suite probably because random defaults connect unrelated parts of the ontology, thereby hindering module extraction. In both suites, Opt's speedup factor (w.r.t. the naive method) is about two, while on average Mod is approximately 87 times faster in the first test suite (max. speedup 125), and 28 times faster in the second (max. speedup 35). On average, the combined method yields a further 13% improvement over Mod alone; the maximum reduction is 31% (2nd suite, Synthetic-DI-rate=25%, DA-rate=15%). The additional conflicts induced by injected disjointness axioms have moderate effects on response time (cf. Table 3). Mod+Opt's average response time across both test suites is 0.7 sec., and the longest Mod+Opt response time

⁵ http://www.geneontology.org

has been 2.09 sec. As a term of comparison, a single classification of the original GO takes approximately 0.4 seconds.

Test suite 1 (CI-to-DI)

	rest suite 1 (CI-to-DI)							
DA	naive	opt	mod	mod+opt				
	29.88							
	32.96			0.32				
	31.50			0.32				
	34.23							
25%	36.47	17.80	0.40	0.34				
30%	37.71	18.09	0.40	0.34				

Test suite 2 (Synth. DIs)

DA	naive	opt	mod	mod+opt
05%	28.20	12.63	0.99	0.84
10%	30.18	13.68	1.04	0.97
15%	32.80	14.90	1.11	1.06
20%	35.68	16.29	1.18	1.10
25%	37.46	17.02	1.25	1.15
30%	38.37	18.79	1.36	1.23

Table 3. Impact of DAs on performance (sec) – DI rate = 15% – priority (1)

CI-to-DI	naive	opt	mod	mod+opt
05%		05.74		
10%	52.82			00.28
15%	81.84	16.56	00.34	00.31
20%	133.62			00.33
25%	193.27	26.42	00.41	00.36

Synth DI	s naive	opt	mod	mod+opt
05%	12.76			
10%	23.72			0.86
15%	34.53			1.21
20%	44.92			1.96
25%	55.92	25.77	3.87	2.46

Table 4. Impact of $|\mathcal{D}|$ on performance (sec) – DA rate = 15% – priority (2)

Table 4 is the analogue of Table 2 given priority (2). With respect to priority (1), the computation time for \mathcal{KB}^{Σ} and query answering in the first test suite grows faster for the naive algorithm, while there are smaller differences for the optimized approaches (the reponse times of the combined approach are almost identical). In the second test suite, the performance of the naive algorithms decreases less dramatically, while the optimized methods seem slightly less effective than in the first test suite. In all cases, the speedups of Mod and Mod-Opt remain well above one order of magnitude. The performance as DAs grow has similar features (see Table 5).

The above test sets are N-free. We carried out a new set of experiments by randomly introducing normality concepts in DIs, within the scope of quantifiers. Specifically, $\exists R.C$ is transformed into $\exists R.N.C$. The response times of the naive algorithm and Mod^7 under priority (1) are listed in Table 6 for increasing values of $|\varSigma|$ (that is directly related to the amount of normality concepts occurring in \mathcal{KB}). We estimate that the values of $|\varSigma|$ considered here are larger than what should be expected in practice, given the specific role of explicit normality concepts, cf. footnote 6. Such values are also much

⁶ So far, all the applicative examples that are not N-free satisfy this restriction, as apparently the only purpose of explicit normality concepts is restricting default role ranges to normal individuals, cf. Ex. 12 and the nomonotonic design pattern in [7, Sec. 3.3].

⁷ In this setting Opt and Mod+Opt are not applicable, in general.

Test	suite	1 ((C)	[_t	n-	D.	n
1000	Suite	1 1		ייי	·U	◡.	_,

rest suite 1 (er to b1)							
DA	naive	opt	mod	mod+opt			
	84.53						
	90.38						
	91.84						
	92.93						
	93.54						
30%	96.37	19.49	0.38	0.33			

Test suite 2 (Synth. DIs)

		,	-	
DA	naive	opt	mod	mod+opt
05%	29.55	14.93	1.28	1.07
	30.81			
	34.54			
	36.79			
	40.86			
30%	43.35	18.74	1.79	1.34

Table 5. Impact of DAs on performance (sec) – DI rate = 15% – priority (2)

larger than in N-free experiments, where $|\varSigma|$ is bounded by the query size. Response times increase accordingly. In most cases, the naive algorithm exceeded the timeout. In the first test suite, Mod remains well below 1 minute; in the second suite it ranges between 100 seconds and 10 minutes. The reason of the higher computation times in the second suite is that the extracted modules are significantly larger, probably due to the random dependencies between concept names introduced by fully synthetic DIs.

	$ \Sigma $	50	100	150	200	250
Test suite 1						
ĺ	naive	1794.37	>30 min.	>30 min.	>30 min.	>30 min.
	mod	2.31	7.26	14.77	25.32	39.22
Test suite 2						
ſ	naive	>30 min.				
	mod	103.4	211.5	327.4	459.2	586.7

Table 6. Impact of normal roles values (sec) – DI rate = 25% DA rate = 15%

6 Conclusions

The module-based and optimistic optimizations introduced here are sound and complete, where the later applies only if the knowledge base is N-free. In our experiments, the combined method (when applicable) and the module-based method make \mathcal{DL}^N reasoning at least one order of magnitude faster (and up to \sim 780 times faster in some case). In most cases, optimized reasoning is compatible with real time \mathcal{DL}^N reasoning. This is the first time such performance is reached over nonmonotonic KBs of this size: more than 20K concept names and over 30K inclusions. Our approach brings technology closer to practical nonmonotonic reasoning with very large KBs. Only the random dependencies introduced by synthetic DIs, combined with numerous restrictions of role ranges to normal individuals, can raise response time over 40 seconds; in most of the other cases, computation time remains below 2 seconds.

⁸ Good results have been obtained also for KBs with ∼5200 inclusions under rational closure semantics [9,8].

We are currently exploring a more aggressive module extraction approach, capable of eliminating some of the normality concepts in Σ and related axioms. Besides improving performance over non-N-free KBs, a more powerful module extractor might enable the application of the combined Mod+Opt method to non-N-free \mathcal{DL}^N knowledge bases, by removing all normality concepts from \mathcal{KB} before Opt is applied.

We are also planning to adopt a different module extractor [23] that is promising to be faster than the OWLAPI implementation.

Last but not least, we are progressively extending the set of experiments by covering the missing cases and by widening the benchmark set, using real ontologies different from GO as well as thoroughly synthetic ontologies.

Acknowledgements: The authors would like to thank the reviewers for their valuable comments and suggestions. This work has been partially supported by the PRIN project Security Horizons.

References

- 1. F. Baader and B. Hollunder. Embedding defaults into terminological knowledge representation formalisms. *J. Autom. Reasoning*, 14(1):149–180, 1995.
- F. Baader and B. Hollunder. Priorities on defaults with prerequisites, and their application in treating specificity in terminological default logic. J. Autom. Reasoning, 15(1):41–68, 1995.
- P. A. Bonatti, M. Faella, and L. Sauro. EL with default attributes and overriding. In Int. Semantic Web Conf. (ISWC 2010), volume 6496 of LNCS, pages 64–79. Springer, 2010.
- 4. P. A. Bonatti, M. Faella, and L. Sauro. Adding default attributes to EL++. In W. Burgard and D. Roth, editors, *AAAI*. AAAI Press, 2011.
- P. A. Bonatti, M. Faella, and L. Sauro. Defeasible inclusions in low-complexity DLs. J. Artif. Intell. Res. (JAIR), 42:719–764, 2011.
- P. A. Bonatti, C. Lutz, and F. Wolter. The complexity of circumscription in DLs. J. Artif. Intell. Res. (JAIR), 35:717–773, 2009.
- P. A. Bonatti, I. M. Petrova, and L. Sauro. A new semantics for overriding in description logics. Artificial Intelligence, 222:1–48, 2015. Available online: http://www.sciencedirect.com/science/article/pii/S0004370215000028.
- 8. G. Casini, T. Meyer, K. Moodley, and R. Nortje. Relevant closure: A new form of defeasible reasoning for description logics. In E. Fermé and J. Leite, editors, *Logics in Artificial Intelligence 14th European Conference, JELIA 2014, Funchal, Madeira, Portugal, September 24-26, 2014. Proceedings*, volume 8761 of *Lecture Notes in Computer Science*, pages 92–106. Springer, 2014.
- G. Casini, T. Meyer, K. Moodley, and I. J. Varzinczak. Towards practical defeasible reasoning for description logics. In T. Eiter, B. Glimm, Y. Kazakov, and M. Krötzsch, editors, Description Logics, volume 1014 of CEUR Workshop Proceedings, pages 587–599. CEUR-WS.org, 2013.
- 10. G. Casini and U. Straccia. Rational closure for defeasible description logics. In T. Janhunen and I. Niemelä, editors, *JELIA*, volume 6341 of *Lecture Notes in Computer Science*, pages 77–90. Springer, 2010.
- 11. G. Casini and U. Straccia. Defeasible inheritance-based description logics. *J. Artif. Intell. Res. (JAIR)*, 48:415–473, 2013.
- 12. F. M. Donini, D. Nardi, and R. Rosati. Description logics of minimal knowledge and negation as failure. *ACM Trans. Comput. Log.*, 3(2):177–225, 2002.

- W. Drabent, T. Eiter, G. Ianni, T. Krennwallner, T. Lukasiewicz, and J. Maluszynski. Hybrid reasoning with rules and ontologies. In F. Bry and J. Maluszynski, editors, *Semantic Techniques for the Web, The REWERSE Perspective*, volume 5500 of *Lecture Notes in Computer Science*, pages 1–49. Springer, 2009.
- 14. B. F., C. D., M. D., D. NARDI, and P. S. P. The description logic handbook, theory, implementation, and applications (2nd edition). In *The Description Logic Handbook*, pages 555–555. Cambridge University Press, CAMBRIDGE, 2010.
- 15. S. Ghilardi, C. Lutz, and F. Wolter. Did I damage my ontology? a case for conservative extensions in description logics. In P. Doherty, J. Mylopoulos, and C. Welty, editors, *Proceedings of the Tenth International Conference on Principles of Knowledge Representation and Reasoning (KR'06)*, pages 187–197. AAAI Press, 2006.
- L. Giordano, V. Gliozzi, N. Olivetti, and G. L. Pozzato. A non-monotonic description logic for reasoning about typicality. *Artif. Intell.*, 195:165–202, 2013.
- L. Giordano, N. Olivetti, V. Gliozzi, and G. L. Pozzato. ALC + T: a preferential extension of description logics. *Fundam. Inform.*, 96(3):341–372, 2009.
- 18. B. C. Grau, I. Horrocks, Y. Kazakov, and U. Sattler. Modular reuse of ontologies: Theory and practice. *J. Artif. Intell. Res. (JAIR)*, 31:273–318, 2008.
- Y. Kazakov and P. Klinov. Incremental reasoning in EL+ without bookkeeping. In *Description Logics*, pages 294–315, 2013.
- R. Kontchakov, F. Wolter, and M. Zakharyaschev. Can you tell the difference between dllite ontologies? In G. Brewka and J. Lang, editors, *Principles of Knowledge Representation and Reasoning: Proceedings of the Eleventh International Conference, KR 2008, Sydney, Australia, September 16-19, 2008*, pages 285–295. AAAI Press, 2008.
- T. Lukasiewicz. Expressive probabilistic description logics. Artif. Intell., 172(6-7):852–883, 2008.
- C. Lutz, D. Walther, and F. Wolter. Conservative extensions in expressive description logics. In M. M. Veloso, editor, *IJCAI 2007, Proceedings of the 20th International Joint Conference on Artificial Intelligence, Hyderabad, India, January 6-12, 2007*, pages 453–458, 2007.
- 23. F. Martin-Recuerda and D. Walther. Axiom dependency hypergraphs for fast modularisation and atomic decomposition. In M. Bienvenu, M. Ortiz, R. Rosati, and M. Simkus, editors, Proceedings of the 27th International Workshop on Description Logics (DL'14), volume 1193 of CEUR Workshop Proceedings, pages 299–310, 2014.
- A. L. Rector. Defaults, context, and knowledge: Alternatives for OWL-indexed knowledge bases. In *Pacific Symposium on Biocomputing*, pages 226–237. World Scientific, 2004.
- 25. E. Sandewall. Defeasible inheritance with doubt index and its axiomatic characterization. *Artif. Intell.*, 174(18):1431–1459, 2010.
- U. Sattler, T. Schneider, and M. Zakharyaschev. Which kind of module should I extract? In B. C. Grau, I. Horrocks, B. Motik, and U. Sattler, editors, *Proceedings of the 22nd International Workshop on Description Logics (DL 2009), Oxford, UK, July 27-30, 2009*, volume 477 of CEUR Workshop Proceedings. CEUR-WS.org, 2009.
- R. Stevens, M. E. Aranguren, K. Wolstencroft, U. Sattler, N. Drummond, M. Horridge, and A. L. Rector. Using OWL to model biological knowledge. *International Journal of Man-Machine Studies*, 65(7):583–594, 2007.
- 28. T. Y. C. Woo and S. S. Lam. Authorizations in distributed systems: A new approach. *Journal of Computer Security*, 2(2-3):107–136, 1993.