Anti-Aliasing on the Web

Jasmine Novak
IBM Research, Almaden
jnovak@us.ibm.com

Prabhakar Raghavan
Verity, Inc
praghava@verity.com

Andrew Tomkins
IBM Research, Almaden
tomkins@almaden.ibm.com

ABSTRACT

It is increasingly common for users to interact with the web using a number of different aliases. This trend is a double-edged sword. On one hand, it is a fundamental building block in approaches to online privacy. On the other hand, there are economic and social consequences to allowing each user an arbitrary number of free aliases. Thus, there is great interest in understanding the fundamental issues in obscuring the identities behind aliases.

However, most work in the area has focused on linking aliases through analysis of lower-level properties of interactions such as network routes. We show that aliases that actively post text on the web can be linked together through analysis of that text. We study a large number of users posting on bulletin boards, and develop algorithms to *anti-alias* those users: we can with a high degree of success identify when two aliases belong to the same individual.

Our results show that such techniques are surprisingly effective, leading us to conclude that guaranteeing privacy among aliases that post actively requires mechanisms that do not yet exist.

Categories and Subject Descriptors

E.4 [**Data**]: Coding and Information Theory; G.3 [**Mathematics of Computing**]: Probability and Statistics; I.2.6 [**Artificial Intelligence**]: Learning; I.2.7 [**Artificial Intelligence**]: Natural Language Processing; I.5.3 [**Pattern Recognition**]: Clustering

General Terms

Algorithms, Measurement, Security, Theory

Keywords

Aliases, Pseudonyms, Alias detection, Privacy, Bulletin boards, Personas

1. INTRODUCTION

In this paper, we study the identification of unique users among a set of online pseudonyms, based on content analysis. Specifically, we consider the problem of reverse engineering the multiple aliases belonging to an individual based on text posted by those aliases in public fora such as bulletin boards, netnews, weblogs, or web pages.

WWW2004, May 17–22, 2004, New York, New York, USA.
ACM 1-58113-844-X/04/0005.

Beginning with Chaum's work on digital pseudonymous identities [3, 4], there has been significant work on mechanisms to provide aliases that cannot be penetrated or linked through examination of network transmission data. Systems such as Onion Routing [16], Crowds [17], Freedom [18], and LPWA [9] have moved network pseudonyms from an academic interest to a reality. However, even after all network artifacts have been removed, certain classes of pseudonyms remain vulnerable to identity detection [14]. This vulnerability stems from the fundamental role of participants in an online world: to provide value, the distinct pseudonyms must engage in interactions that are likely to be information-rich, and are hence susceptible to a new set of attacks whose success properties are not yet well understood.

On the other hand, research in economics and game theory has focused [8] on the social cost resulting from the widespread availability of inexpensive pseudonyms. A user can readily open a hundred email accounts at a free service such as Yahoo! or Hotmail, while masquerading under a hundred different identities at an online marketplace such as eBay or ePinions. The power of a system without enforcement mechanisms or registries comes with the potential for various forms of abuse.

Research in the field of *reputation networks* attempts to devise trust mechanisms that make it unlikely that a mountebank who preys on innocent people (say in an online marketplace) will garner the level of trust needed to command an effective economic premium for goods or services (i.e., if the same merchandise is sold by two individuals of different trust levels, the one with the higher trust value commands a higher price – this effect is visible in online marketplaces such as eBay).

We focus on aliases used by individuals in order to post on online bulletin boards. Users on these boards adopt multiple aliases for many different reasons. In some cases, an old alias has been banned by a moderator, or a password has been forgotten. In others, an old alias has lost the trust of the group, developed bad personal relationships with members of the group, or still exists, but requires an *alter ego* to support his arguments. Some users enjoy creating aliases that can take different sides, or can express different personalities (sometimes from the perspective of different genders). And some aliases allow a user to take on a reasonable or an extreme position in turn. Finally, of course, some users wish to express questionable or socially unacceptable views, or wish to discuss immoral or illegal activities.

Our main contribution is to establish that techniques from data mining and machine learning, when carefully tuned,

can be surprisingly effective at detecting such aliases, to the extent that our perception of the privacy afforded by aliasing mechanisms may be optimistic.

In the language of machine learning, we seek to "cluster" aliases into equivalence classes, where the aliases in a class are all deemed by the system to be the same user. A system to perform such a clustering must address two underlying problems. First, given some characterization of the content authored by an alias, and given new content, it must determine the likelihood that the alias produced the new content. And second, given such a primitive for computing likelihoods, it must determine the most appropriate clustering of aliases into authors. We show that in each case, algorithms tailored to the scope and nature of the domain perform significantly better than classical techniques.

1.1 Summary of Results

First, we studied several mechanisms for ranking a set of authors by likelihood of having generated a particular dataset. Given 100 authors (synthetically) split into 200 aliases, our best similarity measure ranks an author as most similar to her paired alias (out of 199 candidates) 87% of the time. In order to attain this result, we consider a number of different feature sets and similarity measures based on machine learning and information theory.

Next, we explore algorithms for clustering given this combination of features and the notion of similarity. For this we require a measure for comparing two clusterings, to evaluate how well our algorithm (as well as alternatives) perform relative to the ground truth. One of our contributions is the development of a clean 2-dimensional measure motivated by the concepts of precision and recall that are fundamental in information retrieval. We believe that for settings such as ours, this measure is more natural for evaluating the effectiveness of a clustering scheme than traditional measures.

On the previously mentioned set of 100 two-element clusters, we achieve the perfect cluster (i.e., contains all aliases, and no new aliases) 92% of the time if our clustering algorithm is terminated after it produces 100 clusters. We also give results for different distributions of cluster sizes and numbers of aliases.

Finally, we consider the problem of automatically stopping the clustering process at the "right" point – this would be the setting when we do not have a synthetic dataset (with a known target number of alias clusters). We present a clean and natural stopping condition for this problem that requires no outside parameterization. On the benchmark described above, this condition achieves a figure of merit within 2% of the optimal stopping point.

These results are attained using significantly less data per alias than other studies in the area, and achieve error rates that are substantially lower than other published reports.

1.2 Outline of Paper

We begin in Section 2 by covering related work. In Section 4 we develop our similarity measure capturing the likelihood that each author in the system generated a particular collection of text. Next, in Section 5 we describe combining the output of the similarity measure into a clustering of aliases representing the underlying individuals. In Section 6, we describe a case study moving from our analytical domain into a real world bulletin board with no planted clusters.

2. RELATED WORK

The field of author analysis, or *stylometrics*, is best known for its detailed analysis of the works of Shakespeare [2, 22], its success in resolving the authorship of the Federalist Papers [13], and its recent success in determining the author of the popular novel *Primary Colors*. Our problem is somewhat different – rather than determine which of several authors could have written a piece of text, we wish to extract from dozens of online identities a smaller set of underlying authors. Diederich et al. [6] used SVM classifiers to learn authorship of newspaper articles; de Vel et al. [5] and Tsuboi and Matsumoto [19] used the same technique for email messages. Argamon et al. [1] studied matching newspaper articles back to the newspaper. Krsul and Spafford [10] performed author identification on C programs rather than traditional documents, using features such as comments, indentations, case and so forth.

Most similar to our work, Rao and Rohatgi [14] study netnews postings and use PCA and a nearest-neighbor-based clustering algorithm to attain an almost 60% success rate at merging 117 users with two aliases each back into their original classes. They concluded that users are safe from anti-aliasing if fewer than 6500 words of text exist for the user. Our results in contrast indicate that for the data in our corpus, and the algorithms we develop, significantly less text enables significantly higher accuracy.

Grouping aliases into authors fits the general paradigm of agglomerative clustering algorithms [20, 21] from machine learning. Here one begins with a set of entities (say documents) each in its own cluster, then repeatedly *agglomerates* the two "closest" clusters into one – thereby diminishing the total number of clusters by one (at each agglomeration step). An important piece of this process is deciding when to halt the process of agglomeration; see [15, 11] for some discussion.

3. DATA

For our experiments, we gathered postings from message board archives on `http://www.courttv.com`. Posters on the CourtTV message boards tend to be highly engaged, posting frequently and emotionally, and use of multiple pseudonyms is quite common. We first crawled the homepage for the message boards `http://www.courttv.com/message_boards` to get a list of available topics. We then picked several topics of discussion: the Laci Peterson Case, the War in Iraq, and the Kobe Bryant case. Our selections were motivated by the volume of posting on the board, and our assessment of the likelihood that posters on the board would adopt aliases in their postings.

For each topic, we crawled the topic homepage to generate a list of thread URLs, which we then crawled to generate a list of pages of postings. We then crawled and parsed the postings pages. We broke the resulting content into individual posts, and extracted metadata such as the alias of the poster and the date and time of the post.

Our preliminary methodology to evaluate the effectiveness of our algorithms is to gather a large number of posts from a series of aliases, split each alias into smaller sub-aliases, then ask the algorithm to piece the sub-aliases back together (see [14] for an earlier application of this technique). Thus, careful data cleaning is extremely important to guarantee that the algorithm does not "cheat" by using, for example,

the signature at the end of each post as a highly discriminant feature. Thus, we removed signatures, titles, headers from inclusions of other messages, and any quoted content copied from other postings. When considering word counts and misspellings as features, we also removed any HTML tags and lowercased the remaining text. We then scanned a large number of postings by hand to verify that no additional hidden features based on the current alias remained.

Many users included emoticons, or "smilies" in their postings. These gif images were easily identifiable as they were included from a common directory (`http://board1.courrtv.com/smilies`). We included counts of usage of each type of smiley to our set of features. After computing the frequency of words, misspellings, punctuation and emoticons for each posting as described below, we accumulated the results to create a record of features for each user.

Each message board contained a large number of postings (323K postings on the Laci board at the time of our crawl), and a large number of users (3000 on the Laci board), so we had a range in the number of users and number of messages to use in our experiments. While our results improved as we analyzed more messages per alias, we sought to identify authors with a minimal amount of data. Except as noted, for all experiments cited in this paper we used 50 messages per alias totaling an average of 3000 words.

3.1 Terminology

In the following, we will use the term alias or pseudonym interchangeably. Conversely, we will use the term author to refer to the underlying individual, who may employ a single alias or several of them.

4. SIMILARITY MEASURES

Our goal in this section is to develop a *directed* similarity measure to capture the likelihood that the author of the text in one corpus would have generated the text in a second corpus. By directed we refer to fact that these likelihoods are not symmetric between pieces of text. In Section 5, we will use the resulting similarity measure to cluster aliases together.

We begin by considering the appropriate set of features to use. Next, we turn to algorithms for determining similarity given a fixed feature set.

4.1 Feature Extraction

As input from our data gathering process we are given a set of aliases, and for each alias, a set of posts with associated metadata. We refer to the collection of posts as the corpus for that particular alias.

From the corpus for each alias, we must now extract the features that we will use to characterize the corpus going forward. Stylometers argue that there are certain stylistic features that an author cannot choose to obscure, and that these features remain relatively constant from domain to domain for the same author. The two most commonly used feature classes in the stylometry literature are the following: first, numerical measures such as the mean and standard deviation of word, sentence, and paragraph length; and second, the distribution of words within a relatively small number of sets of *function words* (frequent content-free words whose occurrence frequency does not change much from domain to domain[1]). However, these features are typically used to ascribe very large collections of highly ambiguous text to a small number of possible authors (often 2 to 5), while our goal is to map much smaller amounts of more specific language to a much larger population of authors, so we must broaden the range of permissible features.

After some experimentation, we chose to model the feature set representing the corpus of a particular alias by the following four distributions:

Words: After detagging, words are produced by tokenizing on whitespace. We do fold all words into lowercase, but we do not perform stemming.

Misspellings: Words that are not in a large dictionary.

Punctuation: Any of the common punctuation symbols.

Emoticons: Any of the emoticon images allowed by CourtTV.

Function Words: Described above.

Figure 1 shows the results of an experiment comparing each of the four feature distributions. For this experiment, we took 100 aliases with at least 100 posts each from the Laci Peterson message board. We split each alias into two sub-aliases of 50 posts each, broken at random. For each of the resulting 200 sub-aliases, we applied the selected feature extractor to the alias. We then employed the KL similarity measure defined below; this is our best-performing measure, used here to benchmark the different feature sets. For each sub-alias a, we compute the similarity between a and the other 199 sub-aliases, and sort the results. The figure shows the probability that a's matching sub-alias was ranked first out of the 199 candidates.

There are a few messages to take away from the figure. First, words are clearly the most effective feature set, so much so that we have focused entirely on them for the remainder of the discussion. Second, it should be possible to extend our techniques to merge the different feature sets together gracefully, perhaps attaining an even higher overall result, but we have not taken this path. Third, our success probabilities are dramatically greater than in traditional stylometry: in that field, 90% probability of correctly identifying one of five authors given a large amount of text for the classification is considered an excellent result. This is due perhaps in part to our algorithms, but certainly largely due to the fact that personas on the web are much more distinguishable than William Thackeray and Jane Austen.

4.1.1 Other Features

There are a number of additional features that appear powerful, that were beyond our scope to analyze. These include: correlation of posting times; analysis of signature files; clustering of misspellings; references to entities such as people, locations, and organizations; expressed relationships such as daughter, husband, etc.; use of blank lines, indentations and other formatting cues; use of HTML tags such as fonts, colors, or links; and use of capitalization. A comprehensive treatment of these would entail augmenting our feature set with hidden markov models (for temporal

[1]Some examples might include: and, but, which, that, might, this, very, however, and punctuation symbols.

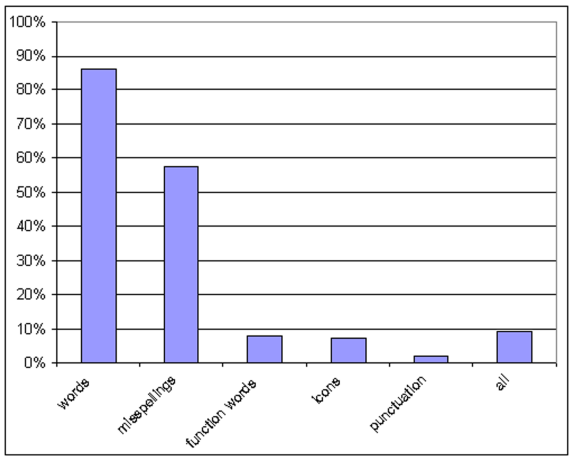

Figure 1: Evaluation of Different Feature Sets.

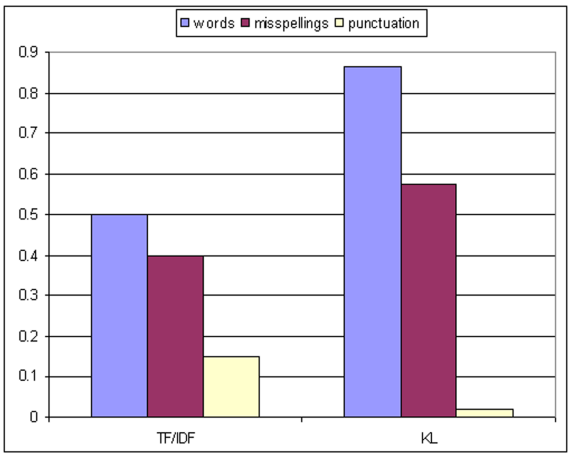

Figure 2: Evaluation of Similarity Algorithms.

features), link analysis (for references) and some entity extraction. Such detailed analyses would likely lower our already low error rates; our goal here is to demonstrate that even our simpler set of features suffice to viably jeopardize privacy derived from aliases.

4.2 Algorithms for Similarity

Let A be the set of all aliases, and let $n = |A|$. Let a be an alias, and p_a be the feature vector for the alias using the word feature set: each dimension in p_a represents a word, and the entry corresponds to the probability of that word in a's corpus, so $\sum_i p_a(i) = 1$. Next, let p_{bg} be the background distribution of word occurrences across all text from all aliases.

We now present three algorithms for similarity. We note that the similarity measure produced need not be symmetric, so $Sim(a, b)$ need not be $Sim(b, a)$. The interpretation is that $Sim(a, b)$ is the likelihood that the text in the corpus of alias a could have been produced by the author of the text for alias b.

4.2.1 Information Retrieval Similarity

This measure is based on the standard information retrieval cosine similarity measure [7]. We define v_a to be a vector corresponding to alias a with appropriate weighting for the measure, as follows: $v_a(i) = p_a(i)/p_{bg}(i)$.

The definition of the measure is:

$$Sim_{\text{TF/IDF}}(a, b) = \frac{v_a \cdot v_b}{|v_a| \cdot |v_b|}.$$

4.2.2 KL Similarity

The KL divergence of two distributions is defined as follows:

$$D(p||q) = \sum_i p_i \frac{\log p_i}{\log q_i}.$$

The KL divergence measures the number of extra bits that must be used to encode distribution p if instead of constructing the best possible code for p, we instead use the optimal code for distribution q. Intuitively, this seems to capture the notion that the author of q might have produced the text of p if that text can be encoded efficiently assuming it was in fact generated by q. The definition of the measure, then, is the following:

$$Sim_{\text{KL}}(a, b) = D(p_a||p_b)$$

where the distributions p_a and p_b are smoothed according to the discussion in Section 4.4, but the measure is computed only on non-zero elements of the original p_a.

This measure also has a probabilistic interpretation in our framework. Consider the corpus of alias $a \in A$ as a sequence of words $w_1, w_2, \ldots, w_n$. The probability that b would have generated that sequence in that order is simply $\Pi_{j\in[1..n]} p_b(w_j)$. Assuming that the corpus of a has size n, and the distinct words are given in W, then the number of occurrences of word i in the corpus is $np_a(i)$, and the total probability that b would generate the corpus is given by

$$\chi = \Pi_{i\in W} p_b(i)^{np_a(i)}.$$

Taking logs, this becomes $\log \chi = n \sum_i p_a(i) \log p_b(i)$. We observe that $D(p_a||p_b) = H(p_a) - \log \chi/n$. The terms $H(p_a)$ and n are both independent of b, so the ranking induced by maximizing the probability of b generating the corpus of a, over all b, is the same as the ranking induced by minimizing the KL divergence.

4.3 Results

Results for these algorithms are shown in Figure 2. They show that Sim_{KL} performs significantly better than does $Sim_{\text{TF/IDF}}$, so we will adopt the KL measure going forward. Using the benchmark described above, and using words as our feature, the algorithm ranks the correct alias first out of 199 possibilities 88% of the time.

4.4 Smoothing

In presenting Figure 2, we must mention one critical modification to the measure: that of *smoothing* the probability distribution defined by each feature vector. This is a standard concern whenever a probabilistic generative model is used to characterize content: how should we deal with two authors a and b who might be the same person, if the sample content from a doesn't use a particular word that b used (and vice versa)? A model without smoothing would assign zero probability to the event that a generated b's output.

There are a number of traditional approaches to smoothing; we use the simple approach of taking a linear combi-

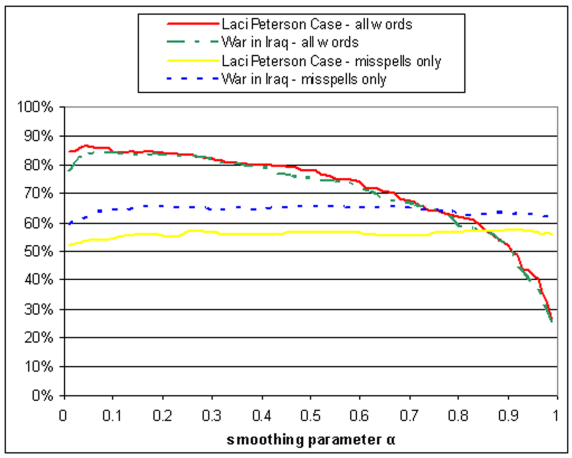

Figure 3: Evaluation of Smoothing Parameter α.

nation of α times a's distribution over features with $1-\alpha$ times the background distribution over features. The only parameter to study is therefore the weight α in the linear combination.

In the literature, values of α in the range of $[0.8, 0.9]$ seemed typical, and so we assumed that these values would be appropriate in our setting as well. To our surprise, the effectiveness of the algorithm increased as the smoothing parameter α dropped toward 0.8, and so we continued to examine its performance as we smoothed the distribution even more. Figure 3 shows the results for word-level features; smoothing with $\alpha = 0.02$ is optimal. Thus, the most effective technique for smoothing a distribution in our setting is to replace 98% of it with the generic background distribution!

The reason is the following. Due to the Zipf law on word usage in natural language [24, 23], each alias of an author will use many words that the other alias does not use. Each such word use in a naively smoothed distribution of large α will contribute a term of $(\alpha p_a(i) + (1-\alpha)p_{bg}(i))/(\alpha \cdot 0 + (1-\alpha)p_{bg}(i))$ to the measure; this term will be large as $p_{bg}(i)$ is tiny for such an infrequent term. Thus, if a particular alias used 17 highly infrequent terms, the most similar author will be the one who used as many of the 17 as possible, and normal differences in the frequencies of other terms would be swamped by this factor. By smoothing heavily, we allow the more frequent terms to contribute more strongly to the overall measure.

5. CLUSTERING ALGORITHMS

Having explored similarity measures, we now turn to clustering algorithms that make use of the entire family of directed similarities between aliases in order to determine meaningful clusters of aliases.

5.1 Definitions

Given a set of aliases A, we define a clustering $\mathcal{C} = \{C_1, \ldots, C_k\}$ over the aliases as any partition of the elements of A.[2] We define the "good" clustering $\mathcal{G}$ to be the correct clustering of the aliases. We define all aliases within the same cluster of $\mathcal{G}$ to be *siblings*. For an alias $a \in A$ we define $c(a)$ to be the cluster of a in $\mathcal{C}$, and $g(a)$ to be the cluster of a in $\mathcal{G}$.

[2]A partition has two properties: every alias belongs to at least one cluster, and no alias belongs to multiple clusters.

5.2 Measures

Our goal is to develop clustering algorithms; therefore settling on a measure to evaluate the quality of such an algorithm is of paramount importance. At its heart, such a measure must compare our proposed clustering to the correct one. However, there is no consensus on a single measure for this task, so we must spend some care in developing the correct framework.

Numerous measures have been proposed to compare clusterings, based typically on comparing how many pairs of objects are clustered together or apart by both clusterings, or by comparing pairs of clusters, or by adopting measures from information theory. The cleanest formulation of which we are aware is given by Meila [12], who proposes addressing many of the concerns with the above methods using a new measure called the *Variation of Information* (VI). Let H and I be the standard entropy and mutual information measures for distribution. Then for two clusterings $\mathcal{C}$ and $\mathcal{G}$, the VI is defined as follows:

$$VI(\mathcal{C}, \mathcal{G}) = H(\mathcal{C}) + H(\mathcal{G}) - 2I(\mathcal{C}, \mathcal{G}) = H(\mathcal{C}|\mathcal{G}) + H(\mathcal{G}|\mathcal{C}).$$

Some useful properties of this measure are:

1. VI is a metric
2. VI is scale-invariant in the sense that each point can be doubled without changing the measure
3. VI is linear in the sense that the VIs computed on induced clusterings of subsets of the points can be combined in the final VI
4. The value of VI is bounded above by the logarithm of the number of items.

Thus, VI is an attractive approach to measuring distance between clusterings, and we adopt it as such. However, the values of VI are difficult to interpret, so we would like to preserve the properties of the measure while allowing the reader to get a better feel for the results. We observe that a high-quality clustering has two properties:

- It places siblings in the same cluster
- It places non-siblings in different clusters.

An algorithm can perform well on the first measure by placing all aliases in the same huge cluster, or can perform well on the second measure by placing each alias in a distinct cluster. We seek algorithms that simultaneously perform well on both measures. We therefore adopt the following definitions. By analogy with the field of information retrieval, we define the *precision* P of a clustering to be the quality of the clustering with respect to property 1, as follows: $P(\mathcal{C}) = \sum_{a \in A} \Pr_{b \in c(a)}[g(b) = g(a)]$. Similarly, we define the *recall* R of a clustering to be the quality of the clustering with respect to property 2, as follows: $R(\mathcal{C}) = \sum_{a \in A} \Pr_{b \in g(a)}[c(b) = c(a)]$. Thus, precision captures the notion that the members of a's cluster are siblings, while recall captures the notion that the siblings of a are in the same cluster as a.

While these two figures of merit share intuition with the measures from information retrieval, they may behave quite

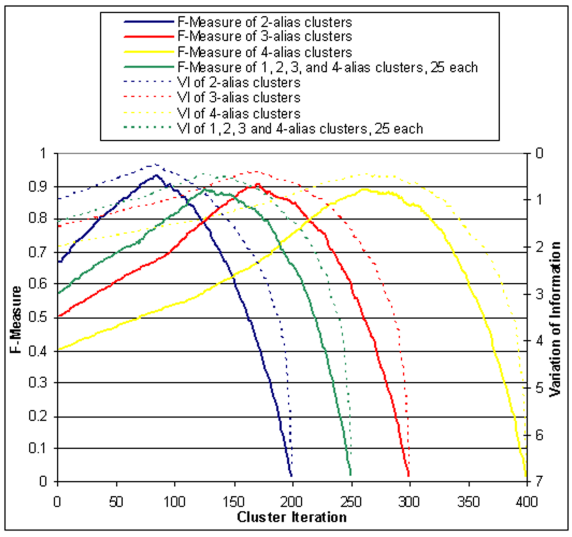

Figure 4: VI versus F-measure for Clusterings.

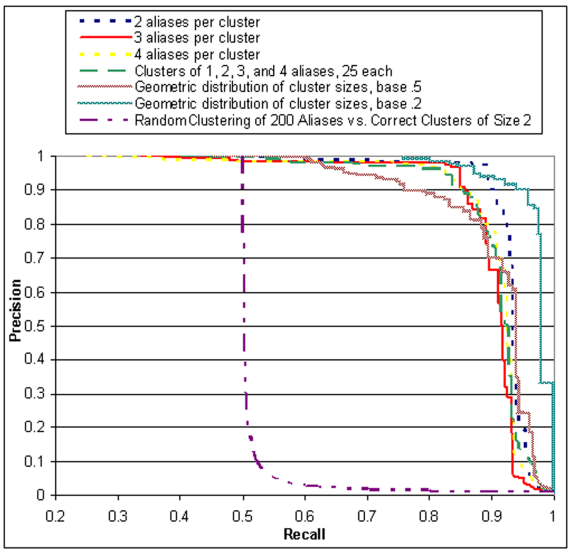

Figure 5: Clustering Using Mutual Ranking.

differently. Most importantly, precision and recall in the context of a clustering are inherently duals of one another in the sense that swapping $\mathcal{G}$ and $\mathcal{C}$ swaps the precision and recall values. Since our context always includes a correct clustering and an algorithmic clustering, we can use the terms with clarity; in general, though, they might more accurately be called the 1-recall (recall of $\mathcal{C}$ with respect to $\mathcal{G}$) and the 2-recall (recall of $\mathcal{G}$ with respect to $\mathcal{C}$), where the i-recall is the $(1-i)$-precision.

Finally, again following the terminology of information retrieval, we define the *F-measure* of a clustering $\mathcal{C}$ as $F = \max_{R,P} 2RP/(R+P)$, where R and P are recall and precision.

In Figure 4 we show in solid lines the F-measure for a number of experiments, and in dotted lines, the VI. The four experiments are described in detail below, but briefly, they cover domains in which the correct number of aliases per cluster is exactly 2, 3, or 4, or a mix of values between 1 and 4. The x axis measures the number of merge operations the clustering algorithm has performed, and the y axis shows the F-measure and VI of the resulting clustering—the scale for F-measure is shown on the left and for VI is shown on the right. As the figure shows, in all cases, the F-measure and the VI track quite closely over the entire range of number of merges performed by the algorithm. Thus, we conclude that F-measure captures the same quality indication as VI for our domain of interest. Henceforth, for clarity, we will adopt precision, recall and F-measure as appropriate for graphing results.

5.3 Mutual Ranking

Having established the measures we will use to evaluate our success, we now move to a limited variant of the clustering problem in which all clusters in the ground truth have size 2. We introduce the "Mutual Ranking" method for clustering, which we will later extend to a more general framework.

The clustering proceeds as follows. We are given a set of aliases A and a directed similarity measure $Sim(a,b)$ as defined in Section 4; the measure is larger (i.e., more similar) if the text of a could have been produced by the author of b. We define $r(a,b)$ to be the rank of b in the sorted list of $Sim(a,\cdot)$; thus, $r(a,b) = 1$ if and only if b is the most likely author (other than a) to have produced the text of a. Thus, $r(a,\cdot)$ is a permutation of $\{c \neq a | c \in A\}$. Mutual ranking iteratively pairs up elements of A greedily according to the measure $r(a,b) + r(b,a)$.

To benchmark this algorithm, we employed the same test set used to evaluate our different feature sets. Recall that we extracted 100 aliases from the Laci Peterson board who produced at least 100 articles, and split them into 200 sub-aliases of 50 posts each, broken at random. We then extracted features using each of our four different feature sets, and set $Sim = Sim_{\text{MLE}}$. We applied 100 steps of mutual ranking, and then measured how many of the resulting 100 clusters were "correct" in that they contained two sub-aliases of the same alias. The results are as follows:

Features	words	misspells	punctuation
Correct Clusters	91	66	12

Figure 5 shows how the precision and recall of the mutual ranking method on this benchmark change as the algorithm clusters more and more aliases. The "sweet spot" of the curve represents the 91 correct clusters shown in the table.

5.4 General Clustering

We now extend the mutual ranking framework to the general clustering problem. We define an interactive scheme for clustering aliases. The scheme, which we call *greedy cohesion*, is given by the following pseudo-code:

Let $\mathcal{C} = \{\{a\} | a \in A\}$ be the "current clustering"
Until stopping_condition(C):
    Pick $C_1, C_2 \in \mathcal{C}$ to minimize cohesion($C_1 \cup C_2$)
    Replace C_1 and C_2 in $\mathcal{C}$ with $C_1 \cup C_2$

The measure depends on the definition of the cohesion of a set of aliases; this is the mutual pairwise average ranking,

or more formally:

$$cohesion(C') = \frac{\sum_{a,b \in C'} r(a,b)}{|C'|(|C'|-1)}$$

Before we consider the stopping condition, we can evaluate how well the scheme chooses clusters to merge. We develop a number of benchmarks for this evaluation, as referenced in our discussion of clustering measures. The benchmarks are:

- Exactly 2, 3, or 4 aliases per cluster: For these three benchmark sets we consider authors who have produced 100, 150, or 200 posts respectively, and from these authors take the appropriate number of 50-post subsets.
- Mixed: This benchmark contains 25 clusters each of size 1, 2, 3, and 4.
- Exponentially decaying with factor $\gamma = 0.2$ or $\gamma = 0.5$: These two benchmarks contain $100(1-\gamma)^i$ clusters of size i, for $i \in [1..4]$.

The results are shown in Figure 4, and the F-measures for the optimal stopping points are shown in Table 1. As the figure shows, for both large clusters and highly skewed clusters, the algorithm performs quite well. Even for the extreme case of 400 authors and 4 aliases per cluster, the F-measure is still in excess of 0.88.

5.5 Comparison to Random

The precision-recall curves we have shown appear reasonable, but how are we to know that any random clustering scheme would not perform as well? We perform the following experiment: given a correct cluster size (for instance, 2), we allow a random clustering algorithm to know the correct size, and to choose a random clustering in which each cluster has size exactly 2. We have added the expected precision/recall of this scheme in Figure 5 to compare it with the actual algorithm. As the figure shows, the recall begins at 0.5 because each singleton cluster contains 1/2 of the siblings. However, until precision has dropped below 0.1, there is no visible improvement in recall—the number of cluster choices is too large, and as we would expect, the scheme performs horribly. Thus, any scheme that manages to move away from the convex form of the random algorithm is gaining some traction in identifying the correct clusters.

5.6 Other Clustering Algorithms

There are many other possible schemes to evaluate the next best cluster merge. A natural idea is to replace cohesion with inter-cluster distance (ICD), and pick the pair of clusters that minimize the inter-cluster distance. ICD is defined as follows:

$$\text{ICD}(C, C') = \frac{\sum_{a \in C, b \in C'} r(a,b)}{|C||C'|}.$$

We evaluated both schemes, and found that each produced similar results.

5.7 Stopping Condition

The agglomerative clustering algorithm defined above can continue merging until there is only one cluster left; this will improve recall at the cost of precision. We must devise a condition the algorithm can apply to terminate merging. We begin with two observations:

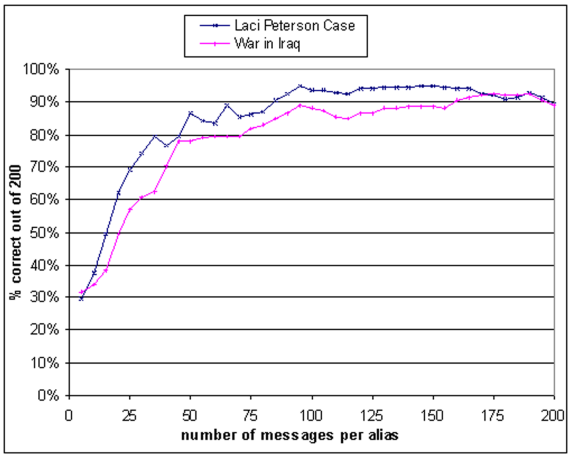

Figure 6: Evaluation of different sized data sets

OBSERVATION 1. *If a clustering has k clusters of size s then the cohesion of any given cluster is no smaller than $s/2$.*

Proof: Each a in the cluster can rank only 1 element first, one element second, and so on; its average rank will be $\left(\sum_{i=1}^{s-1} i\right)/(s-1) = s/2$. Likewise for all elements.

OBSERVATION 2. *If a clustering has k clusters whose average size is s, the average cohesion across all clusters cannot be less than $s/2$.*

Proof: By induction.

Thus, we adopt the following stopping condition. The algorithm stops when it cannot find a merge whose cohesion is within twice the best possible. Formally then, given a clustering problem with $|A| = n$ that has run for t steps, continue if and only if the best attainable cohesion is no more than $\left\lceil \frac{n}{n-t} \right\rceil$.

Table 1 shows the results of this stopping condition.

5.8 Using More User Data

Results using different sized data sets are shown in Figure 6. The figure again considers the running benchmark example of splitting the posts of 100 users into 200 aliases and attempting to re-group. The y axis plots probability of ranking the correct sub-alias top out of 199 candidates. As the figure shows, at 50 messages per alias the results become quite strong, as we have seen before, and as we move toward 100 or 125 messages, we sometimes attain probabilities of correct ranking in excess of 95%. The clustering algorithm typically improves on this probability noticeably.

6. REAL WORLD DATA

We now apply our clustering system in two real world experiments: clustering aliases using postings across multiple topics, and discovering non-synthetic multiple-aliases within a message board.

	Using Stopping Condition		Optimal	
Aliases per Cluster	F-measure	Iterations	F-measure	Iterations
2	.915	.89	.929	85
3	.899	166	.901	172
4	.873	253	.888	261
Mixed	.888	124	.890	125
Geo 0.2	.904	48	.929	37
Geo 0.5	.818	99	.843	135

Table 1: Evaluation of Stopping Condition.

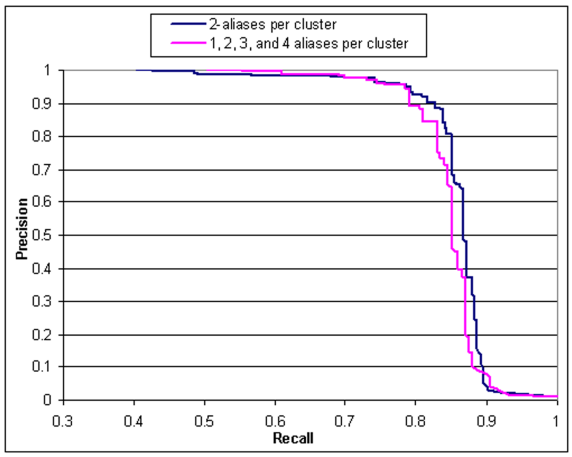

Figure 7: Evaluation of clustering multi-topic data

6.1 Heterogeneous topics

Postings on multiple topics present an additional challenge to clustering aliases. To investigate this problem we set up the following experiments:

- As before, we split the postings for 100 users into 200 sub-aliases. However, for this experiment, the postings for each sub-alias consisted of 25 postings from the Laci Peterson board, and 25 from the War in Iraq message board. We then tried to re-group the sub-aliases using our clustering algorithm.
- Using postings from 100 users, we created 25 clusters each of size 1, 2, 3, and 4 with half of each clusters' postings from the Laci Peterson board, and half from the War in Iraq board. Again we tried to re-group.

Since our algorithm depends on a user's vocabulary being consistent across postings, messages from the same user on different topics tend not be clustered together. To get around this problem, we devised the following scheme to discount vocabulary specific to a particular message board topic:

Let $t(w, b)$ be the tf-idf score for word w within message board b. Remove any word w where $|log[t(w, b)] - log[t(w, a)]| > 2$. For each remaining word w in message board b, if $t(w, a) < t(w, b)$, multiply each user's probability of using w by $t(w, a)/t(w, b)$.

After applying the above, we then used our Mutual Ranking method as previously described. Results from these experiments are shown in Figure 7.

6.2 Non-synthetic cluster discovery

As a final experiment, we applied our anti-aliasing system to a dataset without synthetically derived aliases. Using a fresh set of postings from several topics, we ran our clustering algorithm to identify users who were writing under multiple aliases. To evaluate our recall and precision, we scanned the postings for common patterns and clues that one alias was, or was not, the same user. We quickly realized that a complete analysis was infeasible, especially in identifying which clusters were missing. We opted instead for evaluating the accuracy of a subset of the enumerated clusters.

Using our system, we clustered 400 aliases on the Laci Peterson message board with between 10 and 200 postings each. Our algorithm resulted in 339 clusters, 56 of size two, one of size three, and one of size four. Seven of the clusters appeared to be correct based on the aliases alone, for example: "deli princess" and "deli princess2" or "Anita Mann" and "Anita Newman". These seem to be cases where the user is not attempting to disguise the use of multiple aliases. We then evaluated a sample of 12 non-obvious clusters for accuracy. By our judgment, 9 were correct, 2 were incorrect, and on one cluster we remained undecided.

As an example of the criteria we used in our evaluation, we present the following as evidence that Amadine and ButlerDidIt were indeed two aliases for the same user that were discovered by our algorithms:

- Both use "(s)" excessively. Examples:
  - Amadine:
    "Roxie, thanks very much for the link .... I've bookmarked it and I'm going to purchase the book(s) as soon as I get the chance."
  - ButlerDidIt: *"Scott could have called someone in his family (who live in San Diego) - his brother(s) or his father - and that person(s) could have met him halfway, transferred the body and disposed of it somewhere closer to San Diego - far away from Modesto."*
- Similar vocabulary. Examples:
  - Amadine:
    "Scott's subsequent behavior nailed the lid shut for me"
    "Now let's look at another scenario..."
  - ButlerDidIt:
    "It's hard to imagine such a conversation (and subsequent plan) playing out"
    "And in another scenario..."
- Both use numbered bullets. Examples:

– Amadine:
"I have a few concerns regarding Greta ...
1. Why does she have all these "teasers" to get viewers?
2. Why does she only have defense lawyers?
3. Why does the Fox Network think she's so great?
4. Why does she irritate me so much?"

– ButlerDidIt:
"I've been wondering if anybody knows...
1. If you were trying to weight a body down so it wouldn't float to the top, how much weight in proportion to the body would you need to make sure that it stayed under?...
2. As for prevention it from floating to the surface, if it doesn't go all the way to the bottom, is it likely (provided the weights stay intact) to float at some level below the surface...?
3. Assuming you had enough weight to get the body to the bottom of the ocean (or marina, lake, etc.) floor, is it still likely to get completely consumed by ocean life..."

Some of the criteria we used would be taken into account by our algorithm. For example, AvidReader and BoredInMichigan both dropped apostrophes in conjunctions, as in "dont", "didnt", and "wasnt". They also both used the expression "he, he, he" excessivley.

- AvidReader: *"He He He...sorry I have to intervene here..."*
- BoredInMichigan: *"He He He...Its working"*

Other criteria we used were more subjective. For example AuroraBorealis and dalmationdoggie made spelling mistakes we felt were similar:

- AuroraBorealis: embarrassing
- dalmationdoggie: interested, allot

Our analysis uncovered a drawback to our technique. Users who engage in an intense discussion on a slightly off-topic area, or who focus intently on the same side-topic for a series of posts, tend to get grouped together. One example is ColdWater and ClaraBella who are clustered together by our algorithm. ColdWater is the moderator for a particular thread; the bulk of this user's postings are answers to technical questions about how to use the software. ClaraBella's posts also address the technical aspects of the software as she answers questions posed by a new user of the message board. Addressing this anomaly appears to require a classic combination of traditional text classification methods with ours: while the former would focus on "content-rich" features that focused on the topic(s) of discussion, our methods would focus on features that are symptoms of a given author.

7. CONCLUSION

In this paper, we have shown that matching aliases to authors with accuracy in excess of 90% is practically feasible in online environments. Our techniques are most effective when two aliases of the same author post on the same bulletin board—there is significant cause for concern from a privacy perspective in this arena. Across bulletin boards, or even across sites, however, as the number of posts grows our techniques appear able to uncover aliases with an effectiveness that leads us to suggest that compromise of privacy is a very real possibility.

We have two areas of open problems. The first relates to the algorithm: how can it be improved, and can the techniques used for larger number of authors be applied meaningfully in the stylometric problem domain.

Our second area of open problems is broader but perhaps more critical. Are there meaningful countermeasures for techniques like ours? In particular, can users be given tools or training to make them less susceptible to such attacks? Our algorithms at present have not been optimized to run at web scale, but we have no reason to believe that scale alone will provide an adequate barrier. Our primary suggestion to users is to avoid behaviors that might allow algorithms to make rapid progress in bringing aliases together. Such behaviors would include posting on the same board, using a similar signature file, or mentioning the same people, places, or things. We would recommend avoiding characteristic language, but this is almost impossible to implement. Once a candidate alias has been discovered by a more advanced form of our system, techniques like correlation of posting times and analysis of evolution of discourse and vocabulary could be quite powerful, so in some ways there is safety in keeping personas apart.

But short of making it more difficult for programs to identify aliases, we do not have a suggestion for countering this type of technique, for users who will be entering non-trivial amounts of text under multiple personas which should be kept separate.

8. REFERENCES

[1] S. Argamon, M. Koppel, and G. Avneri. Routing documents according to style. In *Proceedings of First International Workshop on Innovative Information Systems*, 1998.

[2] B. Brainerd. The computer in statistical studies of William Shakespeare. *Computer Studies in the Humanities and Verbal Behavior*, 4(1), 1973.

[3] David Chaum. Untraceable electronic mail. *Communications of the ACM*, 24(2):84–88, February, 1981.

[4] David Chaum. Security without identification: Transaction systems to make big brother obsolete. *Communications of the ACM*, 28(10):1030–1044, October, 1985.

[5] Olivier Y. de Vel, A. Anderson, M. Corney, and George M. Mohay. Mining email content for author identification forensics. *SIGMOD Record*, 30(4):55–64, 2001.

[6] Joachim Diederich, Jrg Kindermann, Edda Leopold, and Gerhard Paass. Authorship attribution with support vector machines.

[7] W. Frakes and R. Baeza-Yates. *Information Retrieval: Data Structures and Algorithms*. Prentice Hall, 1992.

[8] E. Friedman and P. Resnick. The social cost of cheap pseudonyms. *Journal of Economics and Management Strategy*, 10(2):173–199, 2001.

[9] Eran Gabber, Phillip B. Gibbons, David M. Kristol, Yossi Matias, and Alain Mayer. On secure and

pseudonymous client-relationships with multiple servers. *ACM Transactions on Information and System Security*, 2(4):390–415, 1999.

[10] I. Krsul and E. H. Spafford. Authorship analysis: Identifying the author of a program. In *Proc. 18th NIST-NCSC National Information Systems Security Conference*, pages 514–524, 1995.

[11] Hang Li and Naoki Abe. Clustering words with the MDL principle. In *COLING96*, pages 4–9, 1996.

[12] M. Meila. Comparing clusterings. Technical Report 418, UW Statistics Department, 2002.

[13] F. Mosteller and D. Wallace. *Inference and Disputed Authorship: The Federalist*. Addison-Wesley, 1964.

[14] Josyula R. Rao and Pankaj Rohatgi. Can pseudonymity really guarantee privacy? In *Proceedings of the Ninth USENIX Security Symposium*, pages 85–96. USENIX, August 2000.

[15] Edie Rasmussen. *Clustering Algorithms*, chapter 16. Prentice Hall, 1992.

[16] M. Reed and P. Syverson. Onion routing. In *Proceedings of AIPA*, 1999.

[17] M. Reiter and A. Rubin. Anonymous web transactions with crowds. *Communications of the ACM*, 42(2):32–38, 1999.

[18] Zero Knowledge Systems, 2000.

[19] Yuta Tsuboi and Yuji Matsumoto. Authorship identification for heterogeneous documents. Master's thesis, Nara Institute of Science and Technology, 2002.

[20] Ellen M. Voorhees. *The effectiveness and efficiency of agglomerative hierarchic clustering in document retrieval*. PhD thesis, Cornell University, 1986.

[21] P. Willet. Recent trends in hierarchic document clustering: a critical review. *Information Processing and Manaagement*, 24(5):577–597, 1988.

[22] C. Williams. Mendenhall's studies of word-length distribution in the works of Shakespeare and Bacon. *Biometrika*, 62:207–212, 1975.

[23] G. U. Yule. *Statistical Study of Literary Vocabulary*. Cambridge University Press, 1944.

[24] G. K. Zipf. *Human Behavior and the Principle of Least Effort*. Addison-Wesley, 1949.