

Finding Algorithms in Scientific Articles

Sumit Bhatia[†], Prasenjit Mitra^{†*} and C. Lee Giles^{†*}

[†]Department of Computer Science and Engineering, ^{*}College of Information Sciences and Technology
The Pennsylvania State University
University Park, PA-16802, USA
sumit@cse.psu.edu, {pmitra,giles}@ist.psu.edu

ABSTRACT

Algorithms are an integral part of computer science literature. However, none of the current search engines offer specialized algorithm search facility. We describe a vertical search engine that identifies the algorithms present in documents and extracts and indexes the related metadata and textual description of the identified algorithms. This algorithm specific information is then utilized for algorithm ranking in response to user queries. Experimental results show the superiority of our system on other popular search engines.

Categories and Subject Descriptors:H.3.3[Information Search And Retrieval]Information Search and Retrieval – Search Process

General Terms: Algorithms, Experimentation, Design.

Keywords: Algorithms, Pseudocodes, Vertical Search Engines.

1. INTRODUCTION

Algorithms are ubiquitous in Computer Science literature and offer a precise methodology to solve problems. From sorting a few numbers to ranking billions of web pages, from Human Genome project to Economics and Internet, algorithms have influenced each and every aspect of human life [2]. Researchers are busy improving the current algorithms as well as developing new algorithms for new and unsolved problems. The current state-of-the-art search engines however, are not optimized to search for algorithms. They do not distinguish between documents that contain an algorithm and those that do not. As a result, the search results contain a variety of unwanted, irrelevant documents. For example, a query to search for algorithm for *SVM* will return lots of documents that contain *SVM* in them (application papers using *SVM* etc.) even though they do not offer any details about the algorithmic aspects of *SVMs*. Moreover, due to inappropriate ranking schemes the relevant documents might not show up in the top results.

In this paper, we describe a vertical search engine to search for algorithms in scientific documents. It first analyzes a document to check for the presence of an algorithm. If an algorithm is found, the document text is further analyzed to identify sentences that describe the algorithm. In addition, algorithm specific metadata from the document is also extracted and indexed. This algorithm specific information is then utilized to compute the relevance of algorithms to a given user query and the algorithms are presented in decreasing order of their relevance to the user.

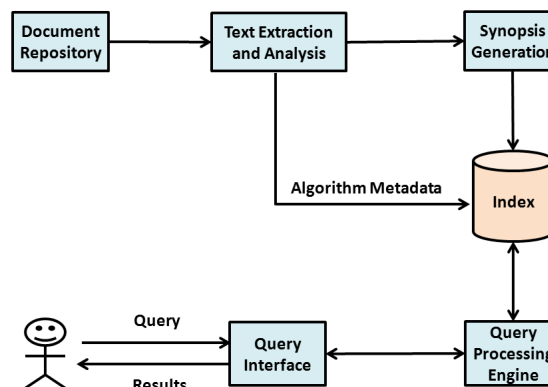


Figure 1: Architecture of the proposed system.

2. SYSTEM ARCHITECTURE

The main components of the proposed system are illustrated in Figure 1 and are described in more detail in the following sections.

2.1 Algorithm Extraction and Indexing

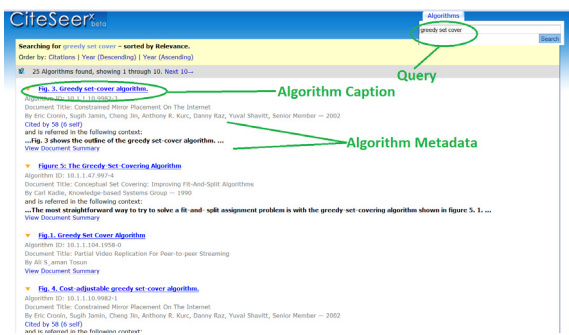
We use the collection of about 1.33 million computer science related documents in CiteSeer^{X1}. All the documents in the collection are in PDF/PostScript format and therefore need to be converted into text format for any further analysis. We experimented with a variety of text extraction tools available and found the performance of PDFTextStream² to be suitable for our needs. It performed best at preserving the sequence of text streams in the order they appeared in the document, especially for documents in double column format that are common in scientific literature.

The document text is then analyzed to check if the document contains an algorithm. Scientific documents in general have a well defined structure. Often algorithms/pseudocodes are described in the form of a stand alone Text-Box, Figure or Table, along with an associated caption and algorithm number. This *algorithm number* is then used to refer to the algorithm in the running text of the document. We call the sentences referring to the algorithm as a *reference sentences*. We can utilize these captions and reference sentences to check for the presence of algorithms in a document. If an algorithm is present, the document text is then further processed to extract the associated *synopsis* – a set of sentences that describes the

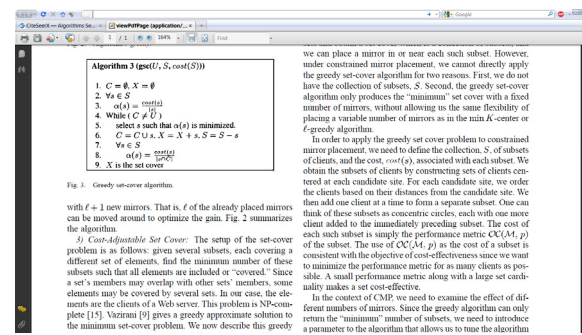
Copyright is held by the author/owner(s).
WWW 2010, April 26–30, 2010, Raleigh, North Carolina, USA.
ACM 978-1-60558-799-8/10/04.

¹<http://citeseerx.ist.psu.edu>

²<http://snowtide.com/PDFTextStream>



(a)



(b)

Figure 2: Screenshots showing (a) results for the query “Greedy Set Cover” and (b) algorithm page displayed on clicking the first result.

algorithm. For further details on algorithm identification and synopsis generation, the interested reader is directed to our previous work [1]. In total, we found 270367 unique algorithms in 112836 documents in the repository. For the documents containing an algorithm, we also extract the document title, author names, publication year and page on which the algorithm is present. We adopt the tools available from the SeerSuite toolkit for this purpose³. All the extracted algorithms from a document and their associated metadata are then indexed using an indexer based on SOLR⁴.

2.2 Query Interface

Our system offers a free text based query interface to the user and the results are returned along with the associated metadata. For each algorithm, a TF-IDF based cosine similarity score is computed between (i) (query, caption), (ii) (query, reference sentence) and (iii) (query, synopsis). The total similarity score for an algorithm is the sum of these three similarity scores. The algorithms are presented to the user in decreasing order of their scores thus obtained. The user interface is implemented using SeerSuite and query processing and ranking is implemented using SOLR. Figure 2(a) shows the screenshot of the result page for the query “Greedy Set Cover”. The top 10 algorithms for the query, along with their associated metadata are presented to the user. The algorithm caption is presented in bold and clicking on it directly takes the user to the PDF page of the concerned document on which the algorithm is present. This is illustrated in Figure 2(b).

3. EXPERIMENTS

We conducted a study using two human annotators to learn how the proposed system and other state-of-the-art search engines fare at satisfying information needs of the users looking for algorithms. We used a set of twenty queries (eg., *greedy set cover*, *graph isomorphism*) and tested them with Google Scholar, Google Web Search and the proposed system. For Google Web Search and Google Scholar, we added the keyword “algorithm” in hope to get more algorithms in the results. A result that returned a valid algorithm was considered as relevant. We report *precision @ 10* and *NDCG @ 10* for comparison.

We admit that this experiment can not be considered completely fair as the other search engines are not fine tuned to search for algorithms. However, we also note that due to the same reason, a user who generally turns to these state of the art, popular search engines to look for algorithms has to work harder to find the desired

³<http://citeseerx.sourceforge.net/>

⁴<http://lucene.apache.org/solr/>

	Precision@10	NDCG@10
Google Web Search	0.41	0.83
Google Scholar	0.44	0.58
Proposed System	0.81	0.94

Table 1: Comparison between the proposed system and other popular search engines.

documents containing relevant algorithms. The results summarized in Table 1 demonstrate the performance gains achieved by the proposed system as compared to other popular search engines. We achieve almost double the precision as compared to other methods as well as an appreciable gain in NDCG values indicating a better ranking of results.

4. CONCLUSIONS AND FUTURE WORK

We described a vertical search engine to search for algorithms in digital documents. A user study demonstrated the performance gains achieved by the system as compared to other popular search engines for algorithm search task. For future work, we plan to investigate various strategies for algorithm ranking. Oftentimes, authors discuss about pros and cons of the algorithms and generally, computational complexity of algorithms is also discussed. We plan to investigate how this information can be utilized for algorithm ranking.

5. ACKNOWLEDGEMENTS

This material is based upon work supported by the National Science Foundation under Grant Nos. 0535656 and 0845487. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the National Science Foundation.

6. REFERENCES

- [1] S. Bhatia, S. Lahiri, and P. Mitra. Generating synopses for document-element search. In *CIKM '09*, pages 2003–2006, 2009.
- [2] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. *Introduction to Algorithms*. MIT Press, Cambridge, MA, 2001.