

PYTHIA: a Framework for the Automated Analysis of Web Hosting Environments

Srdjan Matic
University College London
s.matic@cs.ucl.ac.uk

Gareth Tyson
Queen Mary University of London
gareth.tyson@qmul.ac.uk

Gianluca Stringhini
Boston University
gian@bu.edu

ABSTRACT

A common approach when setting up a website is to utilize third party Web hosting and content delivery networks. Without taking this trend into account, any measurement study inspecting the deployment and operation of websites can be heavily skewed. Unfortunately, the research community lacks generalizable tools that can be used to identify *how* and *where* a given website is hosted. Instead, a number of ad hoc techniques have emerged, e.g., using Autonomous System databases, domain prefixes for CNAME records. In this work we propose PYTHIA, a novel lightweight approach for identifying Web content hosted on third-party infrastructures, including both traditional Web hosts and content delivery networks. Our framework identifies the organization to which a given Web page belongs, and it detects which Web servers are self-hosted and which ones leverage third-party services to provide contents. To test our framework we run it on 40,000 URLs and evaluate its accuracy, both by comparing the results with similar services and with a manually validated groundtruth. Our tool achieves an accuracy of 90% and detects that under 11% of popular domains are self-hosted. We publicly release our tool to allow other researchers to reproduce our findings, and to apply it to their own studies.

CCS CONCEPTS

• **Networks** → *Public Internet*.

KEYWORDS

CDNs, infrastructure, topology, web

ACM Reference Format:

Srdjan Matic, Gareth Tyson, and Gianluca Stringhini. 2019. PYTHIA: a Framework for the Automated Analysis of Web Hosting Environments. In *Proceedings of the 2019 World Wide Web Conference (WWW '19)*, May 13–17, 2019, San Francisco, CA, USA. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3308558.3313664>

1 INTRODUCTION

When deploying a website, companies have the choice to either host it on their own servers, or to offload the content

to third-parties, e.g., Web hosts or Content Delivery Networks (CDNs). Choosing a third party can have multiple advantages, including cost savings, reliability, and the ability to sustain larger amounts of traffic (even during distributed denial of service attacks). We argue that understanding the Web hosting landscape is important for a number of reasons. These range from allowing us to assess how critical certain hosting infrastructures are and to estimate the impact that a network attack could have on the Web [8, 17, 24], to being able to determine who is responsible when incidents (e.g., malware hosting) occur [26]. Despite the importance of the problem, the research community lacks scalable methods to map the hosting landscape. Instead, multiple studies tend to take an ad hoc approach, relying on various assumptions, which differ between papers. Although there are third party services that offer this functionality [18, 30], they do not disclose their methodology, creating concerns for both reproducibility, as well as accuracy of their results.

To fill this gap, this paper *presents an open source tool to the community*,¹ which can determine whether the webpage of an organization is self-hosted or it uses a third-party hosting provider. Our tool, PYTHIA, leverages the HTML code of the webpage, domain information, and the network ownership obtained from RDAP records, to determine if the content is hosted on a third-party infrastructure. This is done by computing the ownership of both the webpage *and* the hosting provider, such that the two can be compared. PYTHIA is built with a modular design and it is capable of obtaining information of the landing webpage even in the presence of complex HTML structures which use redirects.

To evaluate the efficacy of PYTHIA, we run it on 40,000 URLs generated from the Alexa top-10k domains [2]. Our validation process shows that our framework outperforms similar applications available on the Web, and it achieves an accuracy of 90% in detecting when a webpage is hosted by a third party. Furthermore, our measurement reports that over 89% of the popular domains that we inspected, take advantage of third parties for their hosting needs.

PYTHIA is open source and allows the research community to reproduce our findings. We intend this to become a shared community effort, allowing third party researchers to avoid the complexity involved in devising and building their own independent methodologies for this commonly encountered task.

This paper is published under the Creative Commons Attribution 4.0 International (CC-BY 4.0) license. Authors reserve their rights to disseminate the work on their personal and corporate Web sites with the appropriate attribution.

WWW '19, May 13–17, 2019, San Francisco, CA, USA

© 2019 IW3C2 (International World Wide Web Conference Committee), published under Creative Commons CC-BY 4.0 License.

ACM ISBN 978-1-4503-6674-8/19/05.

<https://doi.org/10.1145/3308558.3313664>

¹The source code of PYTHIA is available at <https://bitbucket.org/srdjanmatic/pythia.git>

2 BACKGROUND

Understanding and measuring the Web hosting ecosystem is a complex endeavor. To complete this task, we need both information about the ownership of domains and the ownership of the IP addresses where webpages are hosted. In this section we introduce the concepts on which our approach is based, and the type of data that we retrieve to determine whether webpages are self-hosted or not.

2.1 Third-Party Hosting

Content Delivery Networks (CDNs) and Hosting Services are two popular mechanisms for delivering content to end users on behalf of other organizations. By offloading the task of serving the content of a website to third parties, these solutions are designed to provide better availability, scalability, faster content loads, redundancy, and enhanced security. These technologies have become so widespread that according to recent statistics, more than 60% of the most visited websites use CDNs to serve content to their users [4]. In this work we study the deployment of any kind of solution that delivers Web content for third parties, and for this reason we use the term *hosting* to refer to the *network where Web servers offering a service are based*.

There are countless papers that have explored the hosting patterns of websites, each taking a slightly different approach. A common approach is to launch large-scale distributed measurements [1, 9, 25], which perform DNS queries around the world to retrieve and classify DNS responses. This, unfortunately, is extremely complex and costly; furthermore, it cannot alone confirm if the infrastructure is third-party operated without further inspection. Calder et al. [6] utilized the EDNS-0 Client Subnet extension to simulate distributed queries towards Google's CDN. Although it revealed a large number of servers, all were operated by Google rather than third-parties. These techniques also do not work well for Anycast CDNs [7], which do not necessarily return DNS responses containing redirects. Another strategy employed is to utilize domain prefix lists, which map CNAME responses to their respective CDNs [21]. These, however, are limited to CDNs that exclusively rely on CNAME redirects (e.g., this excludes Bing). Furthermore, the list requires constant maintenance to remain up-to-date. Lastly, some studies utilize IP address to Autonomous System (AS) mappings [14] or metadata encoded into DNS records [3]; these, however are vulnerable to misattributing ownership, e.g., when a CDN places a cache in a third-party network. Such techniques have also been complimented with manually curated AS annotations, which stipulate the type of AS [19]. Again, these suffer from both manual annotation errors and require substantial upkeep. We argue that these diverse ad hoc techniques are driven by the lack of a standardized tool within the community, which can provide metadata on website hosting patterns.

2.2 The RDAP Protocol

To acquire information about a domain's ownership we use the Registration Data Access Protocol (RDAP) [22]. This

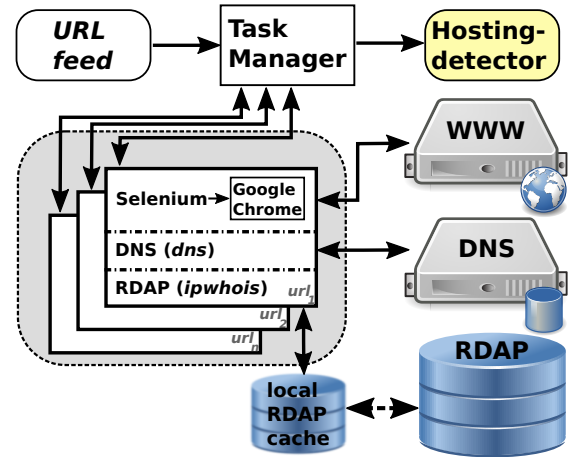


Figure 1: An overview of PYTHIA. The task manager processes URLs from a queue and instantiates a *crawler* task for each URL. The crawler collects information about a webpage and the environment from which the page was retrieved. The output of the crawler is stored in a dedicated data structure, which is later analyzed with the module that checks if content was served from a self-hosted network environment.

protocol was designed to replace the WHOIS [16] protocol as the authoritative source for registering information about IP addresses, ASes and domain names. While its predecessor retrieved free text content, RDAP leverages a RESTful interface to deliver the data in a machine-readable JSON format. This simplifies the parsing process and allows us to easily extract information (e.g., the type of entity to which a range of IP addresses has been assigned, or the description of an AS). In this paper, we use the RDAP protocol to retrieve information about the ownership of an IP address to which a domain resolved.

3 OVERVIEW OF PYTHIA

PYTHIA is entirely written in Python, and Figure 1 provides an overview of its system components. In a pipeline, PYTHIA performs the following two steps, (i) *Data Collection*: it takes a URL feed, and renders the list of websites, recording detailed information on all resources loaded, before complementing the data with RDAP records for the IP address hosting the website; and (ii) *Hosting Detector*: it then passes this information through a hosting detector, which decides if the owner of the webpage is the same as the owner of the hosting infrastructure where the server is located. The outcome is a structured JSON file that details if the website is self-hosted or operated by a third-party. In addition to this, PYTHIA also provides information on the ownership of both the webpage and the hosting service.

3.1 Data Collection

We first present the methods we use to collect the necessary data to infer ownership. This includes Web data, DNS information and, finally, RDAP records for all domains loaded.

Web Data Collection. Upon receiving a list of URLs, the crawler obtains Web content through Selenium,² a popular framework used for testing Web applications. We instrument Selenium to take a URL and to render it within a fully fledged instance of the Google Chrome browser. After the page has been loaded and rendered, our module outputs the retrieved HTML, in addition to a list of all the HTTP requests/responses (URLs) that were generated during the process. Each request/response is accompanied by metadata including the HTTP status code, and HTTP headers (e.g., server, content-type).

An important part of the crawling process is to determine when a webpage has finished loading. Since we do not know how long this process will take, we use an adaptive mechanism that leverages the information logged by the Web browser. We continuously monitor the browser logs and we consider the page loaded once all the requests to external resources have received their corresponding response. At the same time, we also set a hard timeout, after which we will close the browser session independently if the loading succeeded or not.

As part of the process, we also follow all the redirects that occur while loading an URL. This includes not just CNAME or HTTP redirects, but also those triggered by the refresh meta tag or a script. However, being primarily a tool for developers who need to check their own webpages, Selenium does not provide access to the browser internals. Hence, PYTHIA extracts information about the redirection chain from the browser log. After filtering out all the request to external resources, we identify the landing page and URLs on which the browser terminated the navigation.

Domain Resolution. Since Web browser logs do not contain information about domain resolution, PYTHIA launches DNS queries for all domains encountered (at a modest cost of only around 3.2% additional overhead per webpage). An advantage of this systems is that we can use our own DNS server and we do not need to rely on the built-in mechanisms of proprietary resolvers, which might be hardcoded in the browser.

IP Ownership. Using the DNS results, we then determine the “owner” of the IP space where the content is hosted. PYTHIA uses the RDAP protocol to find the network prefix to which an IP address belongs to, and to identify the owner of that range. Our framework uses a local RDAP cache to overcome rate-limiting issues of RDAP servers and to avoid querying ranges for which we already have fresh information. After successfully completing the RDAP resolutions, all the data generated by the above steps is stored in a JSON data structure.

²<https://www.seleniumhq.org/>

3.2 Detecting Hosting Infrastructures

The next step is to use the above information to detect if a website is self-hosted, or whether it uses on a third-party infrastructure. Our tool identifies the organization/company behind a domain name and a webpage, and searches for evidence that the page owner is the same as the owner of the network prefix or the AS hosting the Web server. We do not differentiate among various types of hosting services (i.e., VPS, CDN, or generic web hosting) and we do not use any precompiled list of popular or known hosting services. Instead we extract our information from the URL and the HTML retrieved from the landing page, and we match this data with the RDAP response of the IP address that is hosting the web server. This process allows us to detect the third-party network infrastructures even in the presence of CDN caches located at ISPs: even if we did not identify correctly the provider, our algorithm will detect a mismatch in the ownership of the webpage and the IP range.

To identify the organization that owns a webpage we use both the information from the URL and the HTML code. In particular, from the URL component we extract the Effective Second Level Domain (ESLD), and from the HTML we use the content of the <title> tag. Before retrieving any ownership information for a RDAP response, we first filter unnecessary details from the data such as “comments” or the “symbolic name of the network”, which can contain references to the owner of the webpage even when the IP range is assigned to a completely different organization. After this step, each string contained in the *HTML title* or the *RDAP fields* has its leading space delimiters removed, is cleaned from punctuation characters and stop words, converted into lower case, and finally is split into tokens on space delimiters. The DNS system does not allow domain names to contain space delimiters and it is common have domains, such as “bankofamerica.com”, where the ESLD is a combination of multiple words. To overcome this issue, the ESLD string follows the same cleaning process of the title and the RDAP, with the only difference in the tokenization, which is performed following the technique described in [23].

This process results in a series of string tokens that represent ownership features of both the webpage and the domain/IP address hosting it. The next step is to compare these tokens to see if they correspond. Our algorithm does six checks: four with the strings contained in the title/ESLD/RDAP, and an additional two with the tokenized versions of those strings. First, the algorithm verifies if the HTML title or the ESLD appears as a sub-string in any of the RDAP fields. Subsequently it repeats the same procedure with each string in the RDAP fields by comparing it both with the HTML title and the ESLD. The output of this process is binary and if the algorithm finds a match, it concludes that the owner of the webpage is also the owner of the network. As a final step, the algorithm checks for the presence of common tokens among the lists of tokens obtained from the HTML title/ESLD and the RDAP information. In this case a single match is not enough to conclude that the same

organization owns both the webpage and the network, and we require that the common tokens represent at least 50% of the overall number of tokens in the shortest list.

4 VALIDATION AND EVALUATION

PYTHIA is intended to be both accurate and straightforward to use for the community. To validate its capabilities we next run it over multiple datasets.

4.1 Validating the Data Collection

Our first goal is to test the efficacy of our crawler in collecting the necessary data to perform the hosting classification. Hence, we run PYTHIA over a series of URL lists. A first dataset includes 10,000 unique domain names obtained from a snapshot of the “Alexa top 10,000 websites” (TOP-10K) on 1st of May 2018. A second dataset (TOP-20K-WWW) is an extended version of the previous one, where domains are extended with the “www.” prefix. Finally, a third dataset (TOP-40K-URLs) includes all the entries from TOP-20K-WWW expanded with “http://” and “https://” prefixes.

Data Collection. We run PYTHIA over the TOP-40K-URLs to collect information about their home pages. We split our dataset into chunks of 350 elements and we process each chunk separately. Each element of the chunk is a unique domain name, which is crawled both with and without the “www.” prefix and with the two protocols that PYTHIA supports (HTTP and HTTPS). This means that when we successfully crawl an entire chunk, we obtain information for 1400 unique URLs. We refer to the initial URL from which we begin our crawling, and that we load in the browser, as “starting URL”; similarly, the URL on which the crawl terminates is called “landing URL”. Once a chunk is processed, PYTHIA waits for 120 seconds before switching to the next one. Each chunk is analyzed using 20 parallel instances of our Crawler module, which uses a maximum timeout of 60 seconds while waiting for a webpage to finish loading. Note that when collecting IP Ownership information from RDAP, we randomize waiting timeouts, with of a maximum of 90 seconds, before retrying a query that triggered an exception; after 3 consecutive exceptions, the module marks an IP as “no info available” before switching to the next one. The entire data collection process took place from a single machine, although we note it is possible to split the dataset and run parallel instances of PYTHIA on different machines.

Data Collection Performance. The overall process of downloading the HTML, resolving DNS names and collecting RDAP data, took 35 hours to complete for the TOP-40K-URLs. Figure 2 presents the Cumulative Distribution Function (CDF) of the increase of number successful crawls across time. To be able to use fresh entries from our local RDAP cache, we crawled the 4 starting URLs linked to each domain at the same time. Due to this choice, the CDFs of each “URL format” have very similar shapes. Hence *four* distributions in Figure 2 are almost stacked on the top of each other, and the red line connecting the values is the average across




Figure 2: Cumulative distribution functions of the successful crawls with different URL formats.

those distributions. After 24 hours, we had crawled only 50% of the URLs, and in the last 1/3 of the time, we obtained the information for the remaining half of the dataset. The dataset contained shuffled entries, which did not follow any ranking depending on the “popularity of a domain” and the likelihood to have an “unreachable URL” at the beginning or at the end the crawl is the same. We explain the spike in the increase of the number of downloads after the 24th hour with the presence of our RDAP cache. As we will later show, the majority of URLs/domains use a third-party hosting provider and as times passes we observe an increase of the number of RDAP queries which can be resolved with our local cache. Those local resolutions increase our crawling speed allowing us to allow us to gather information for the same amount of elements, in half of the time.

In Table 1 we summarize the results of the data collection process using PYTHIA. The first thing to notice is that 85% of the starting URLs are successfully reached. Overall, our crawler visited 1,736,929 external URLs, which were retrieved from 54,410 different domains. This suggests an average factor of 42 URLs per “starting URL”. It is therefore clear that each HTML page contains a considerable number of external resources, although it should be noted that this not only includes links to script and images, but also *redirects* from a starting to a landing webpage. Independently of the protocol and the presence of the “www.” prefix, the 6th column in the Table shows that redirects are extremely popular. On average we pass via 2.4 intermediate URLs before reaching a landing page. For this estimate we *only* use redirects that happen when loading a “starting URL” in our browser, and we excluded any redirects triggered by the external resources embedded in the HTML of the landing webpage. In general, redirects seem to be more popular for the HTTP protocol, but the average difference with HTTPS is minimal.

Related to the redirect phenomenon, we also notice that more than 60% of the crawls observed a “change of the domain” name among the “starting URL” and the “landing URL”. This happens with only 1/3 of that frequency value if we crawl domains without the “www.” prefix. The reason is that often the redirection happens from one domain to the same domain expanded with the “www.” prefix. A similar

Format	URLs (DOMAINS)									IPs		
	All	STARTING	CRAWLS				LANDING					
	Tot.	Tot.	COMPLETED	DOMAIN CHANGE	PROTOCOL CHANGE	AVG. REDIRECTS	NON-TRIGGERING EXCEPTIONS	SELF-HOST	3RD-HOST	Tot.	NON-TRIGGERING EXCEPTIONS	LANDING DOMAINS
http	874,574 (46,679)	10,000 (10,000)	9,204	6,114	6,701	2.6	8,968 (8,897)	941 (935)	8,027 (7,962)	34,332	33,728	8,178
http + www	885,035 (43,359)	10,000 (10,000)	9,280	2,606	6,627	2.4	9,033 (8,962)	977 (971)	8,056 (7,991)	32,496	31,973	8,214
https	745,939 (40,696)	10,000 (10,000)	8,215	5,147	580	2.3	8,037 (7,989)	877 (873)	7,160 (7,116)	30,640	30,217	7,353
https + www	794,449 (39,787)	10,000 (10,000)	8,673	2,229	491	2.2	8,462 (8,404)	918 (913)	7,544 (7,491)	30,177	29,781	7,683
All	1,736,929 (54,410)	40,000 (20,000)	35,372	16,096	14,399	2.4	13,940 (11,253)	1,559 (1,220)	12,381 (10,033)	38,092	37,492	9,188

Table 1: Results of running PYTHIA on the top-40k-URLs dataset. The values in the brackets indicates the unique number of elements for each entry.

trend is observable for the “change of the protocol”, when crawling URLs with HTTP (which get upgraded to HTTPS).

The overall number of unique IPs of the landing pages is slightly less than 10,000, and it reflects the fact that crawling the same domain with the four different formats, most of the time will lead to the same landing URL/domain. On average there are 1.22 domain per each landing IP (comparison of columns 7 and 12 in Table 1). This is explained by the presence of large hosting providers with many different customers. The same argument explains why we observe a similar relationship of 1.43 domains per each IP, when considering the dataset of all URLs. Finally, the results of PYTHIA indicate that 89% of the landing URLs are served from a third-party hosting infrastructure which does not belong to the owner of the webpage. In the following section we illustrate how we tested the accuracy of our classification, by using a manually validated groundtruth and by comparing with similar applications.

4.2 Classification Validation

We next validate the efficacy of our tool by compiling a groundtruth classification, and comparing it against PYTHIA.

Compiling a Comparative Dataset. To the best of our knowledge, no groundtruth dataset exists regarding Web hosting. To build this, we randomly select 324 domain names to manually annotate. These are taken from the TOP-10K dataset, crawled with the HTTP protocol. For each of these domains, we load the landing webpage in a browser and use search engines to check if the owner of the IP prefix is an organization offering Web hosting or CDN services to its customers.

We note that 324 domains are not enough to evaluate our tool. Thus, we also collect equivalent data from a variety of public tools that allow users to “discover who is hosting a website”. This allows us to compare our results against their outputs. Example of those services include HOSTING-COMPASS.COM which can detect who is hosting an ESLD, or HOSTINGDETECTOR.COM and WHAT’S MY CDN? which allow more fine grained queries including “www.” as prefix to the ESLD [11, 12, 31]. We choose to use those three application because they are free Web-based services that do not require any registration. We query those services with the URLs from our TOP-10K and TOP-20K-WWW datasets,

TOOL/SERVICE	SELF-HOSTING		3RD-PARTY HOSTING		F1-SCORE
	TP	FN	TP	FN	
PYTHIA	26	3	279	16	0.73
hostingcompass	27	2	102	193	0.21
hostingdetector	12	17	239	56	0.25
whatsmycdn	27	2	127	168	0.24
cdnfinder	28	1	65	230	0.2

Table 2: Performance comparison of PYTHIA and other applications on our manually validated groundtruth.

depending on the service. As the services mentioned above do not provide any detail about the methodology they use to detect hosting providers, in addition to those Web applications, we also use CDNFINDER, an open-source project which aims to detect the usage of CDNs within websites. The tool uses PHANTOMJS and a hard-coded list of hostnames to load a webpage and detect the presence of external resources which are hosted on a CDN [20, 28]. Analogously to the Web services, we downloaded the tool and run it on our TOP-40K-URLs dataset. In total, this results in 5 datasets to compare PYTHIA against.

Comparison with Manual Annotations. Table 2 contains the results of comparing PYTHIA against the above online services, using the manual annotations as the groundtruth. Our algorithm was specifically designed for detecting the presence of “self-hosting” environments. Consequently, a domain will be flagged as “hosted on third-parties” in any situation where the webpage owner differs from the owner of the network prefix (e.g., a private Web server run at at home, where the broadband ISP is the owner of the network prefix). Despite this limitation, in the binary classification problem where a website is either self-hosted or hosted on a third-party service, PYTHIA still outperforms all the other services that we tested, achieving an F1-score which is almost three times larger than the average for the other services. Indeed, on our manual groundtruth we observe an accuracy of over 95%, even when, instead of verifying self-hosting, we focus on the complementary problem of detecting the presence of third-party hosting providers. CDNFINDER performs well in detecting self hosting, but has a very high false positive rate when classifying domains as third-party hosting. Similarly, HOSTINGDETECTOR achieves the highest accuracy in detecting external hosting

SERVICE	DOMAIN	WWW. + DOMAIN	HTTP + DOMAIN	HTTPS + DOMAIN	HTTP + WWW. + DOMAIN	HTTPS + WWW. + DOMAIN
HOSTINGCOMPASS	3,481 (3,283)	-	-	-	-	-
HOSTINGDETECTOR	3,229 (2,879)	5,607 (4,973)	-	-	-	-
WHATSMYCDN	978 (963)	3,268 (3,202)	-	-	-	-
CDNFINDER	-	-	59 (55)	403 (395)	149 (144)	1,849 (1,700)

Table 3: Comparison of PYTHIA with similar services/applications when evaluated on all of our datasets. The values in the brackets indicates results obtained by PYTHIA.

services, at a cost of an extremely high false negative rate (59%), when a single organization is in control of both the webpage and the network prefix.

Comparison with Similar Services. To further test the accuracy of our framework, we compare our results with the four tools/services mentioned earlier. The results of this comparison are shown in Table 3. The goal of this is to show that PYTHIA achieves similar results to other applications. To this end, we narrow our goal to identify all domains which are hosted on third-party network infrastructures. As mentioned in the previous sections, PYTHIA follows any kind of redirect. Hence, Table 1 only presents the classification using the “landing URLs/domains”. Since we do not know what are the exact capabilities of the four services that we tested, and how they handle redirects, we decided to back-propagate the results of our classification from the “landing URL” to the corresponding “starting URL”, and “starting domain”, from where the navigation started. In this way we are able to compare our results with each one of those services and verify that our framework has a detection rate close to those of the other services.

For almost all of the domains inspected, PYTHIA achieves an accuracy of around 90%, and it identifies a third-party-hoster every time one of the other four services detects its presence. Since the highest number of misclassified services originate from the sets of domains analyzed with HOSTINGDETECTOR, we sampled 20 domains without the “www.” prefix and another 20 with the “www.” prefix. We then manually verify if they are actually hosted on a third-party infrastructure. For 27 out of 40 cases, HOSTINGDETECTOR failed to identify self-hosted domains and PYTHIA correctly labeled those as “self-hosting”. Ten of those cases were domains of large universities with their own network prefixes. For another 7 cases, the landing page is the home page of large hosting services such as “Google”, “Salesforce” or “1and1”. PYTHIA correctly labeled these as self-hosted. For four domains our PYTHIA did not succeed in downloading the RDAP information, and PYTHIA could not classify those domains. The remaining 9 domains were hosted on a third-party infrastructure, but we did not detect them. According to those results, we conclude that the accuracy of PYTHIA is inline with similar services we compared to.

5 RELATED WORK

A significant amount of research has been done in the field of Content Delivery Networks and cloud computing. Krishnamurthy et al. [15] were the first to analyze the rise of CDNs and the benefits that they provide to end-users. After

them several studies investigated this trend [1, 7, 13, 25]. Similar work has tried to uncover cloud usage patterns and which Web services are running on a cloud-associated IP address [10, 29].

Our techniques relies on a mix of methodologies, particularly exploiting RDAP data. There have been a small set of past papers that rely on similar data. For example, Cai et al. proposed to combine WHOIS information with the ASN in order to generate a comprehensive AS-to-organization mapping [5]. Tajalizadehkhoob et al. were the first ones to explore the identification of hosting providers by combining passive DNS with WHOIS information [27]. Unfortunately their approach leverages a classification of 2,000 ASes to filter out organization such as ISPs, education and government. This list has limited size and is manually generated, and this raises concerns about its reliability across time. Contrary to previous studies, our work does not use any precompiled list of organization names and it focuses on identifying self-hosting environments. PYTHIA allows other researchers to reproduce our results and it does not require any manual analysis or a priori knowledge of network prefixes or the ASes in charge for routing the network traffic.

6 CONCLUSION

In this work we presented PYTHIA, a tool for collecting information about a webpage and the environment where the page is hosted. Our framework extracts information from the retrieved HTML, the DNS and the ownership information associated to a network prefix. PYTHIA then exploits this data to infer if the website is self-hosted, or is reliant on a third party operator, e.g., a Content Delivery Network. We tested PYTHIA on 40,000 URLs and compared the results with similar applications that detect the presence of known hosting providers. Our framework is accurate and outperforms all other applications, when tested on a manually validated groundtruth. PYTHIA is released as open source and is built in a modular way, which gives the possibility to integrate it with new capabilities and extensions.

7 ACKNOWLEDGEMENTS

This research was supported by the PETRAS IoT hub (EP-SRC grant EP/N023242/1) All opinions, findings and conclusions, or recommendations expressed in this material are those of the and do not necessarily reflect the views of the sponsors.

REFERENCES

- [1] B. Ager, W. Mühlbauer, G. Smaragdakis, and S. Uhlig. Web content cartography. In *Proceedings of the 2011 ACM SIGCOMM conference on Internet measurement conference*, pages 585–600. ACM, 2011.
- [2] Alexa. Top 1M Sites, 2018. <https://www.alexa.com/topsites>.
- [3] T. Böttger, F. Cuadrado, G. Tyson, I. Castro, and S. Uhlig. Open connect everywhere: A glimpse at the internet ecosystem through the lens of the netflix cdn. *SIGCOMM Comput. Commun. Rev.*, 48(1), 2018.
- [4] BuiltWith. DN technologies Web Usage Statistics. <https://trends.builtwith.com/cdn>, 2018. [Online; accessed 14-May-2018].
- [5] X. Cai, J. Heidemann, B. Krishnamurthy, and W. Willinger. Towards an AS-to-Organization Map. In *ACM Internet Measurement Conference*, 2010.
- [6] M. Calder, X. Fan, Z. Hu, E. Katz-Bassett, J. Heidemann, and R. Govindan. Mapping the expansion of google's serving infrastructure. In *Proceedings of the 2013 conference on Internet measurement conference*, pages 313–326. ACM, 2013.
- [7] M. Calder, A. Flavel, E. Katz-Bassett, R. Mahajan, and J. Padhye. Analyzing the performance of an anycast cdn. In *Proceedings of the 2015 Internet Measurement Conference*, pages 531–537. ACM, 2015.
- [8] A. Delignat-Lavaud and K. Bhargavan. Network-based origin confusion attacks against https virtual hosting. In *Proceedings of the 24th International Conference on World Wide Web*, pages 227–237, 2015.
- [9] R. Fanou, G. Tyson, P. Francois, and A. Sathiaselan. Pushing the frontier: Exploring the african web ecosystem. In *Proceedings of the 25th International Conference on World Wide Web*. International World Wide Web Conferences Steering Committee, 2016.
- [10] K. He, A. Fisher, L. Wang, A. Gember, A. Akella, and T. Ristenpart. Next stop, the cloud: Understanding modern web service deployment in ec2 and azure. In *Proceedings of the 2013 conference on Internet measurement conference*, pages 177–190. ACM, 2013.
- [11] HostingCompass. We analyze ip address and web hosting! <https://www.hostingcompass.com>.
- [12] HostingDetector.com. Web Hosting Lookup - Find out who is hosting any website. <https://hostingdetector.com/>.
- [13] C. Huang, A. Wang, J. Li, and K. W. Ross. Measuring and evaluating large-scale cdns. In *ACM IMC*, volume 8, pages 15–29, 2008.
- [14] D. Ibosiola, B. Steer, A. Garcia-Recuero, G. Stringhini, S. Uhlig, and G. Tyson. Movie pirates of the caribbean: Exploring illegal streaming cyberlockers. *AAAI International Conference on Web Blogs and Social Media (ICWSM)*, 2018.
- [15] B. Krishnamurthy, C. Wills, and Y. Zhang. On the use and performance of content distribution networks. In *Proceedings of the 1st ACM SIGCOMM Workshop on Internet Measurement*, pages 169–182. ACM, 2001.
- [16] D. L. RFC 3912, whois protocol specification. Technical report, IETF — Internet Engineering Task Force, 2004.
- [17] J. Liang, J. Jiang, H. Duan, K. Li, T. Wan, and J. Wu. When HTTPS Meets CDN: A Case of Authentication in Delegated Service. In *IEEE Symposium on Security and Privacy*, 2014.
- [18] NetCraft. Hosting Provider Server Count. <https://www.netcraft.com/internet-data-mining/hosting-provider-server-count/>, 2018. [Online; accessed 24-May-2018].
- [19] A. Noroozian, M. Korczyński, C. H. Gañán, D. Makita, K. Yoshioka, and M. van Eeten. Who gets the boot? analyzing victimization by ddos-as-a-service. In *International Symposium on Research in Attacks, Intrusions, and Defenses*, pages 368–389. Springer, 2016.
- [20] phantomjs.org. PhantomJS - Scriptable Headless Browser. <http://phantomjs.org/>.
- [21] S. Quirin, J. Jonas, H. Oliver, C. Luca, and C. Gorg. Structure and Stability of Internet Top Lists. In *Proceedings of the 19th International Conference on Passive and Active Measurement*. Springer, 2018.
- [22] H. Scott and K. N. RFC 7482, registration data access protocol (rdap) query format. Technical report, IETF — Internet Engineering Task Force, 2015.
- [23] T. Segaran and J. Hammerbacher. *Beautiful data: the stories behind elegant data solutions*. "O'Reilly Media, Inc.", 2009.
- [24] M. Simenovski, G. Pellegrino, C. Rossow, and M. Backes. Who controls the internet? analyzing global threats using property graph traversals. In *Proceedings of the 2017 Web Conference*, 2017.
- [25] A.-J. Su, D. R. Choffnes, A. Kuzmanovic, and F. E. Bustamante. Drafting behind akamai: Inferring network conditions based on cdn redirections. *IEEE/ACM Transactions on Networking (TON)*, 17(6):1752–1765, 2009.
- [26] S. Tajalizadehkhoob, C. Gañán, A. Noroozian, and M. v. Eeten. The role of hosting providers in fighting command and control infrastructure of financial malware. In *Proceedings of the 2017 ACM on Asia Conference on Computer and Communications Security*, pages 575–586. ACM, 2017.
- [27] S. Tajalizadehkhoob, M. Korczyński, A. Noroozian, C. Gañán, and M. van Eeten. Apples, oranges and hosting providers: Heterogeneity and security in the hosting market. In *Network Operations and Management Symposium (NOMS), 2016 IEEE/IFIP*, pages 289–297. IEEE, 2016.
- [28] TurboBytes. GitHub - turbobytes/cdnfinder: Webapp and cli-tool to detect CDN usage of websites. <https://github.com/turbobytes/cdnfinder>.
- [29] L. Wang, A. Nappa, J. Caballero, T. Ristenpart, and A. Akella. WhoWas: A Platform for Measuring Web Deployments on IaaS Clouds. In *ACM Internet Measurement Conference*, 2014.
- [30] WebHosting. <https://webhosting.info/>. <https://webhosting.info>, 2018. [Online; accessed 24-May-2018].
- [31] whatsmycdn.com. What's My CDN? <http://www.whatsmycdn.com/>.